

# Large-Scale Visual Recognition With Deep Learning

Marc'Aurelio Ranzato

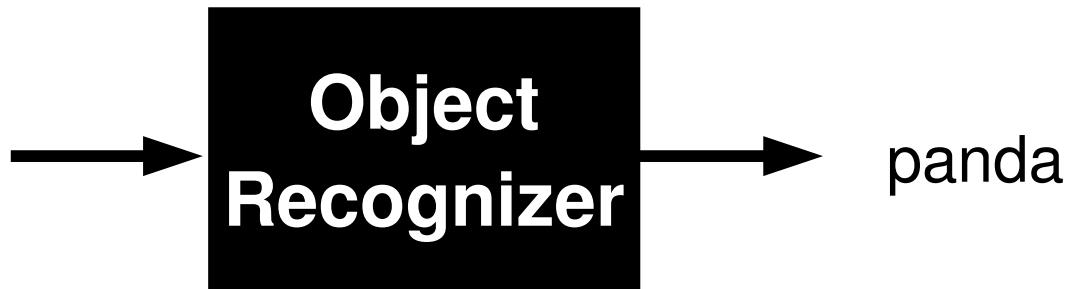


[ranzato@google.com](mailto:ranzato@google.com)

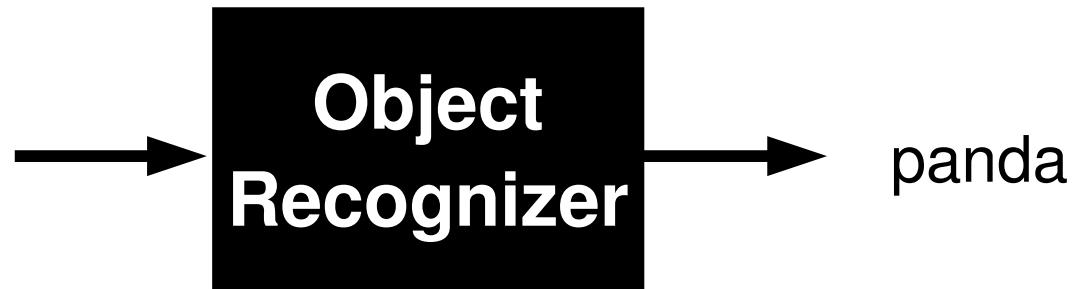
[www.cs.toronto.edu/~ranzato](http://www.cs.toronto.edu/~ranzato)

*Sunday 23 June 2013*

# Why Is Recognition Hard?



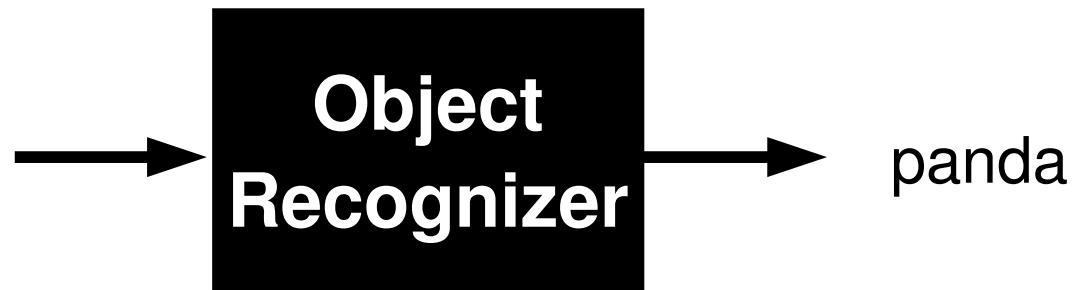
# Why Is Recognition Hard?



Pose



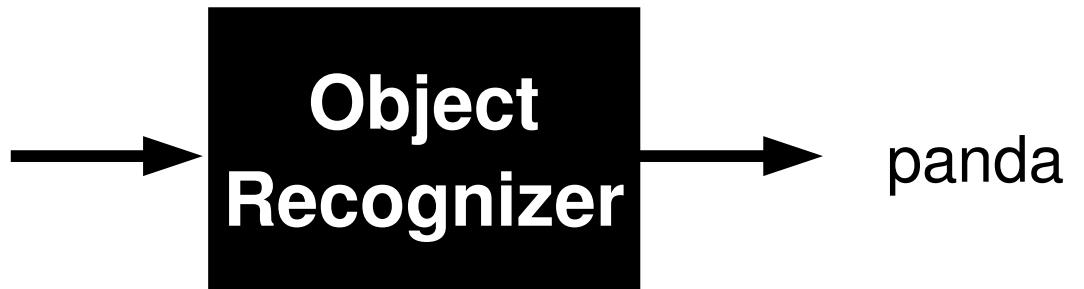
# Why Is Recognition Hard?



Occlusion



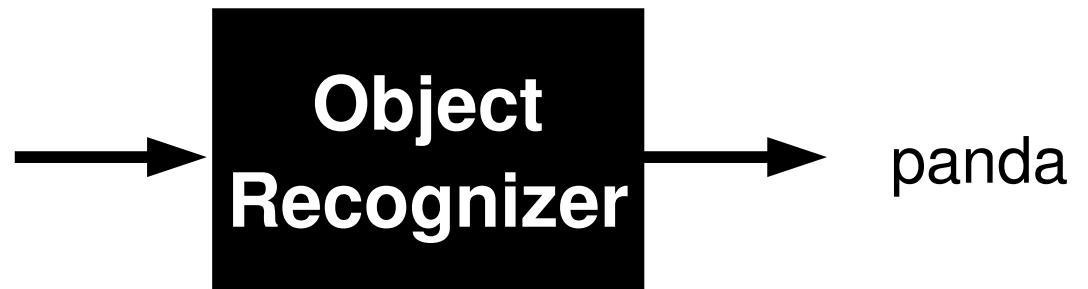
# Why Is Recognition Hard?



Multiple  
objects



# Why Is Recognition Hard?



Inter-class  
similarity



# Ideal Features

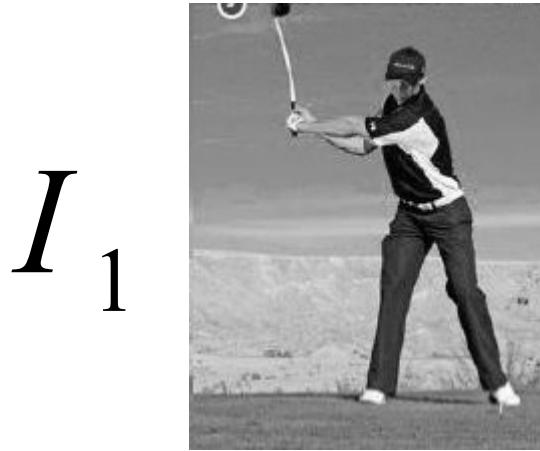


**Ideal  
Feature  
Extractor**

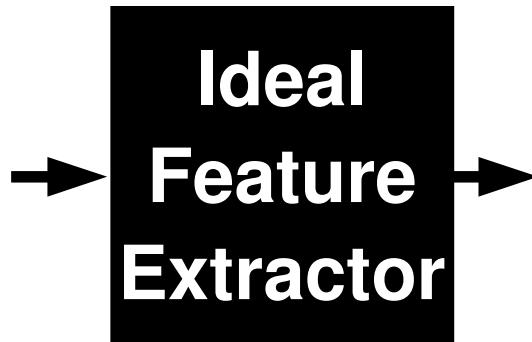
- window, top-left
- clock, top-middle
- shelf, left
- drawing,middle
- statue, bottom left
- ...
- hat, bottom right

**Q.:** What objects are in the image? Where is the clock?  
What is on the top of the table? ...

# Ideal Features Are Non-Linear

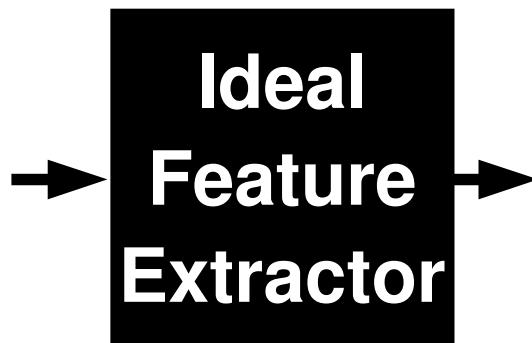


$I_1$



- club, **angle = 90**
- man, frontal pose
- ...

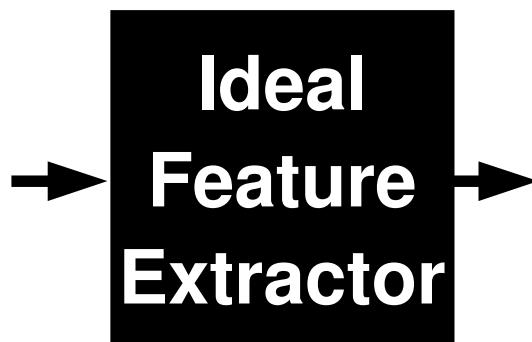
?



- club, **angle = 270**
- man, frontal pose
- ...



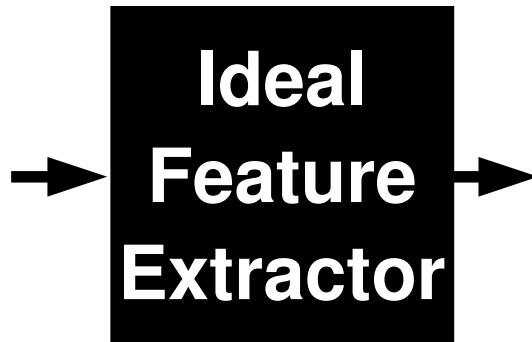
$I_2$



- club, **angle = 360**
- man, side pose
- ...

# Ideal Features Are Non-Linear

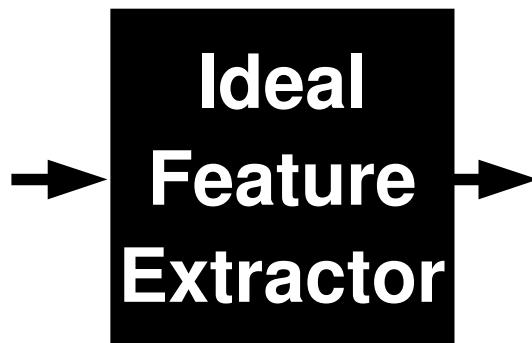
$I_1$



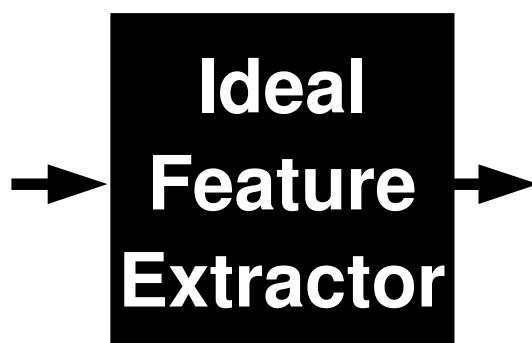
- club, **angle = 90**
- man, frontal pose
- ...

**INPUT IS  
NOT THE  
AVERAGE!**

$I_2$

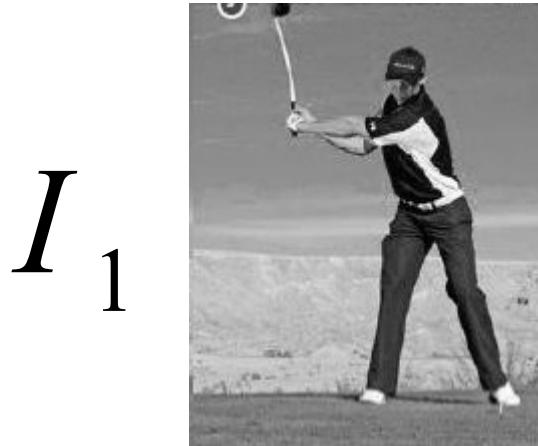


- club, **angle = 270**
- man, frontal pose
- ...

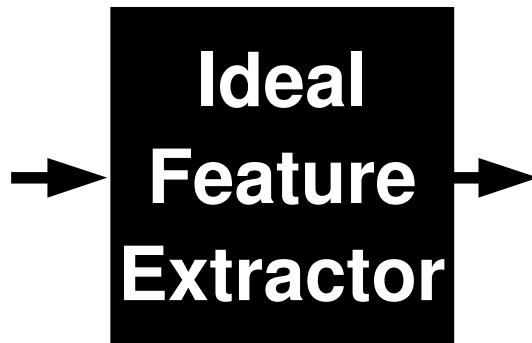


- club, **angle = 360**
- man, side pose
- ...

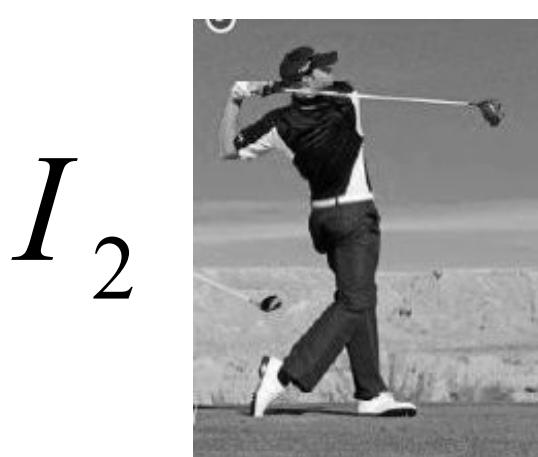
# Ideal Features Are Non-Linear



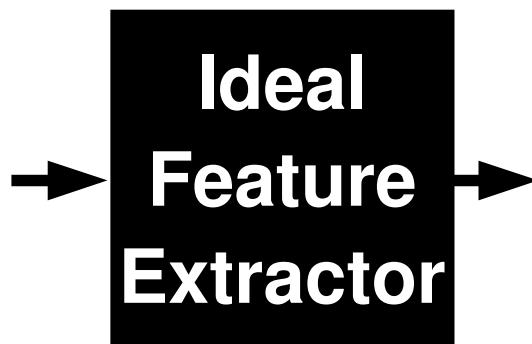
$I_1$



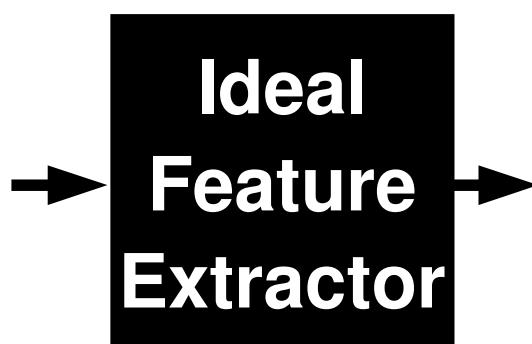
- club, **angle = 90**
- man, frontal pose
- ...



$I_2$

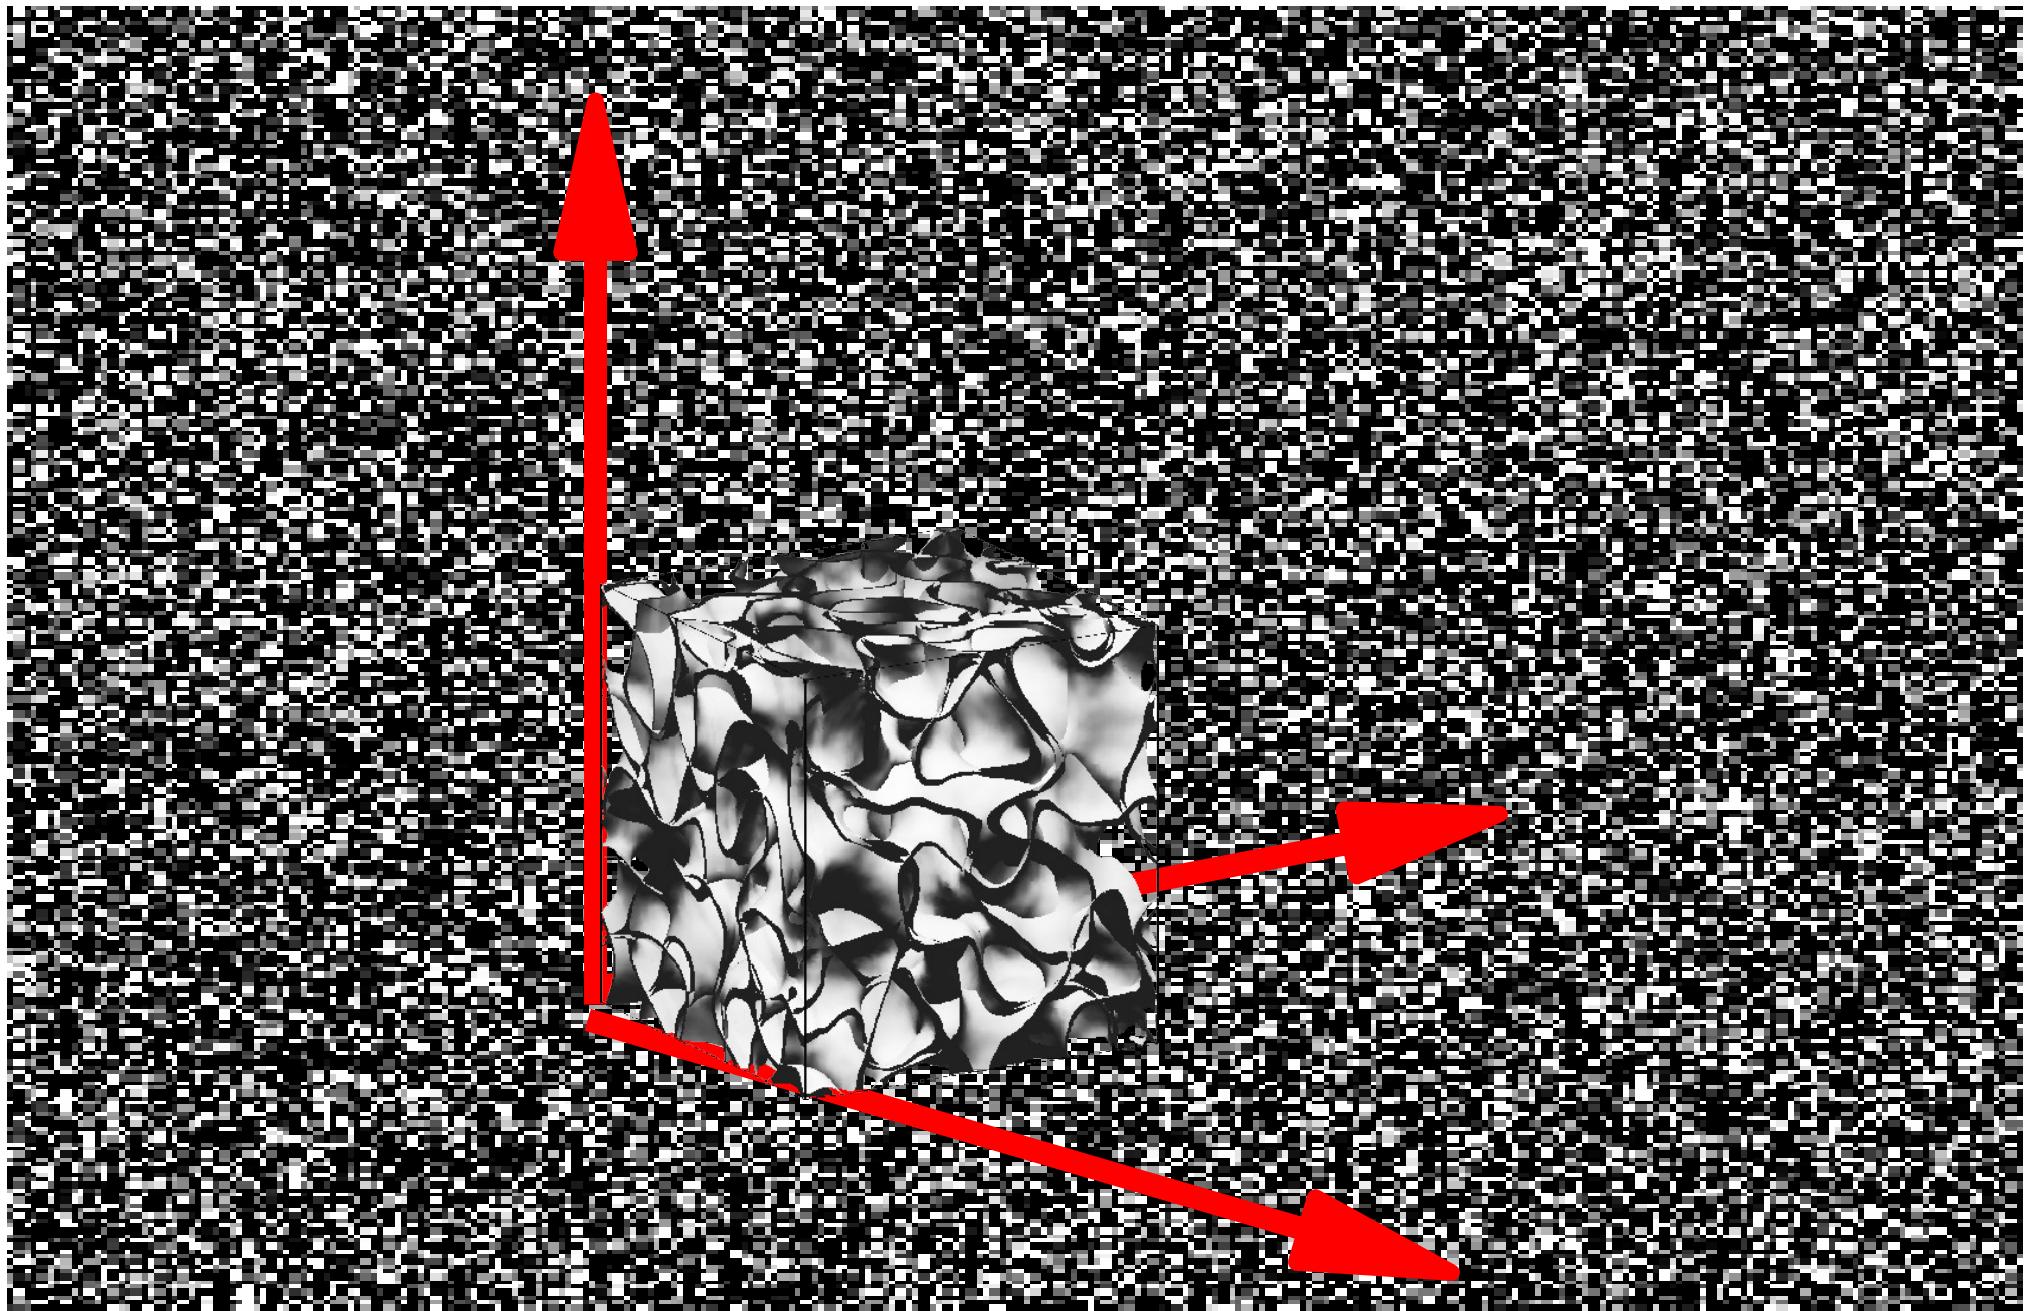


- club, **angle = 270**
- man, frontal pose
- ...



- club, **angle = 360**
- man, side pose
- ...

# The Manifold of Natural Images

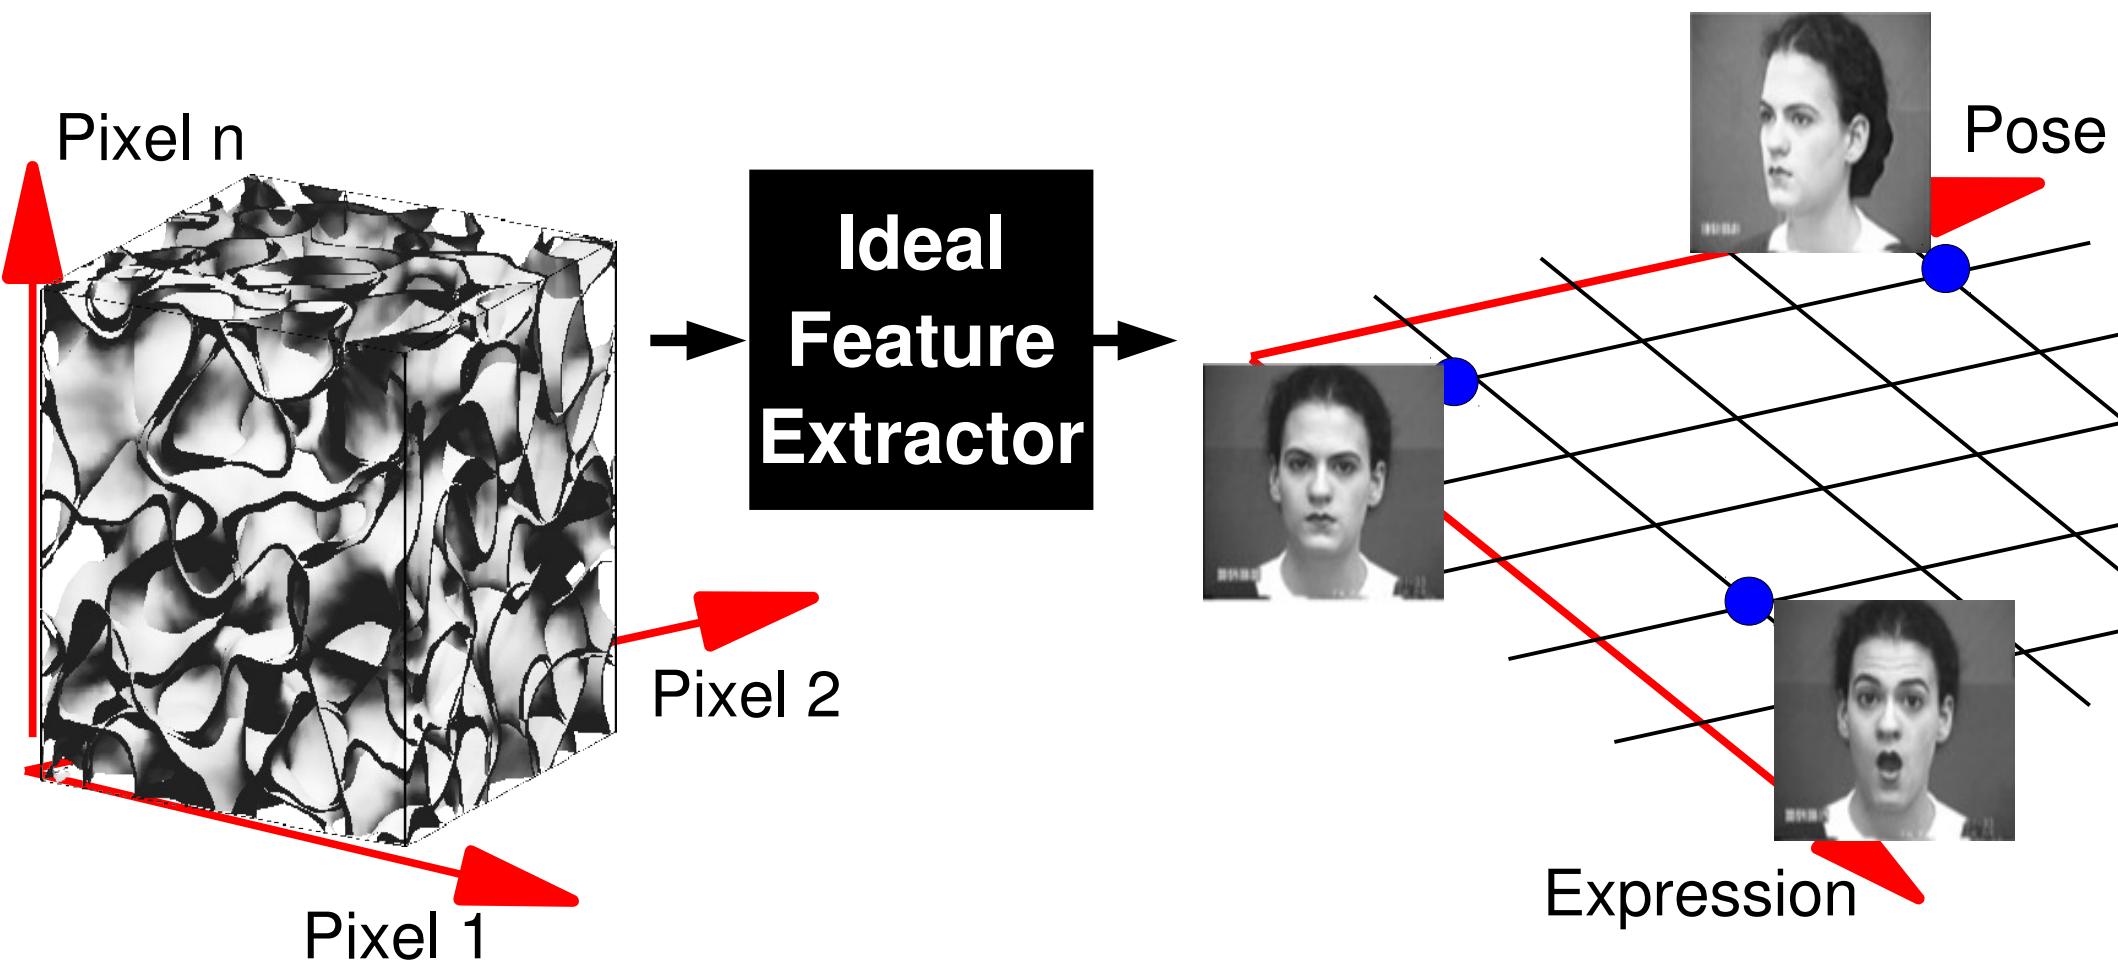


# The Manifold of Natural Images

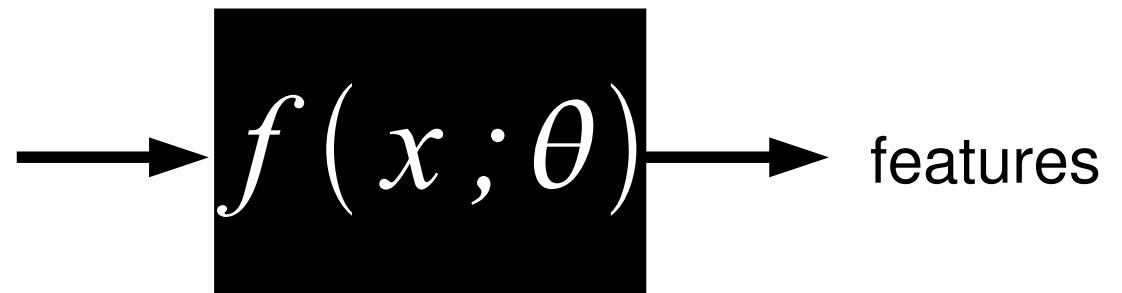
We need to linearize the manifold: learn non-linear features!



# Ideal Feature Extraction



# Learning Non-Linear Features



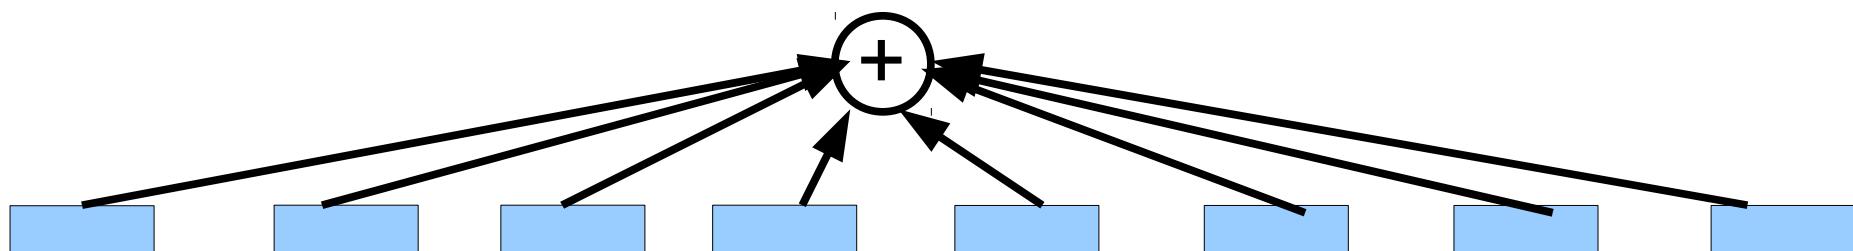
**Q.:** which class of non-linear functions shall we consider?

# Learning Non-Linear Features

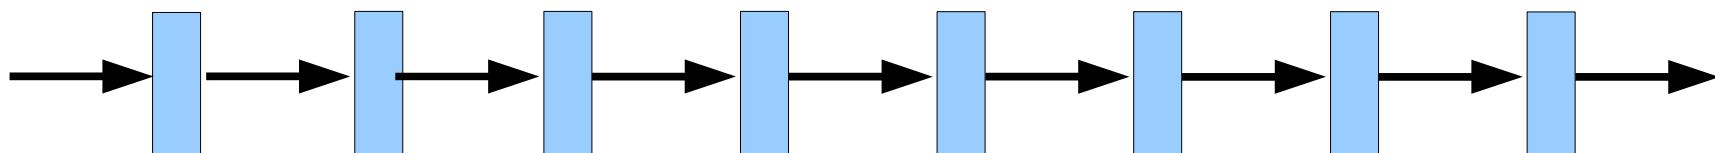
Given a dictionary of simple non-linear functions:  $g_1, \dots, g_n$

**Proposal #1: linear combination**

$$f(x) \approx \sum_j g_j$$



**Proposal #2: composition**  $f(x) \approx g_1(g_2(\dots g_n(x)\dots))$



# Learning Non-Linear Features

Given a dictionary of simple non-linear functions:  $g_1, \dots, g_n$

## Proposal #1: linear combination

$$f(x) \approx \sum_j g_j$$

- Kernel learning
- Boosting
- ...



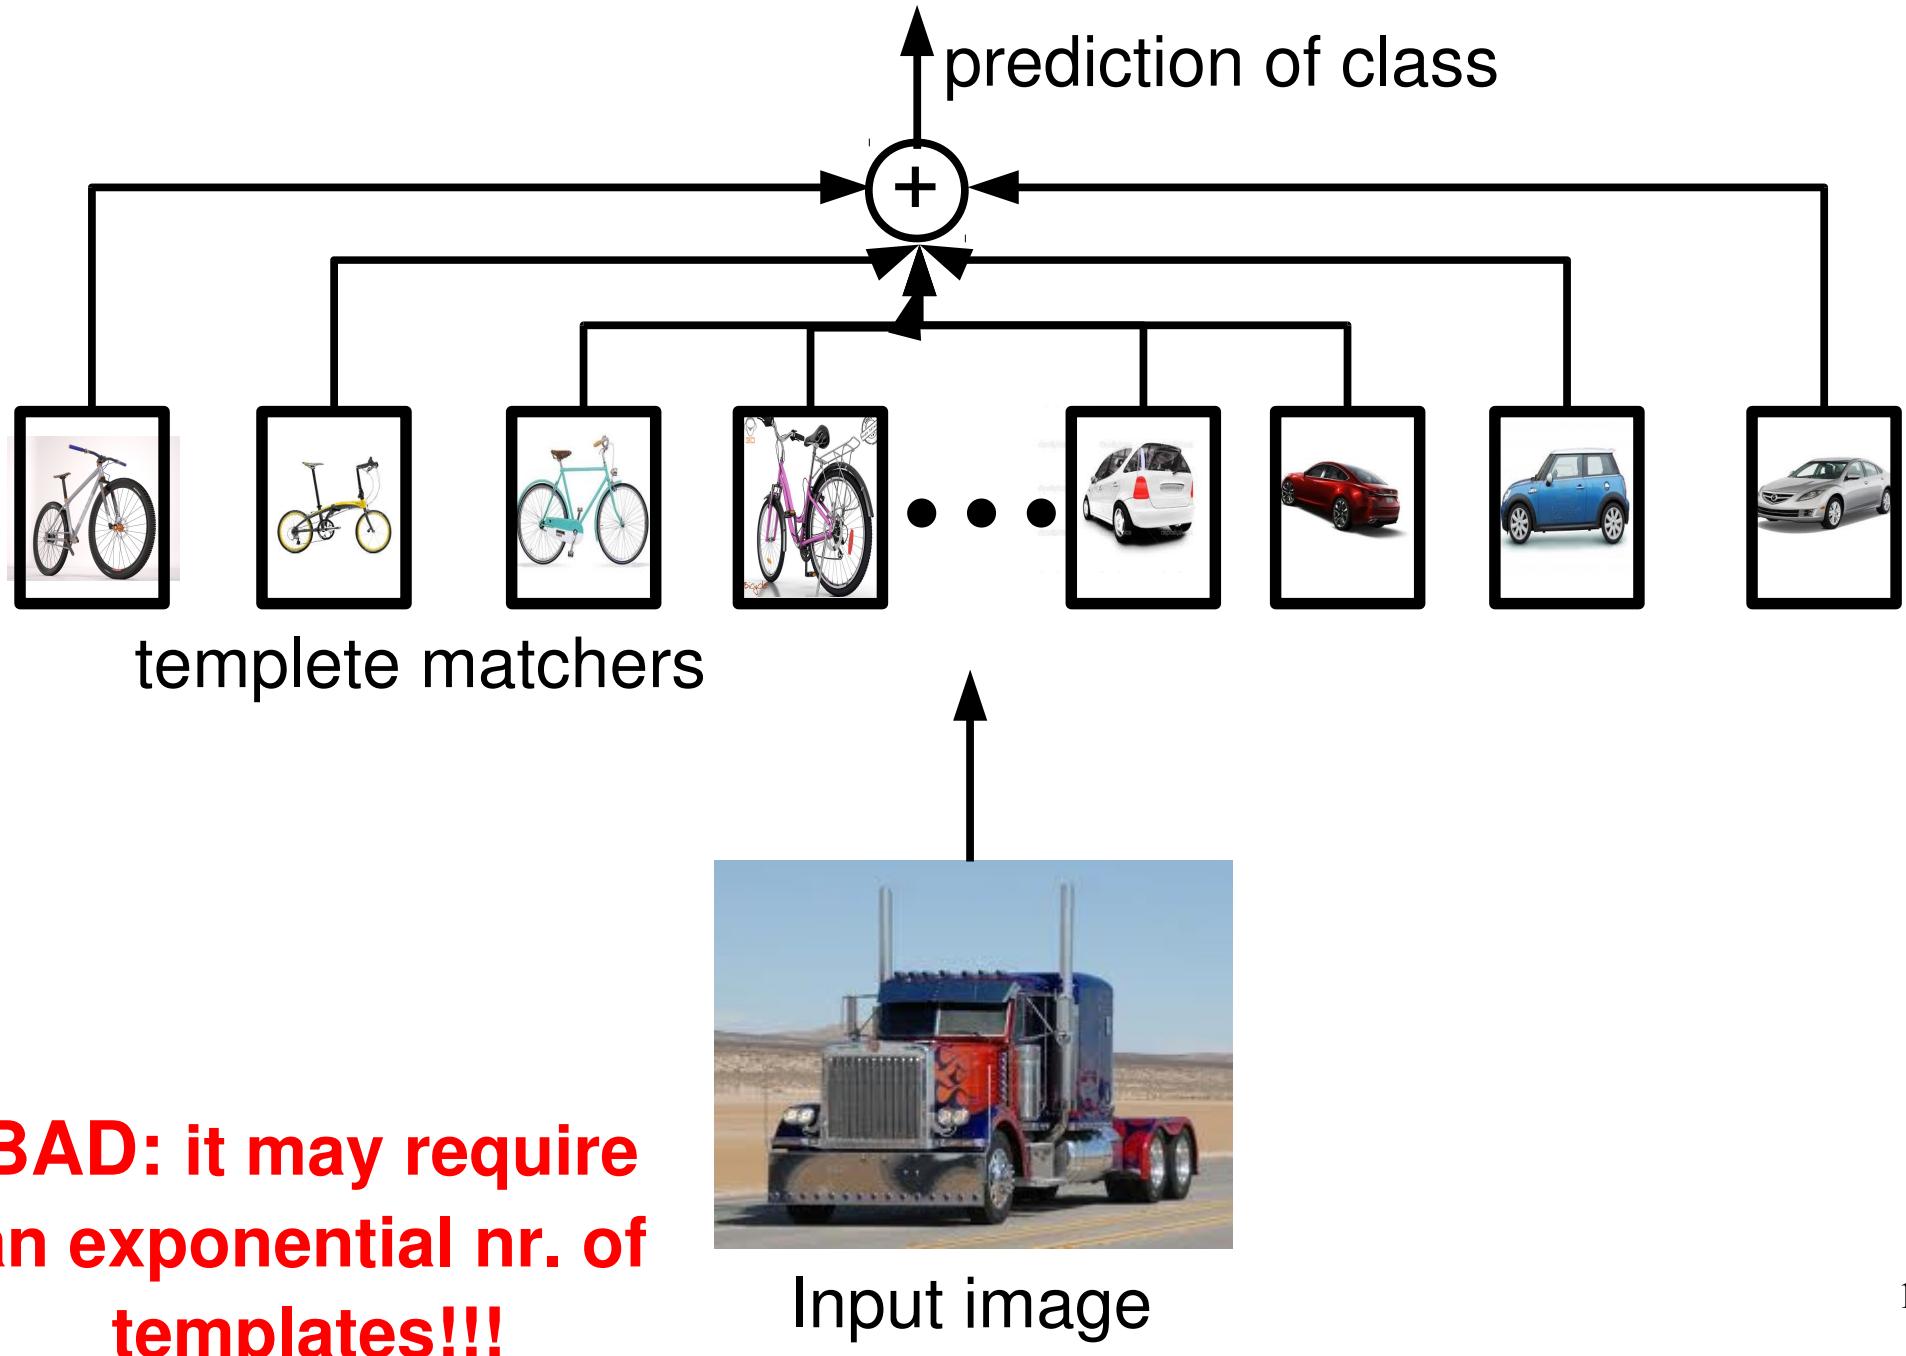
## Proposal #2: composition

$$f(x) \approx g_1(g_2(\dots g_n(x)\dots))$$

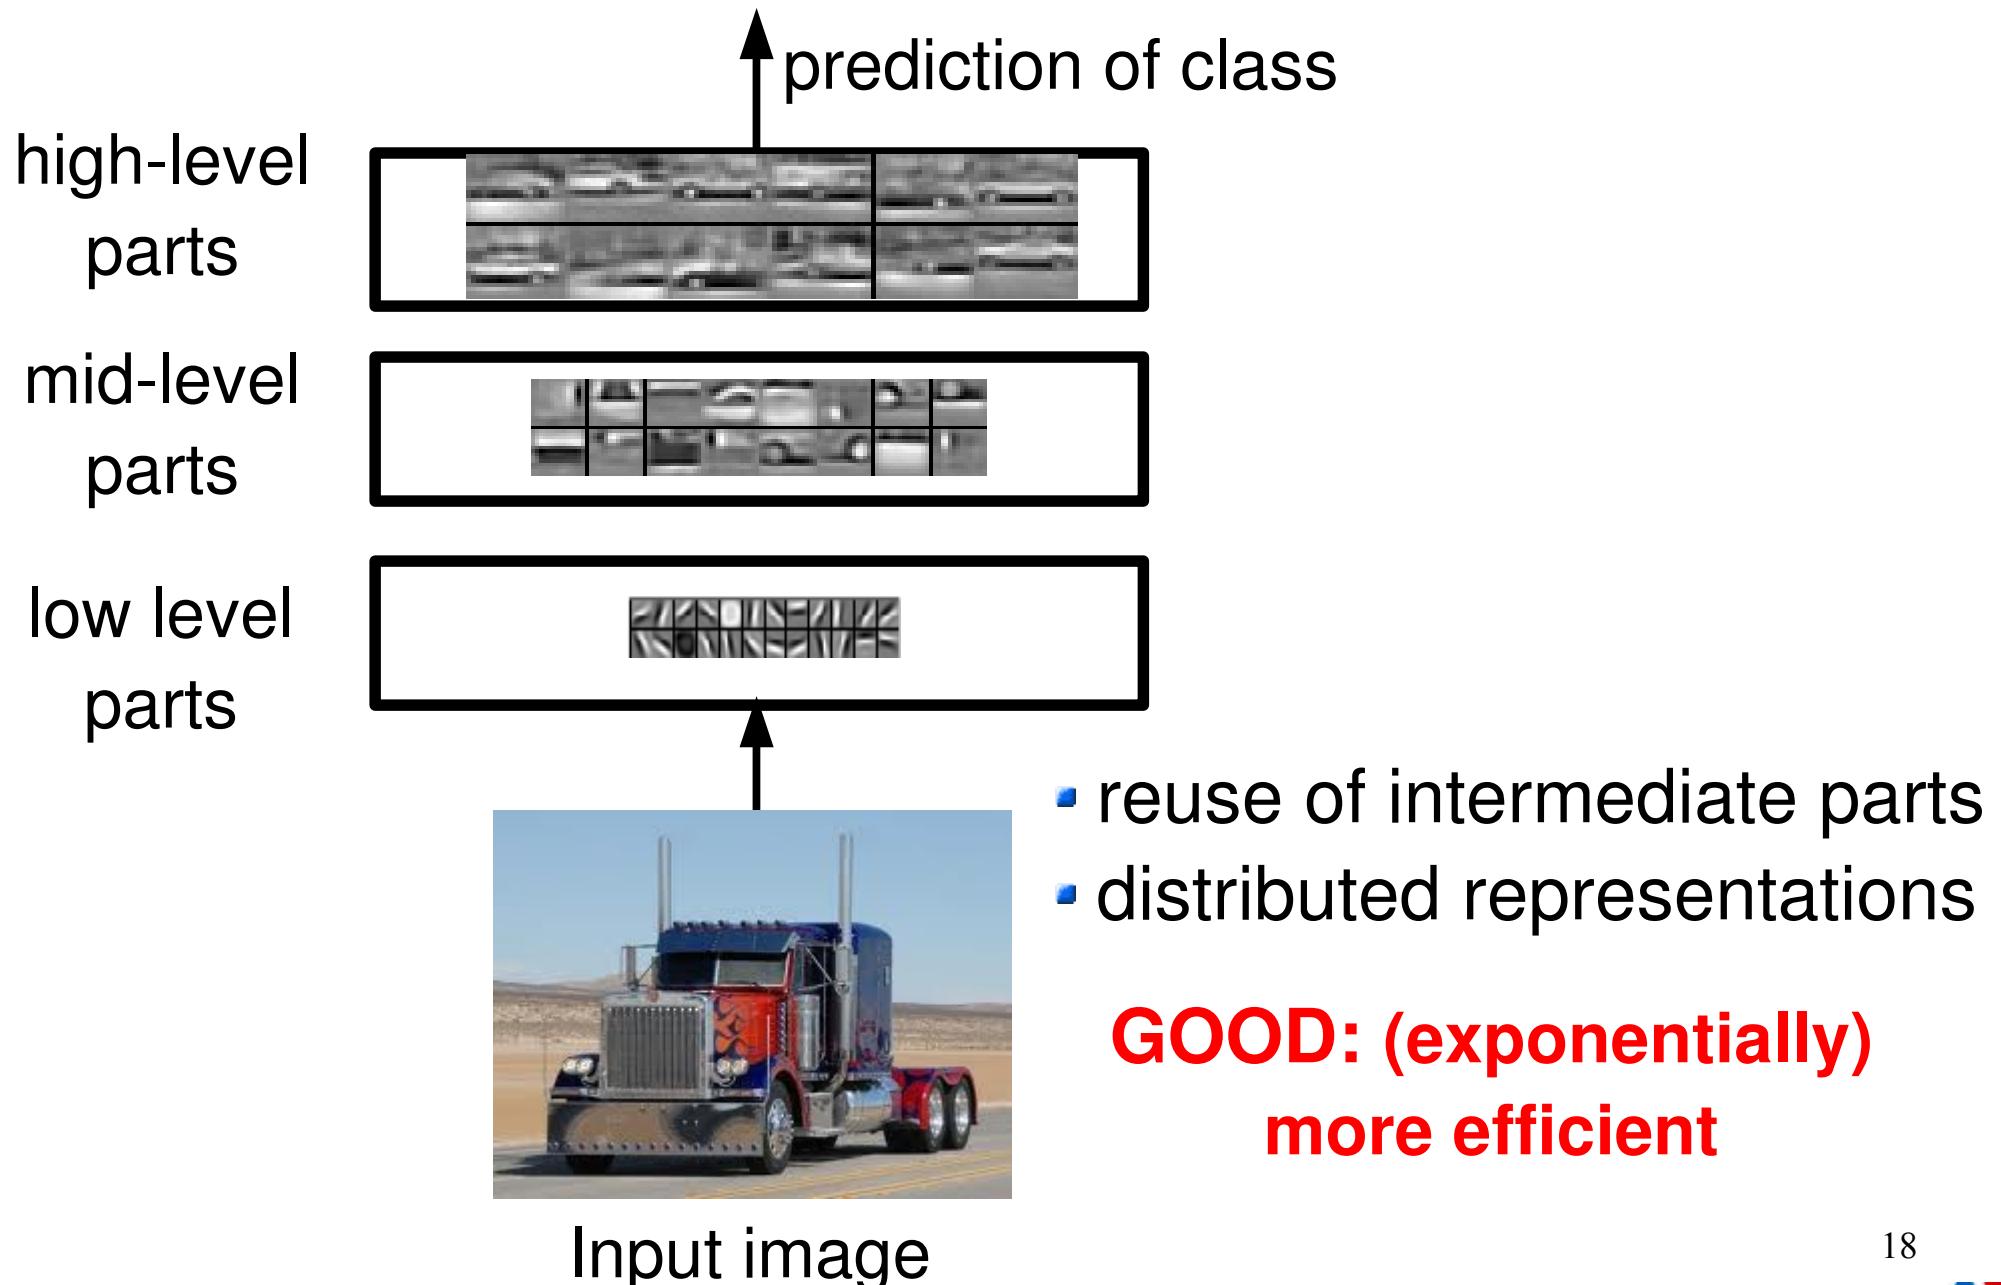
- Deep learning
- Scattering networks (wavelet cascade)
- S.C. Zhou & D. Mumford “grammar”



# Linear Combination

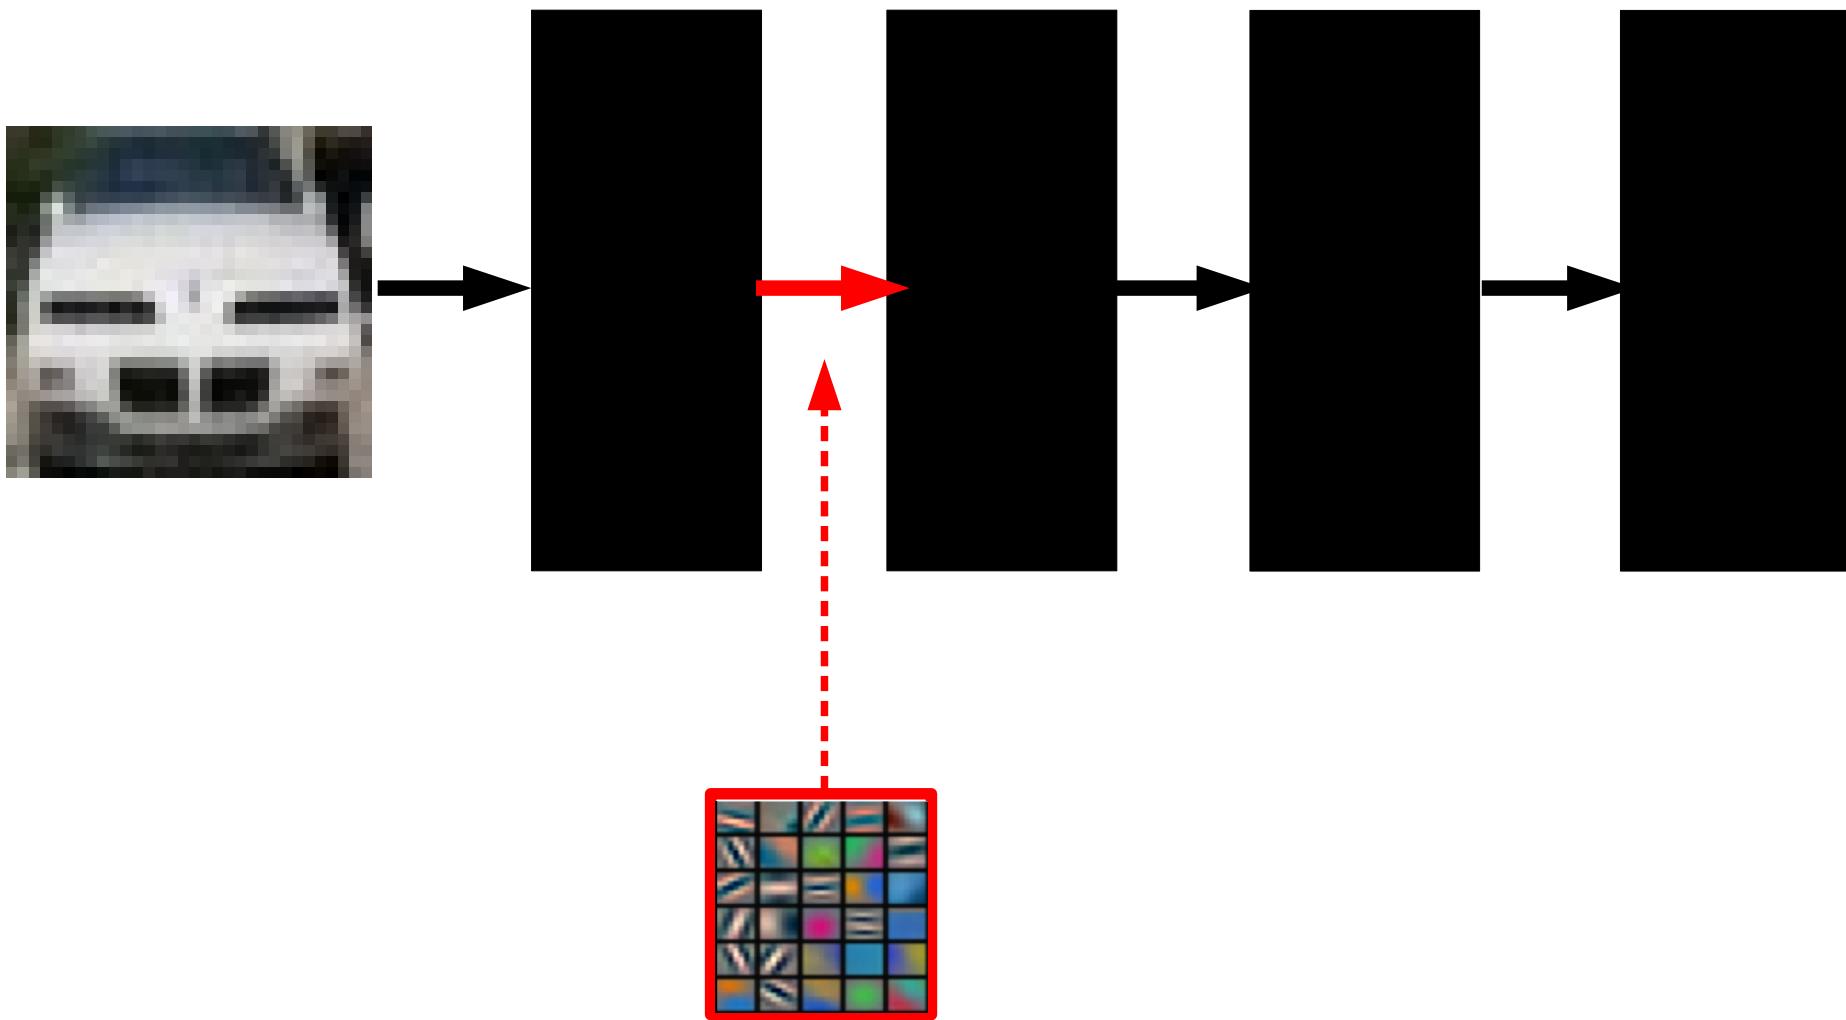


# Composition



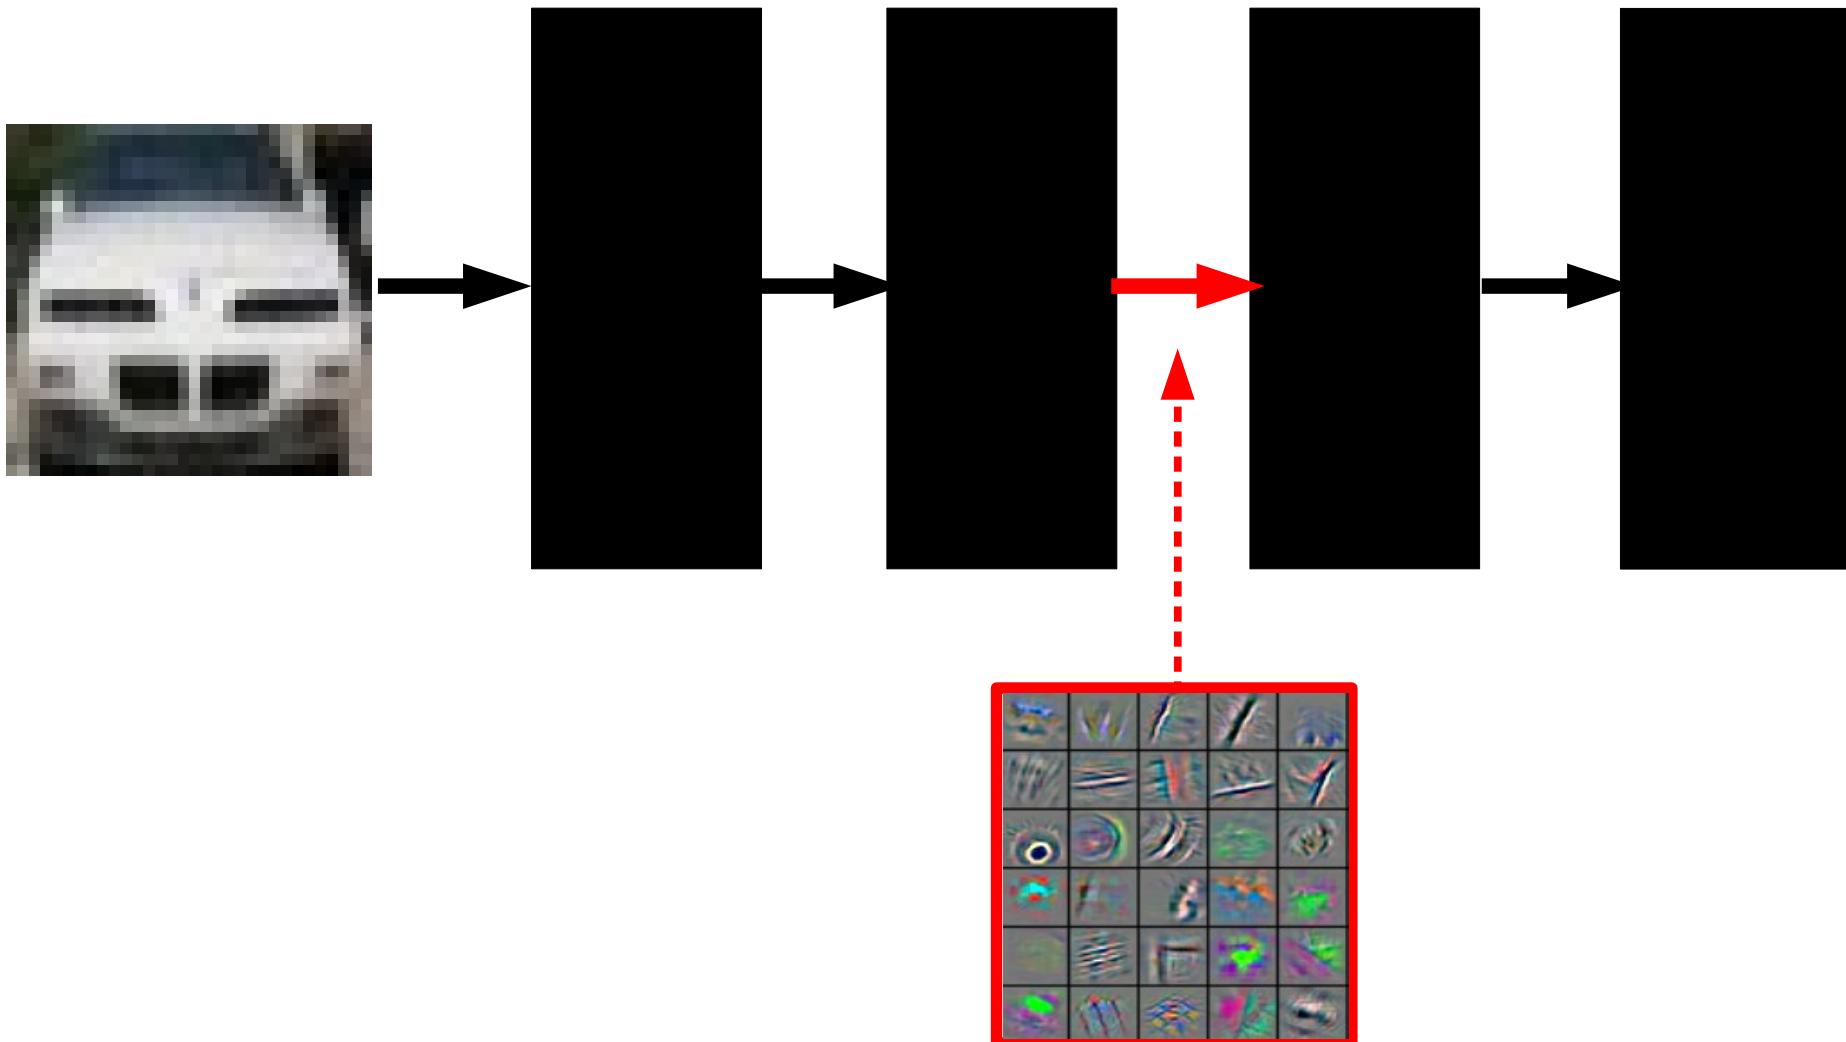
# The Big Advantage of Deep Learning

Efficiency: intermediate concepts can be re-used



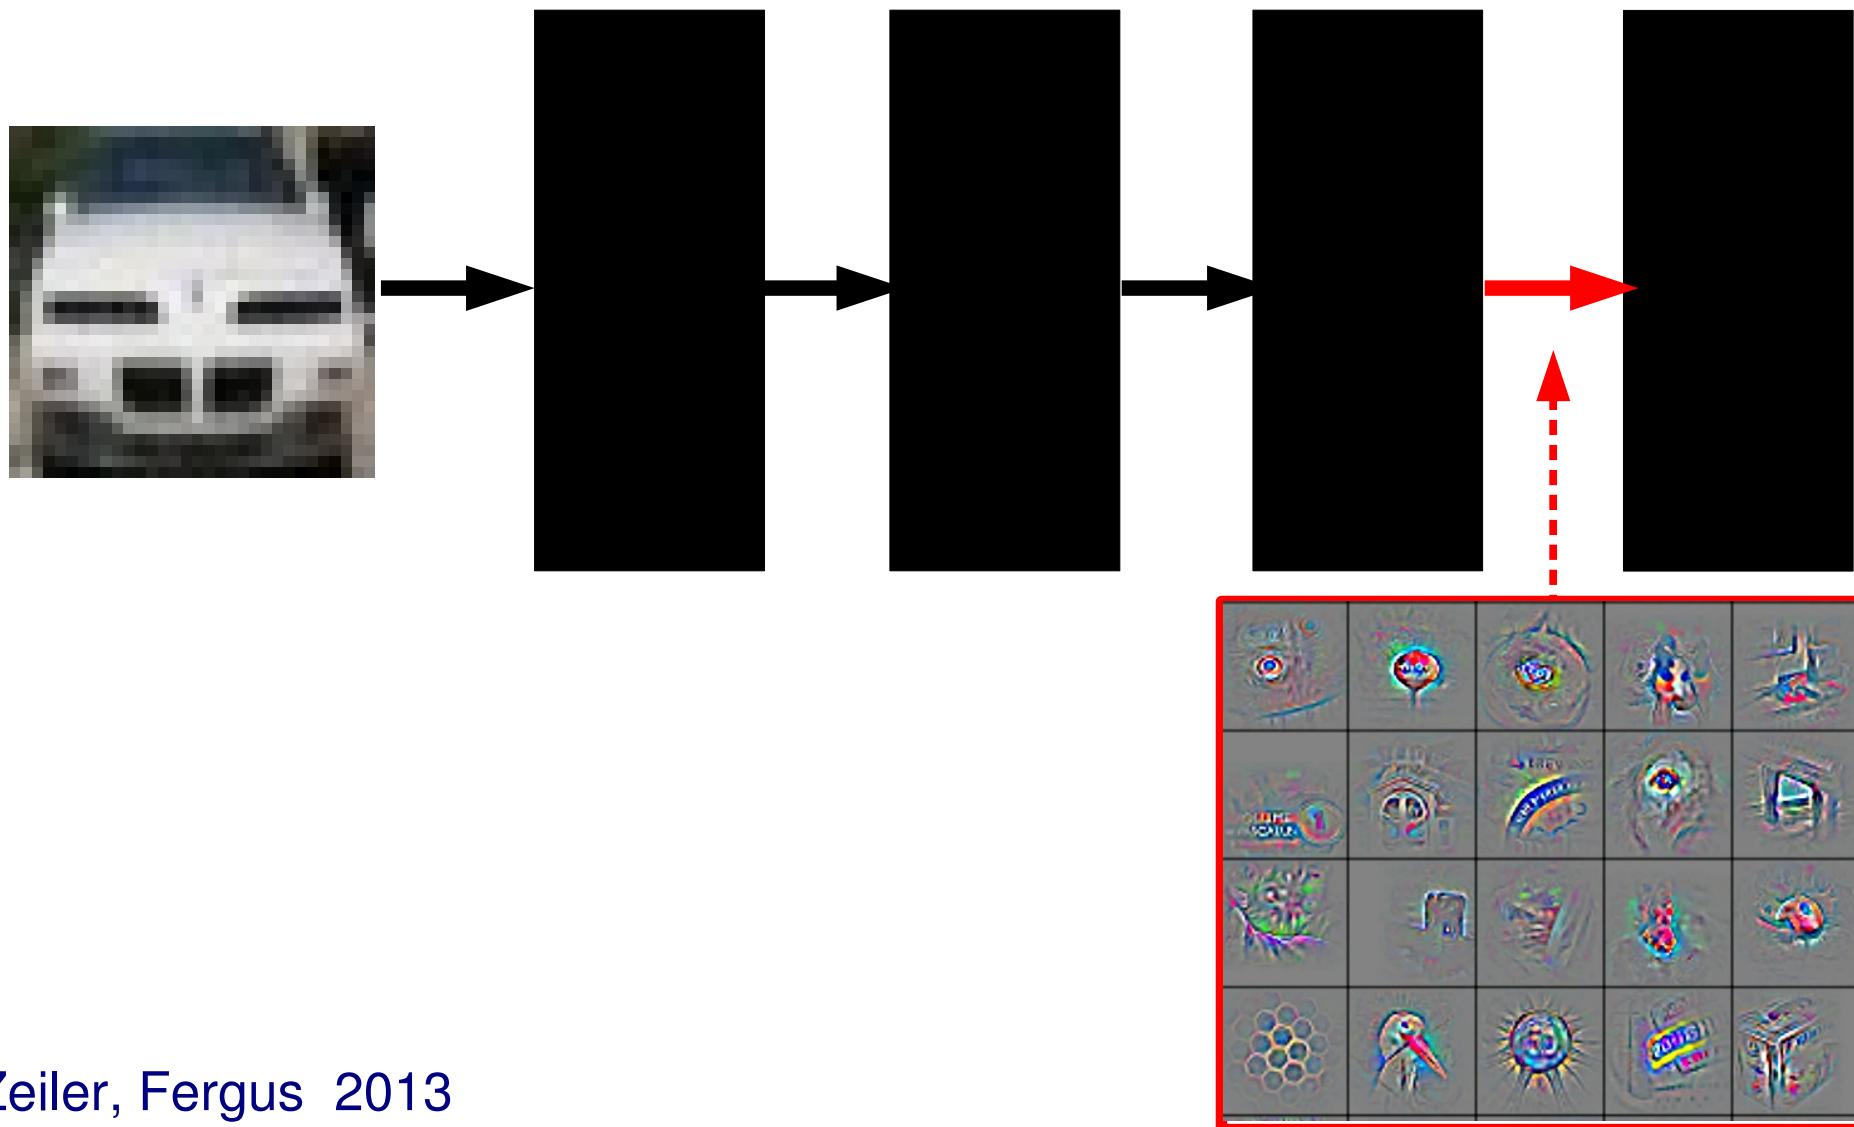
# The Big Advantage of Deep Learning

Efficiency: intermediate concepts can be re-used



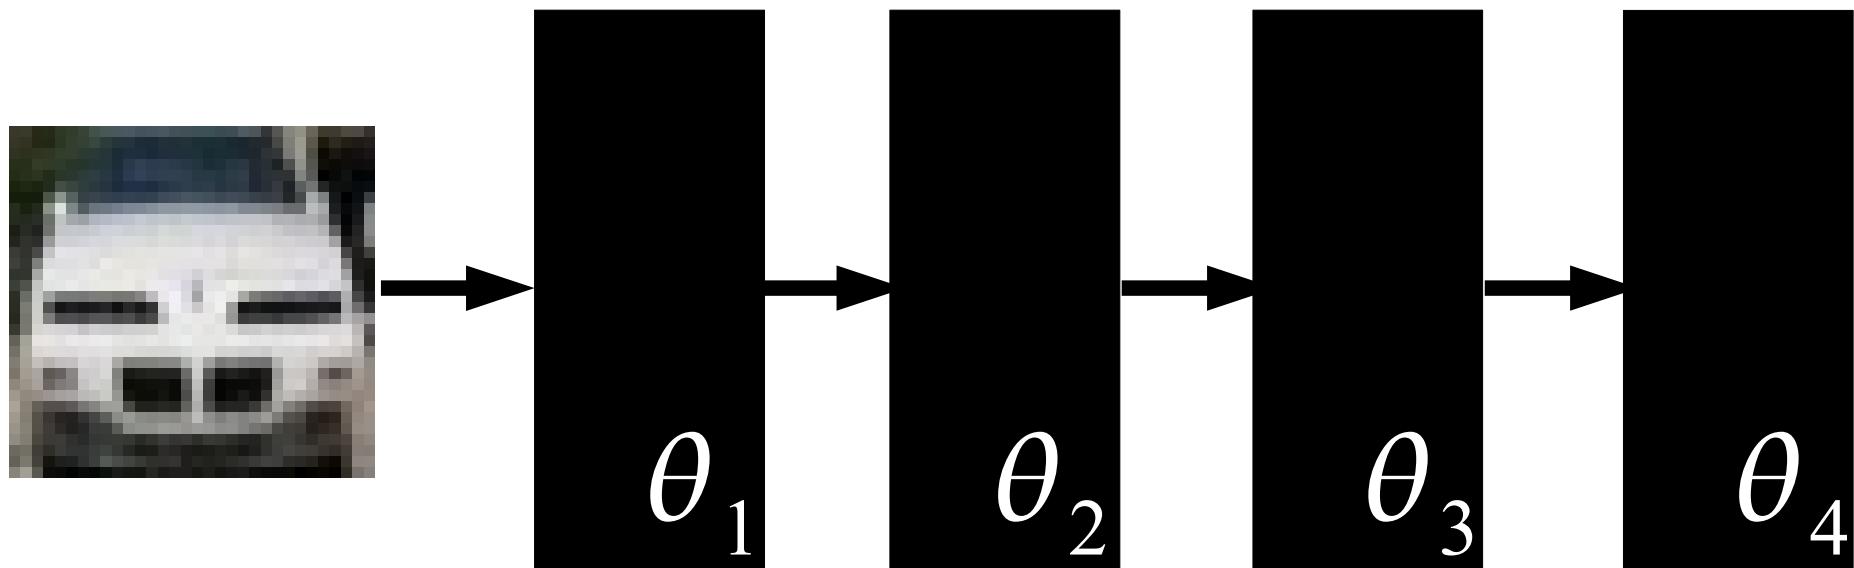
# The Big Advantage of Deep Learning

Efficiency: intermediate concepts can be re-used



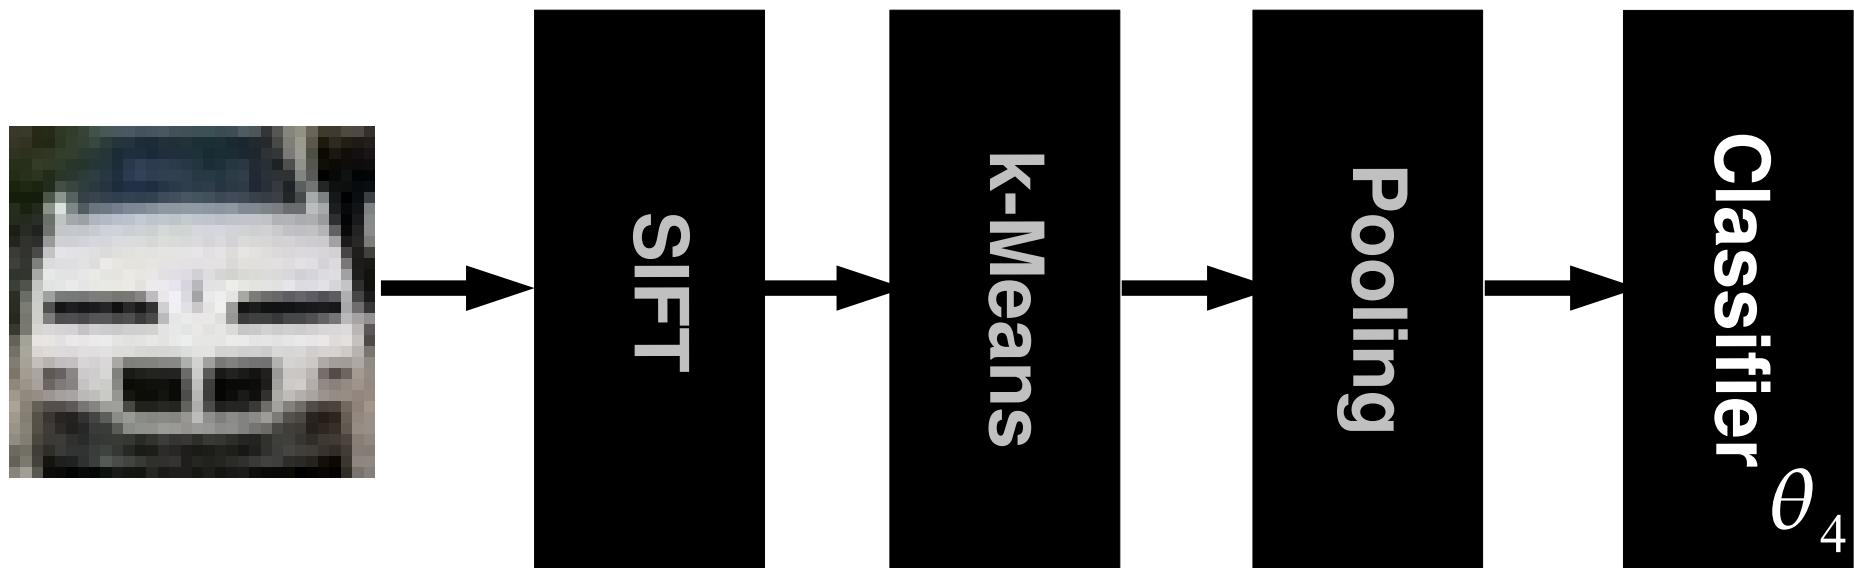
# A Potential Problem with Deep Learning

Optimization is difficult: non-convex, non-linear system



# A Potential Problem with Deep Learning

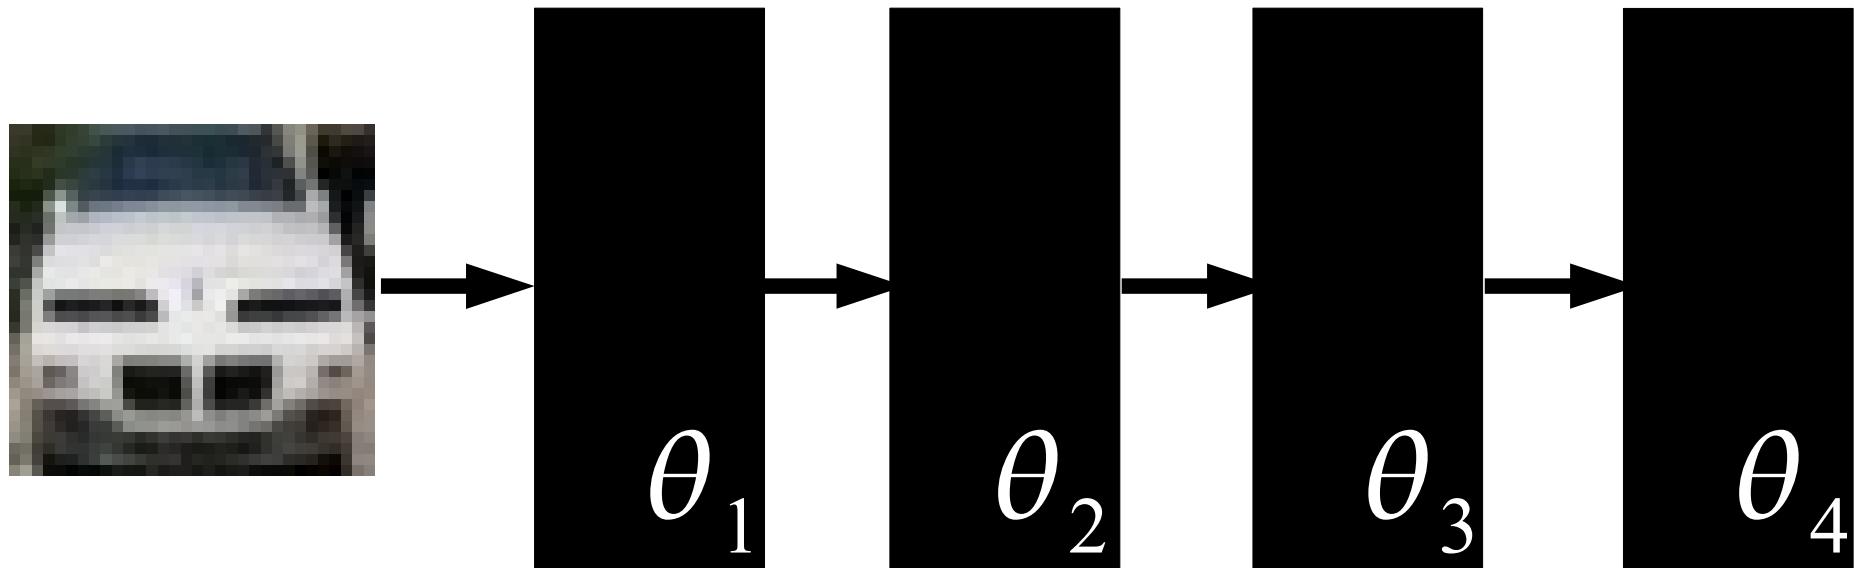
Optimization is difficult: non-convex, non-linear system



**Solution #1:** freeze first N-1 layer (engineer the features)  
It makes it **shallow!**

# A Potential Problem with Deep Learning

Optimization is difficult: non-convex, non-linear system



**Solution #2:** live with it!

It will converge to a local minimum.

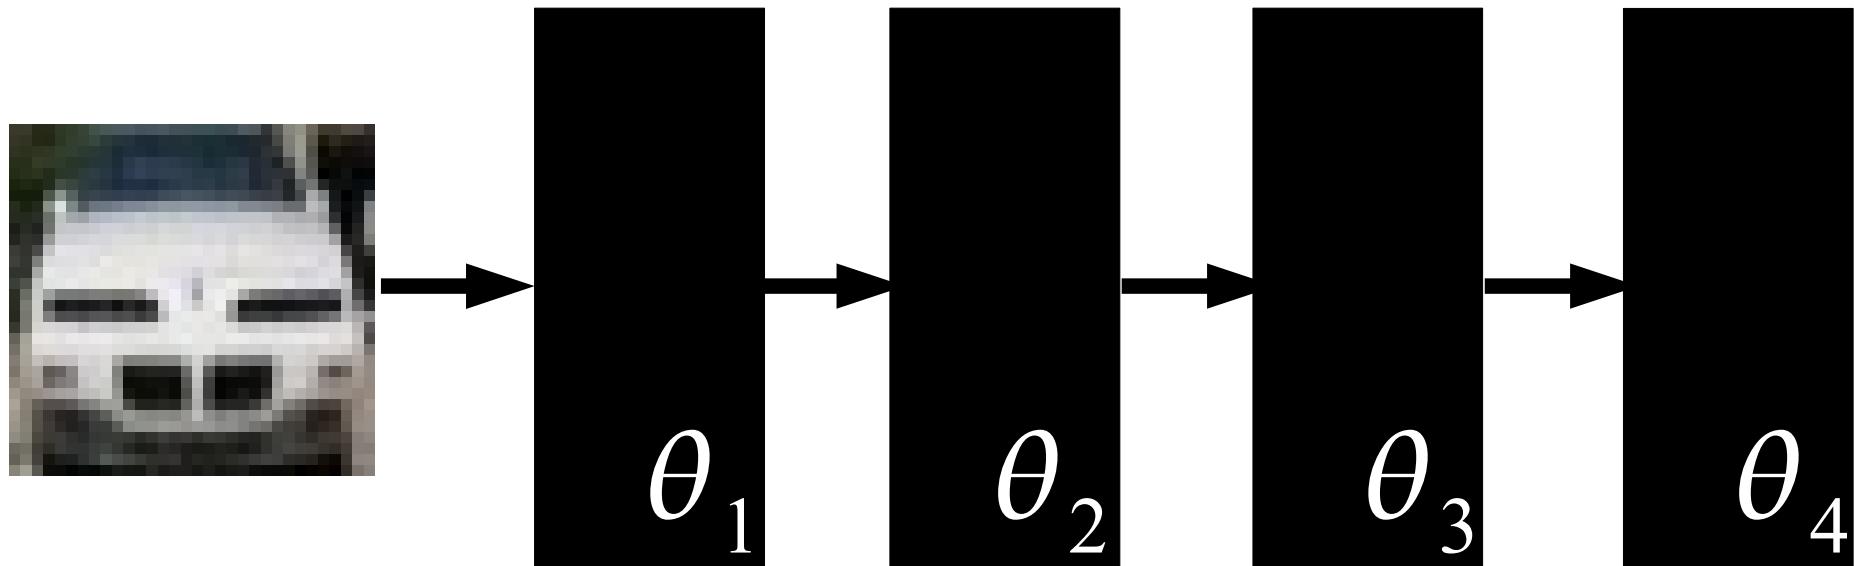
It is much more powerful!!

*Given lots of data, engineer less and learn more!!*

24

# Deep Learning in Practice

Optimization is easy, need to know a few tricks of the trade.

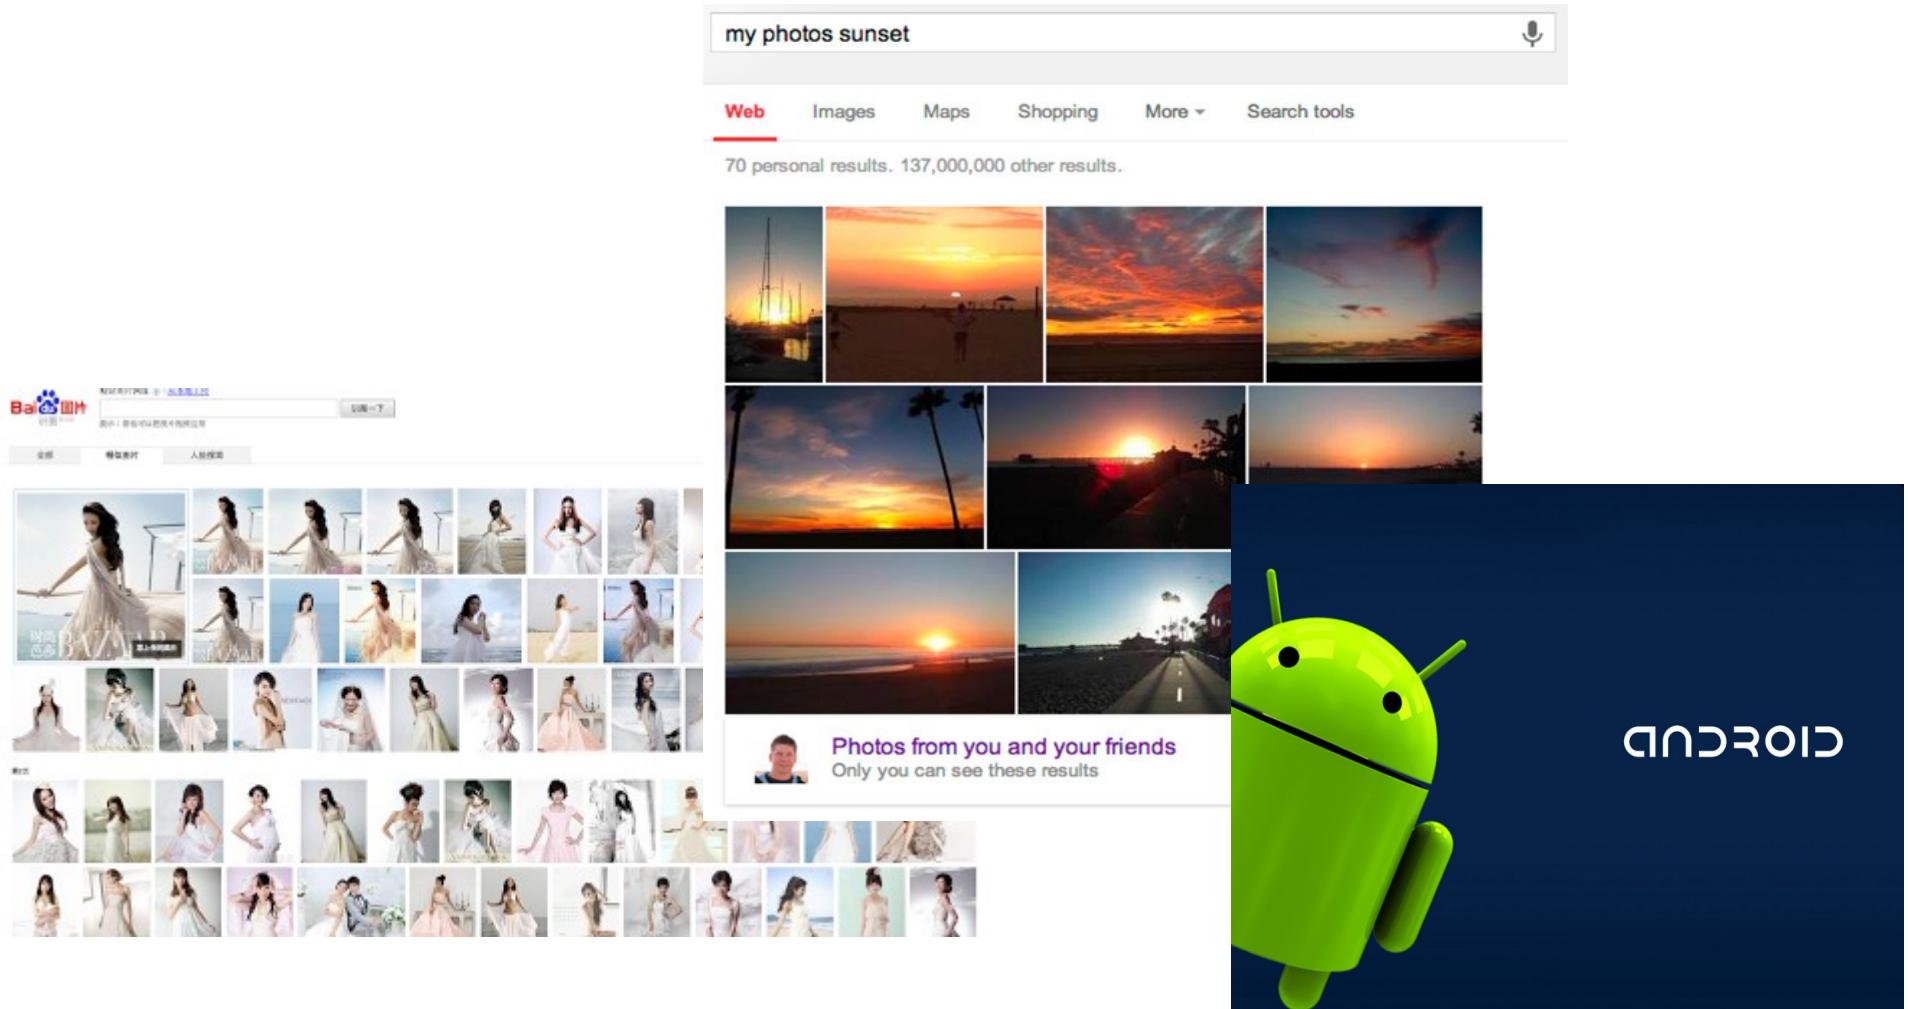


**Q:** What's the feature extractor? And what's the classifier?

**A:** No distinction, end-to-end learning!

# Deep Learning in Practice

It works very well in practice:



# KEY IDEAS: WHY DEEP LEARNING

- We need non-linear system
- We need to learn it from data
- Build feature hierarchies (function composition)
- End-to-end learning

# Outline

- Motivation
- Deep Learning: The Big Picture
- From neural nets to convolutional nets
- Applications
- A practical guide

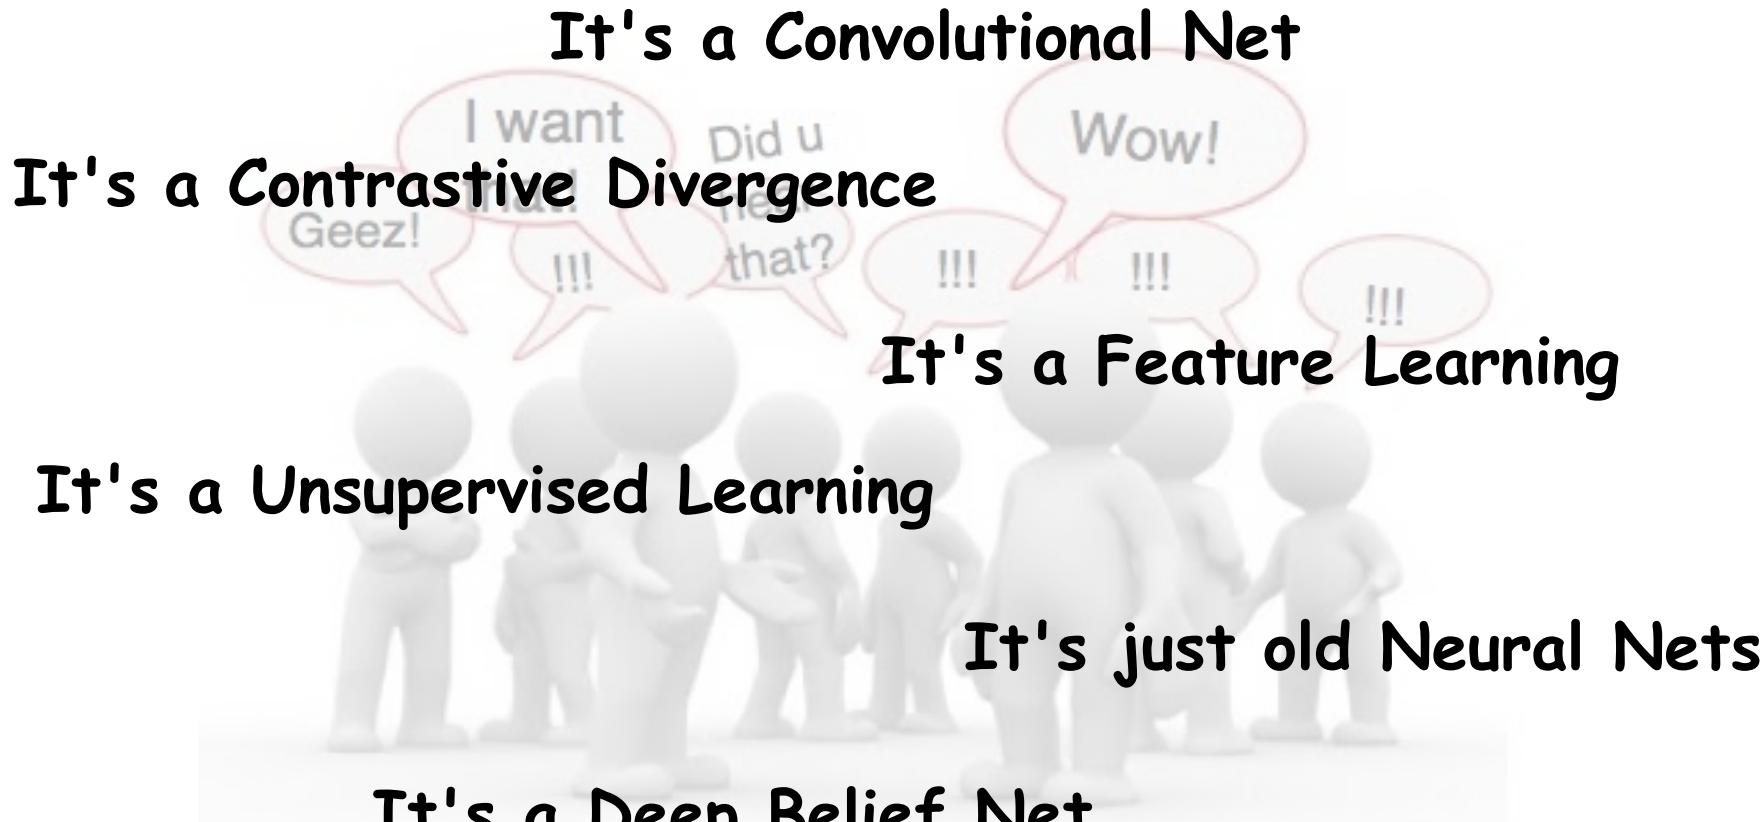
# Outline

- Motivation
- Deep Learning: The Big Picture
- From neural nets to convolutional nets
- Applications
- A practical guide

# What Is Deep Learning?



# Buzz Words



# (My) Definition

**A Deep Learning method is:** a method which makes predictions by using a sequence of non-linear processing stages. The resulting intermediate representations can be interpreted as feature hierarchies and the whole system is jointly learned from data.

Some deep learning methods are probabilistic, others are loss-based, some are supervised, other unsupervised...

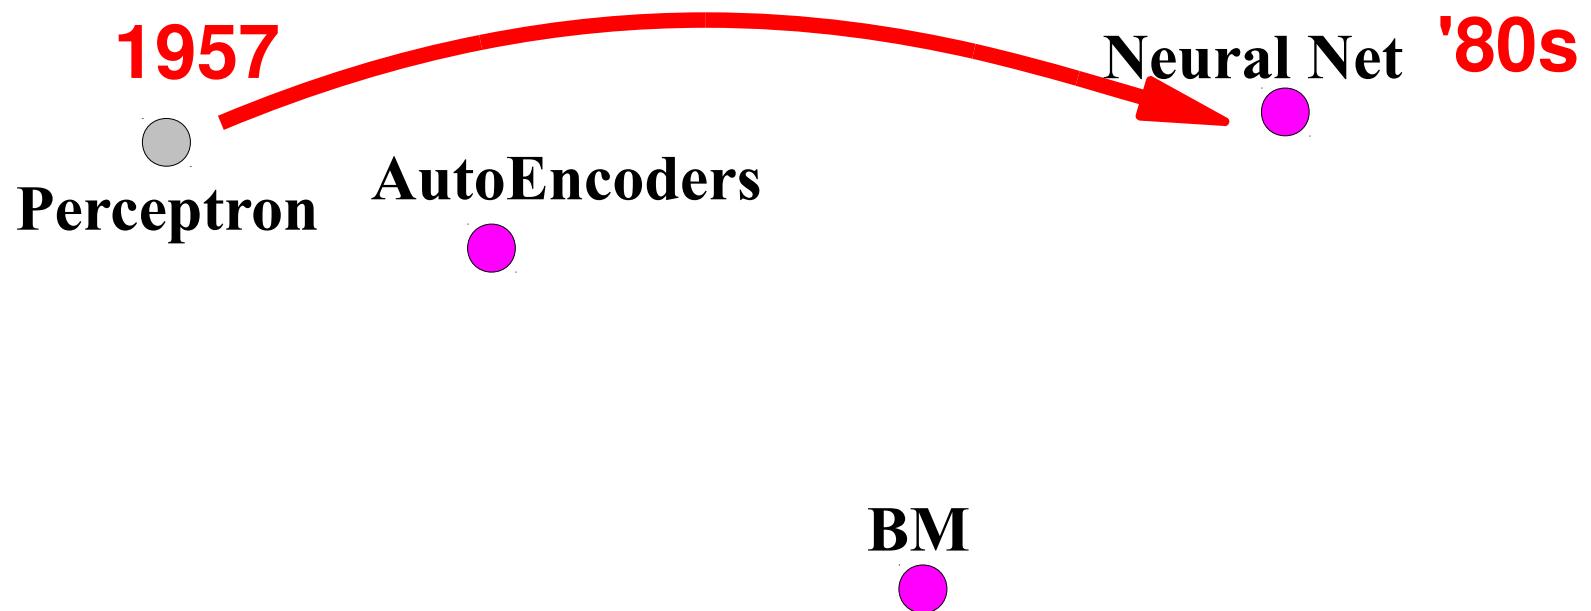
It's a large family!

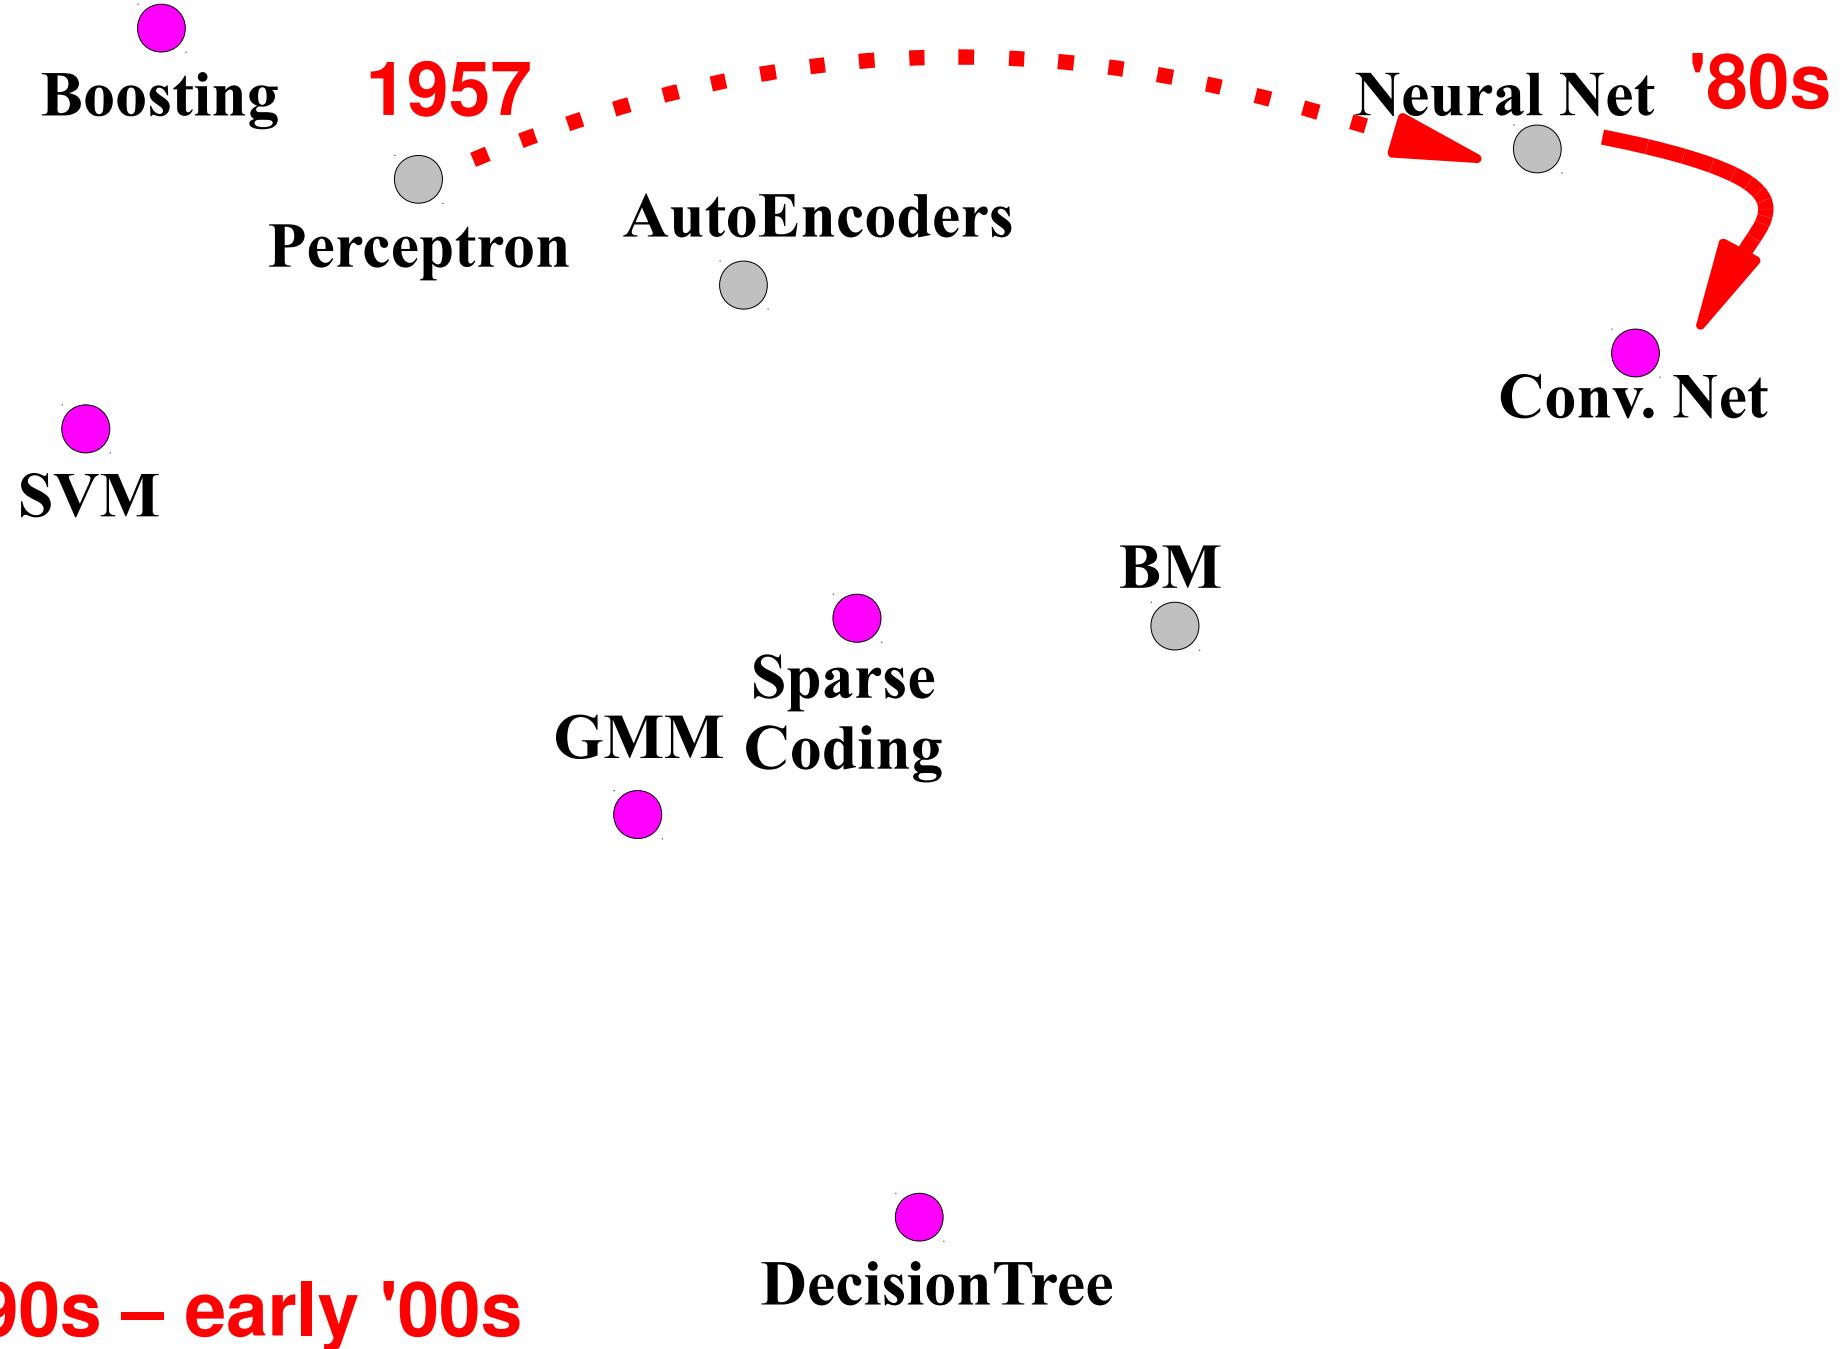
**1957**

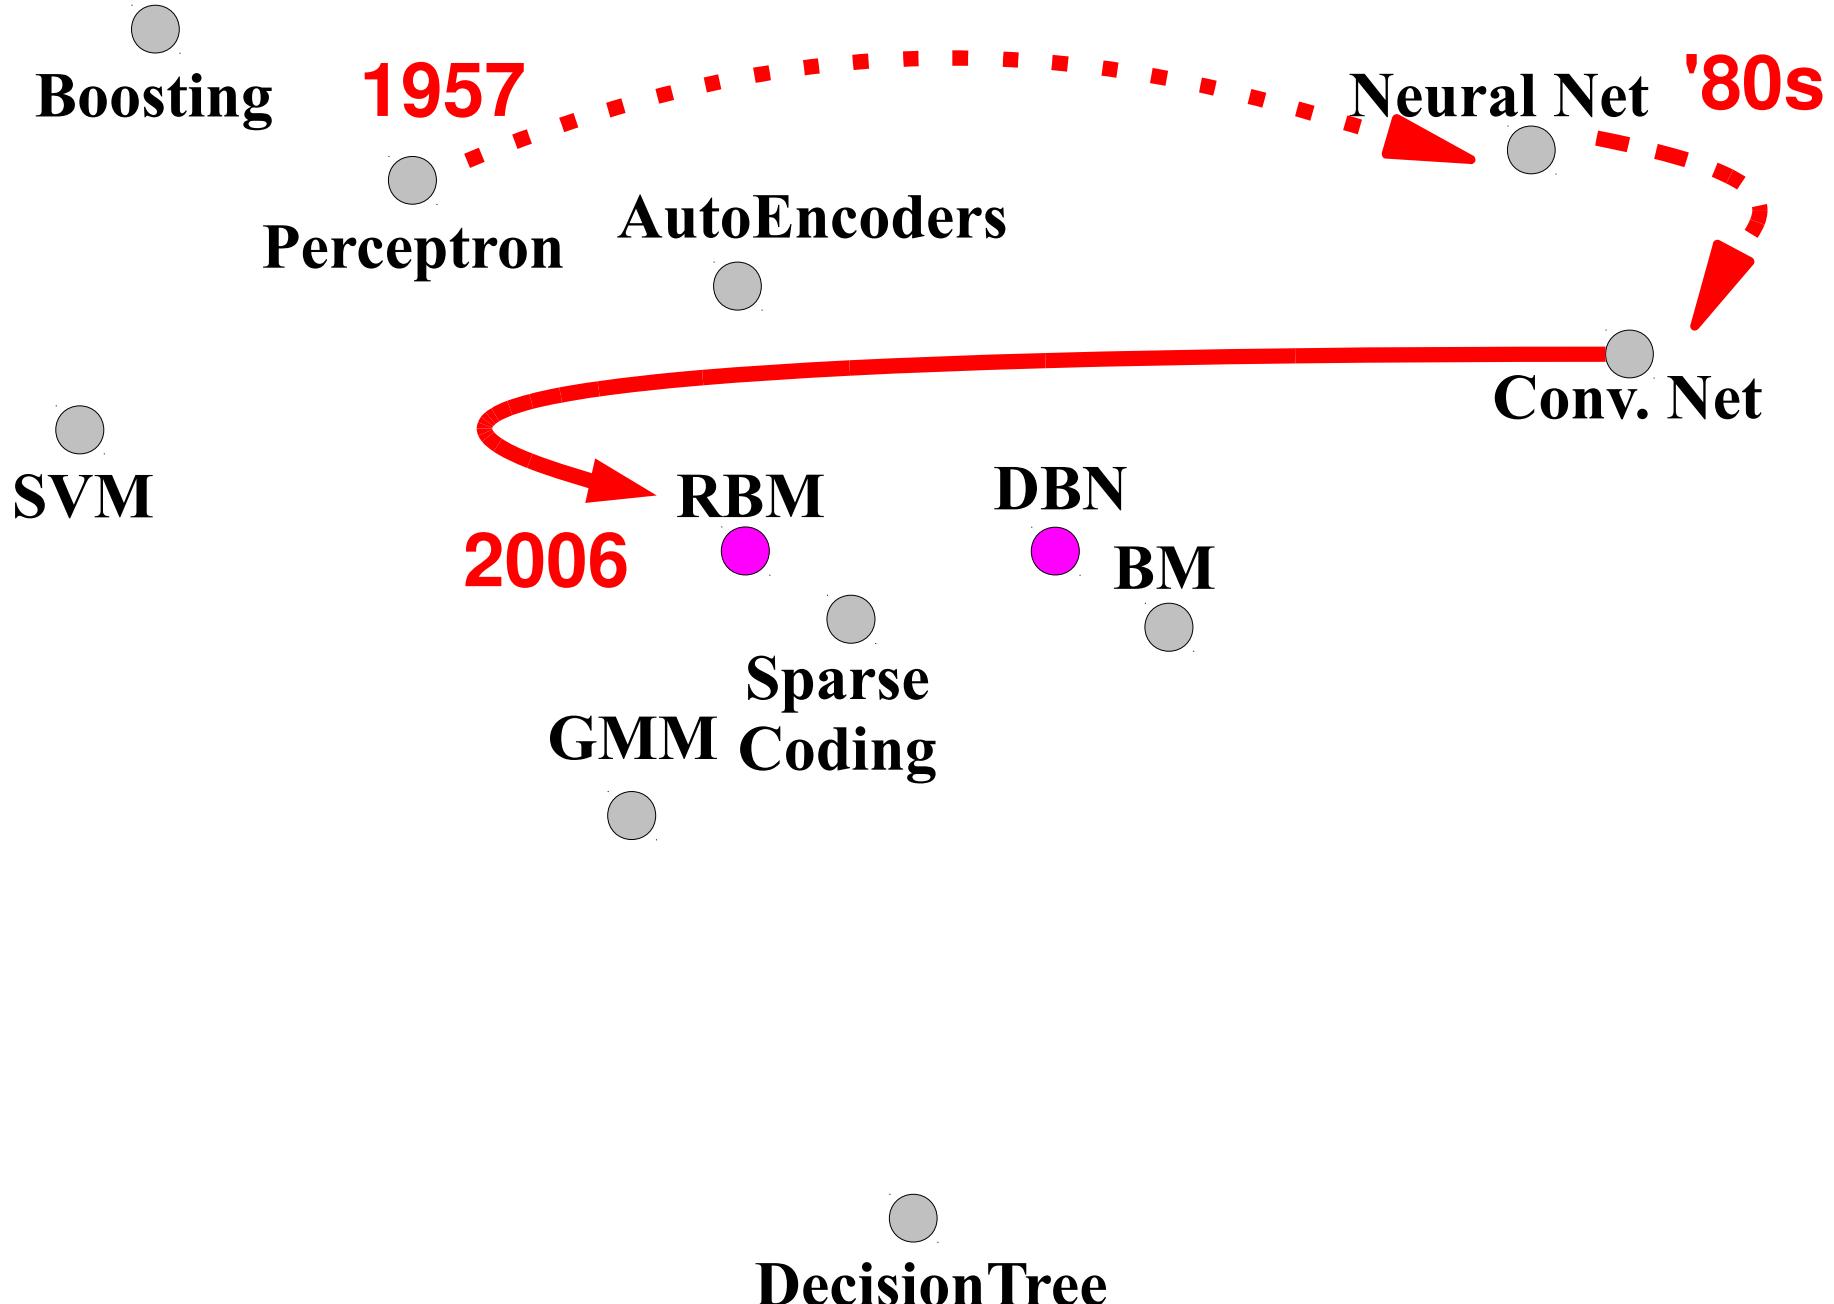


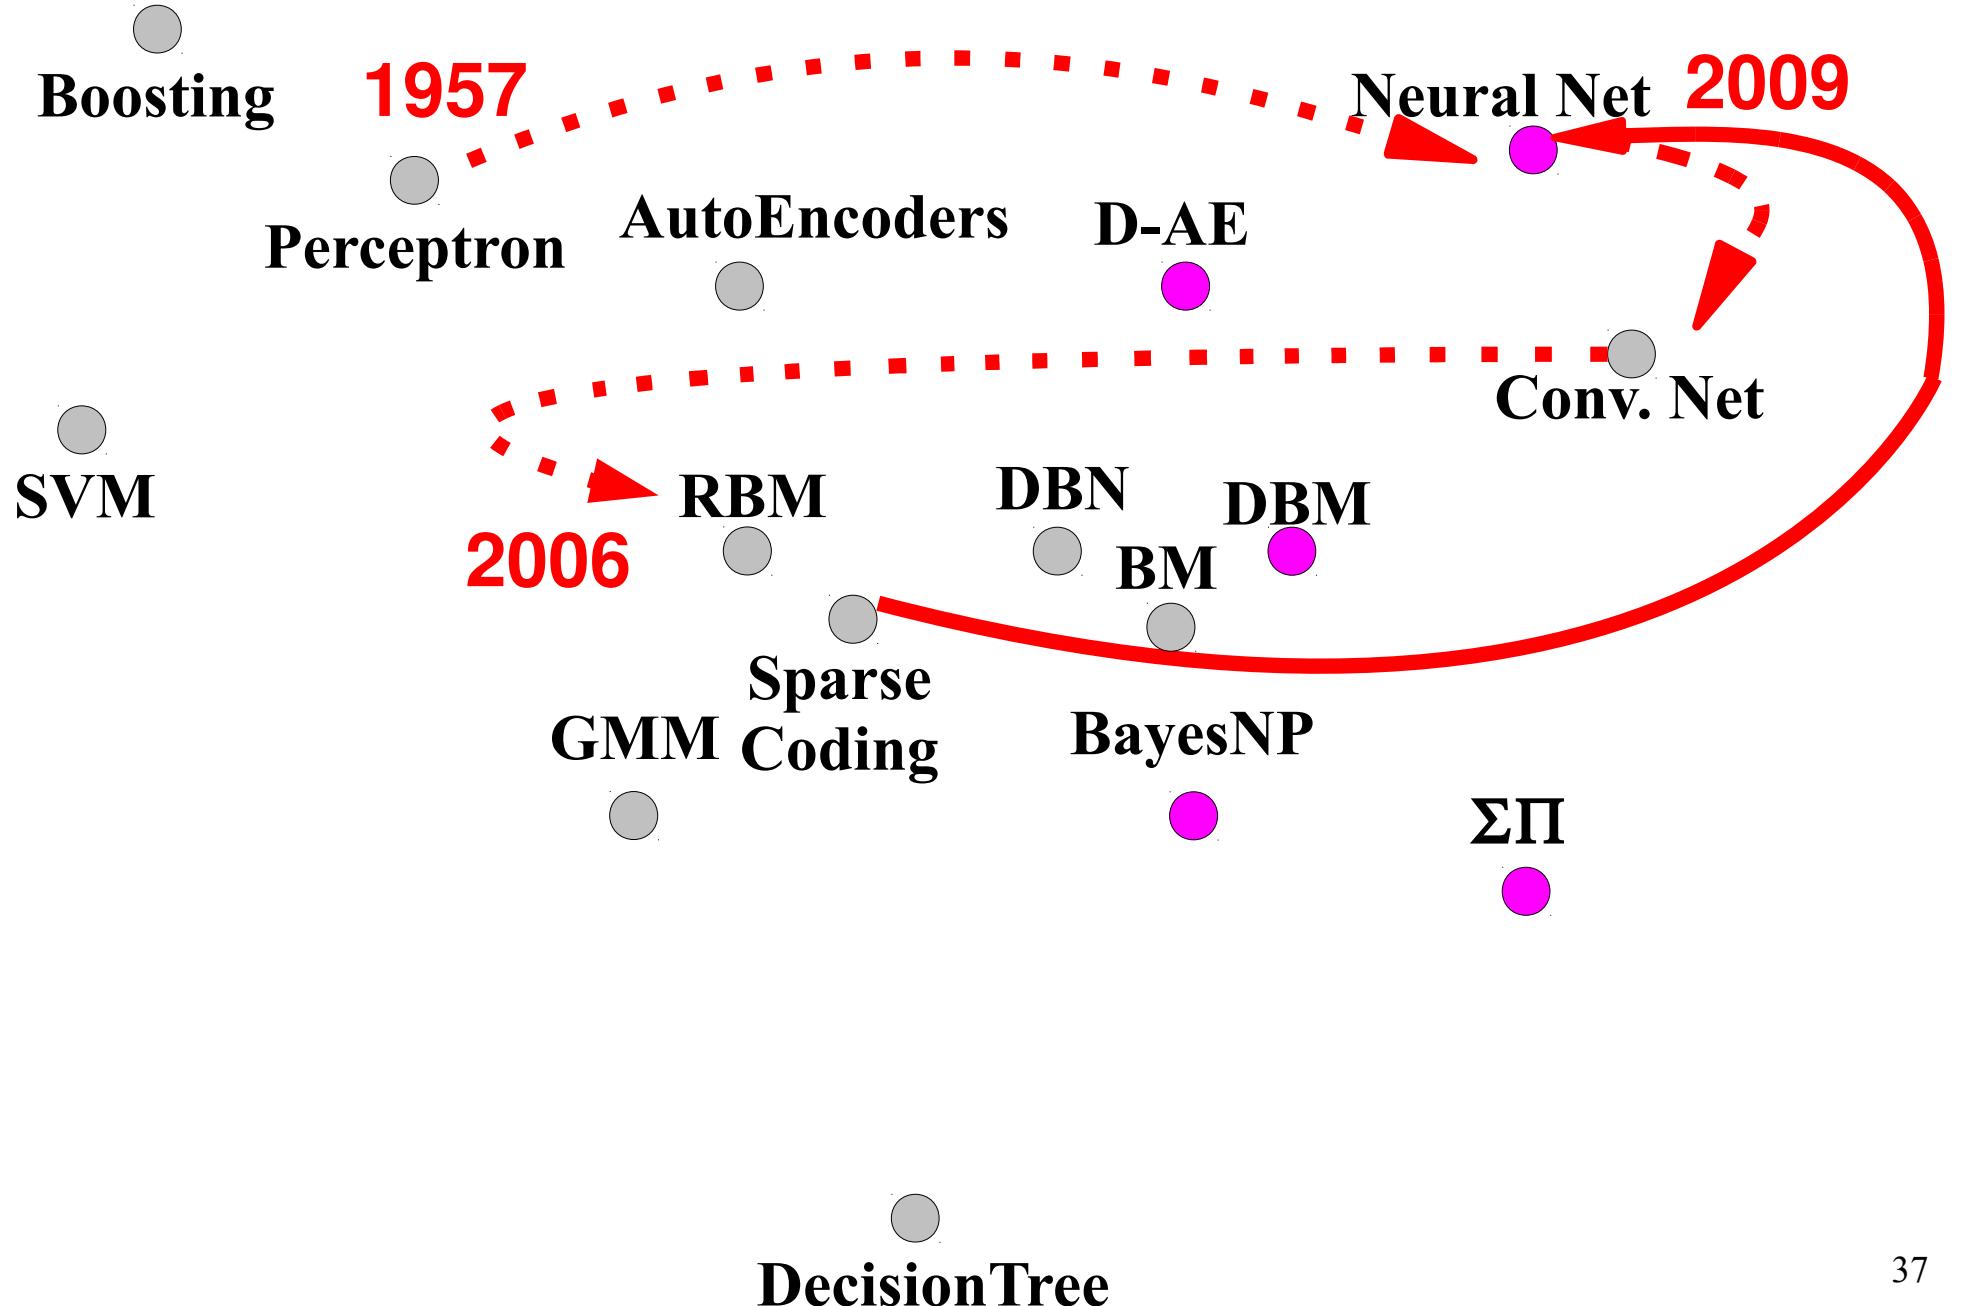
**Perceptron**

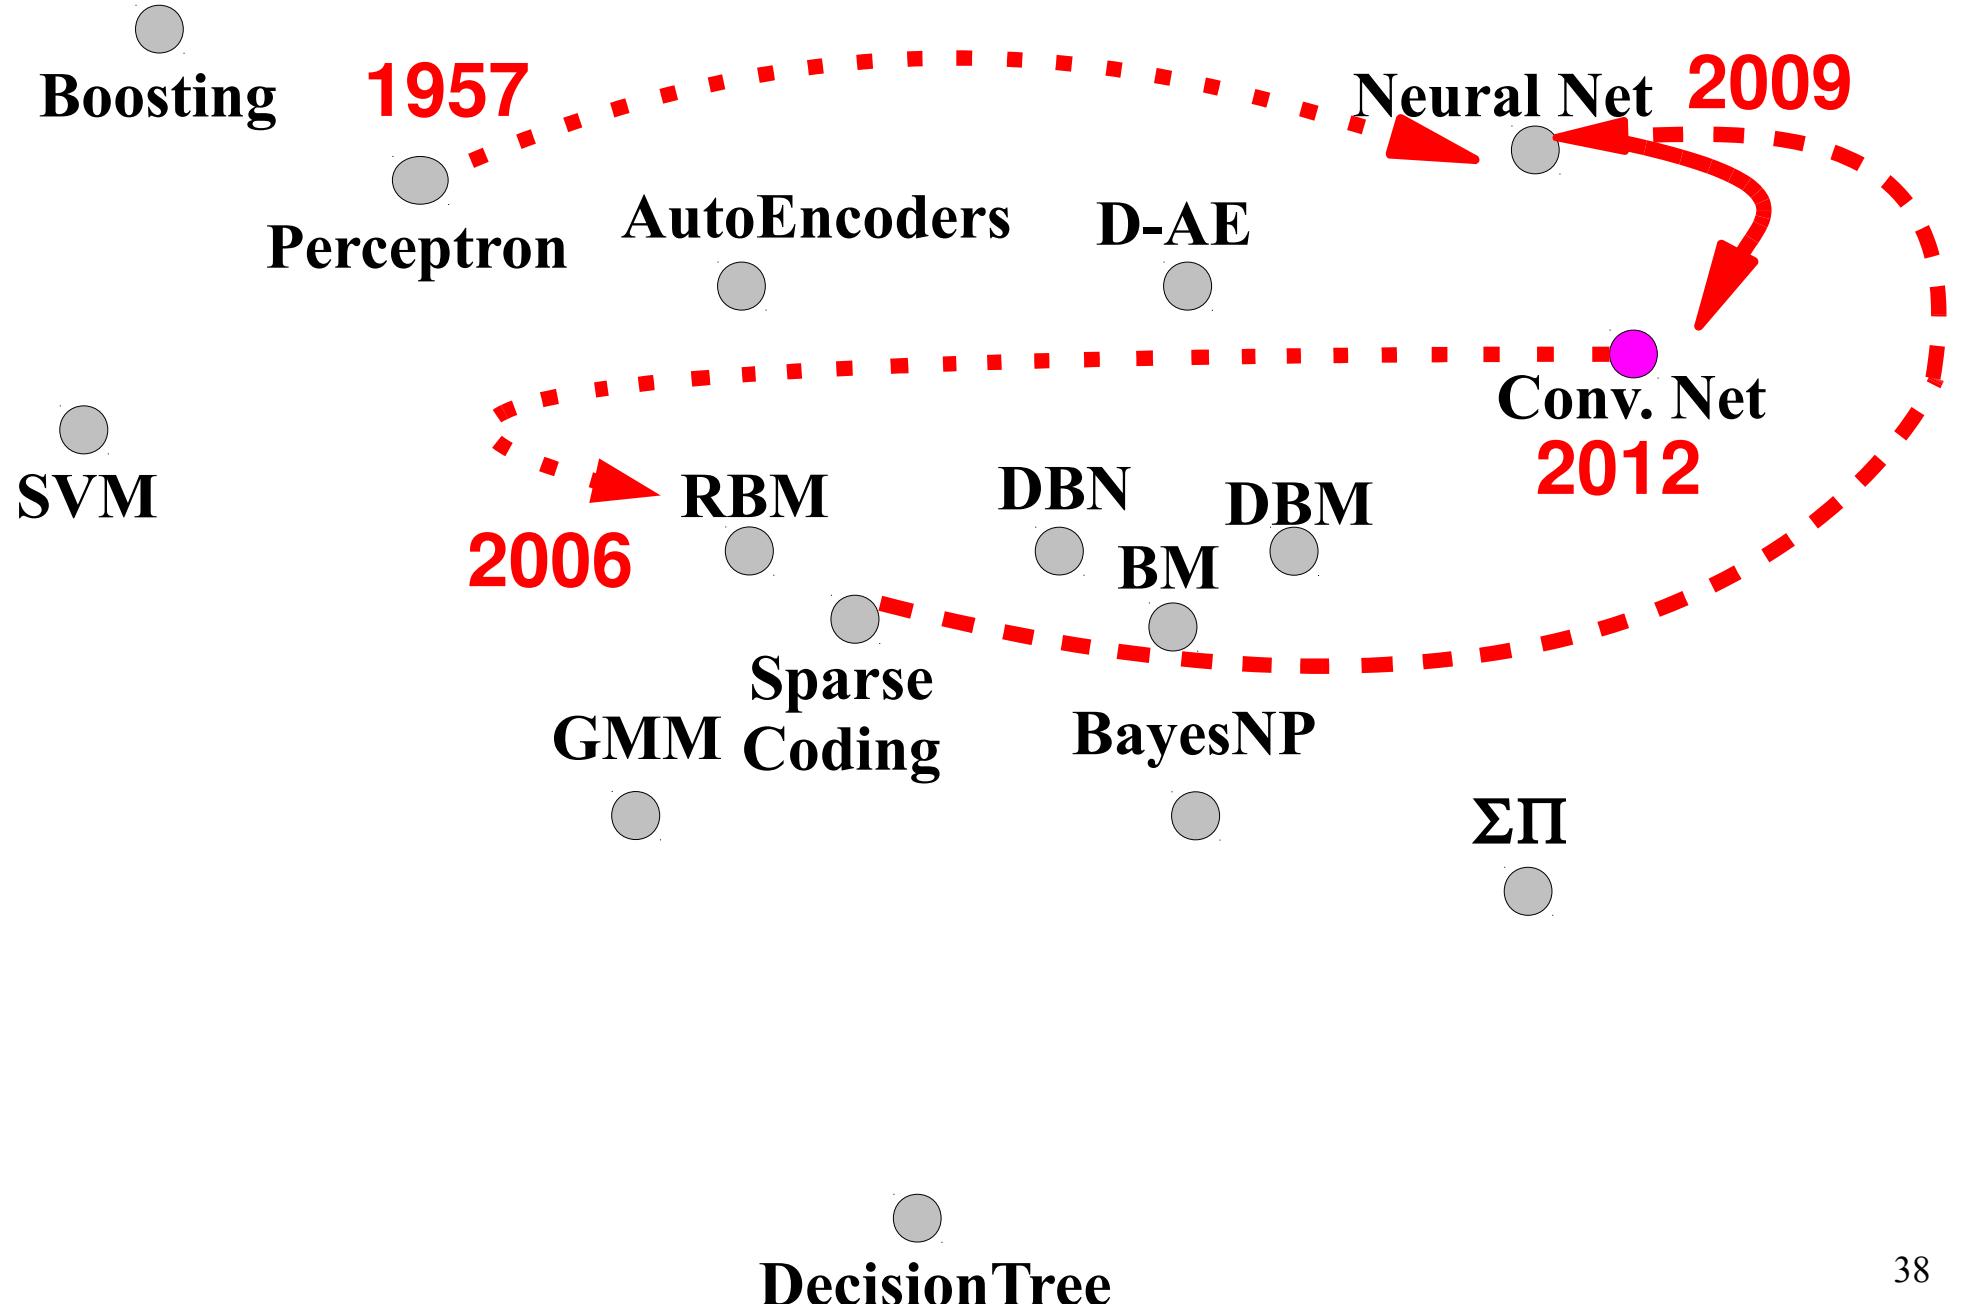
# **THE SPACE OF MACHINE LEARNING METHODS**





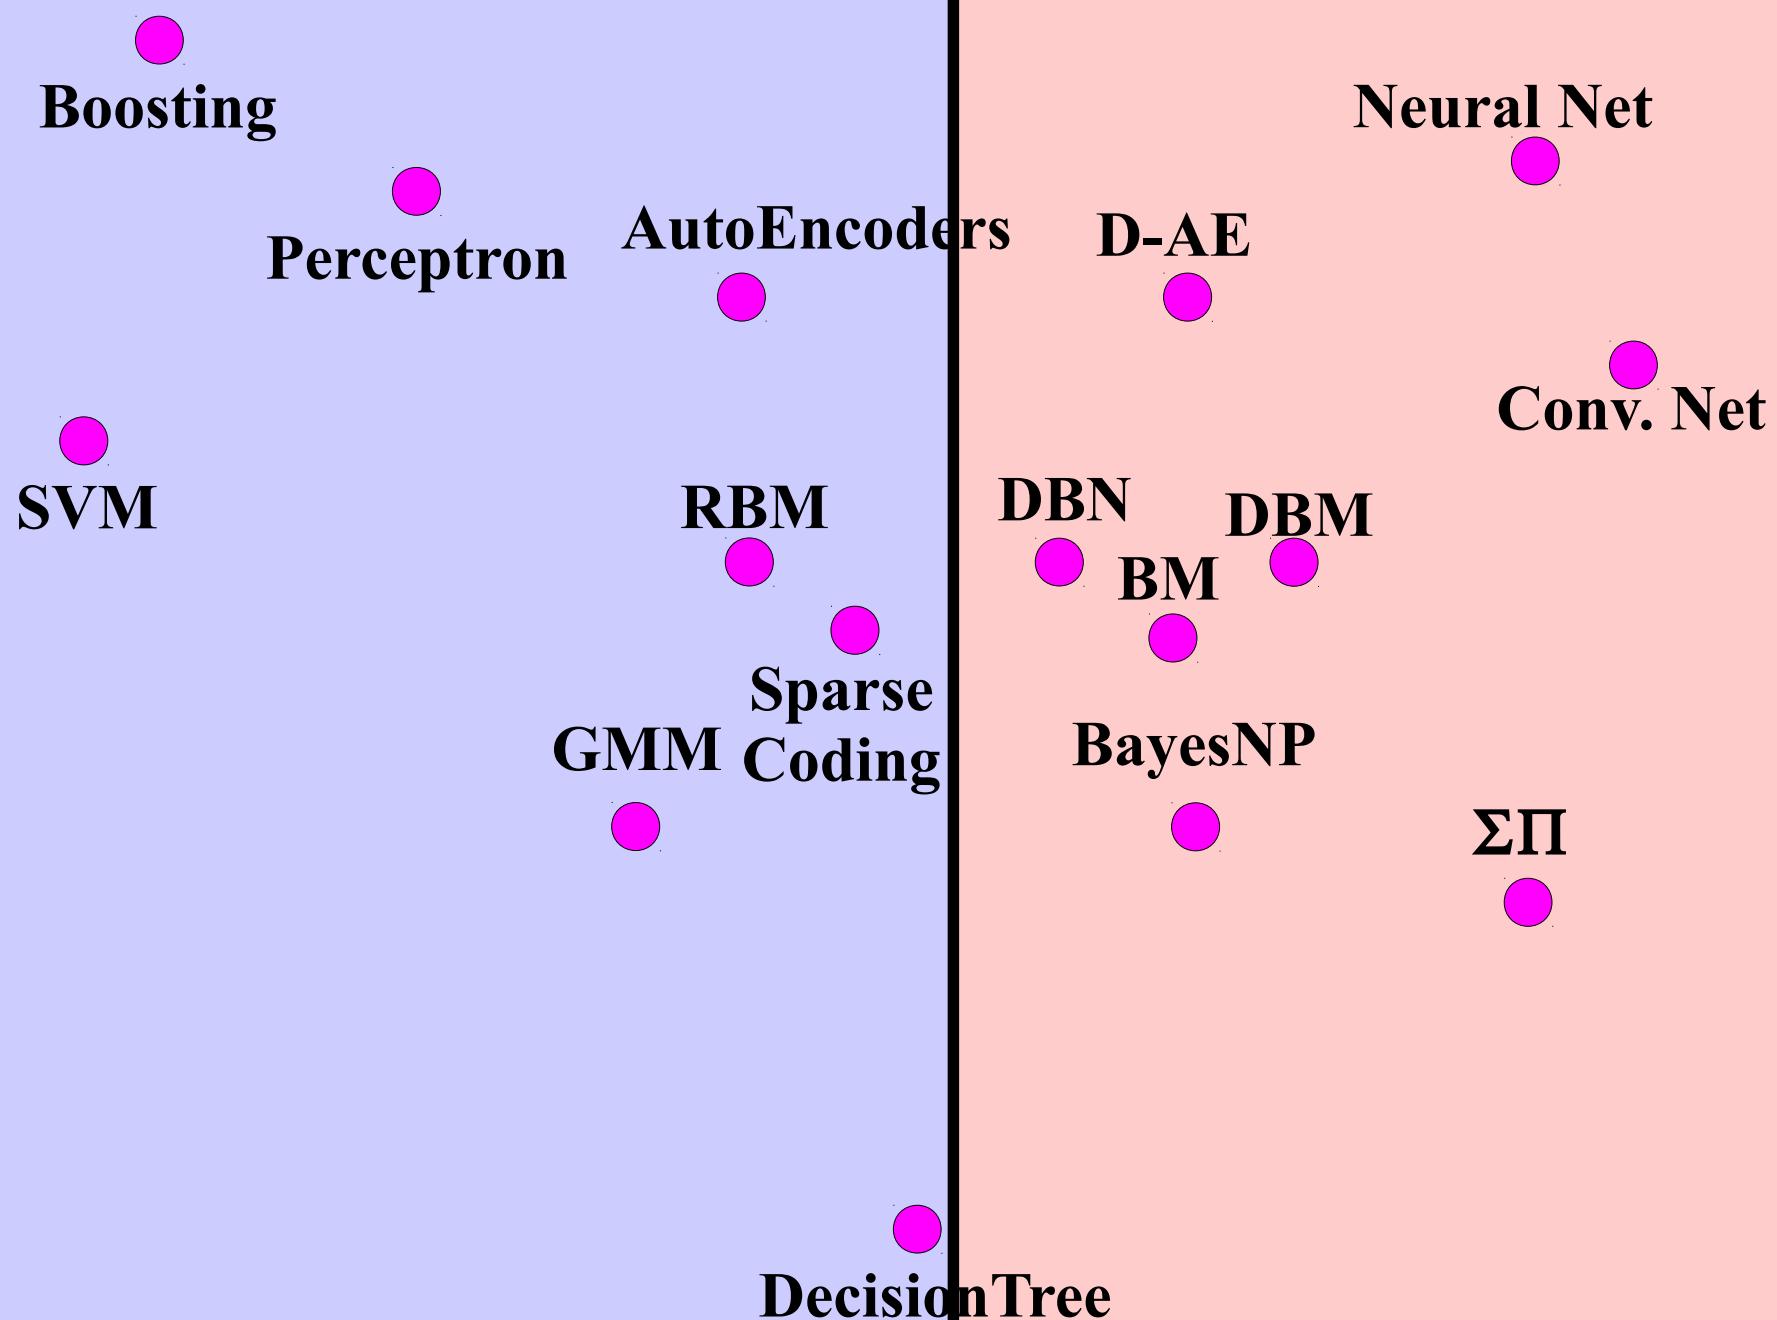






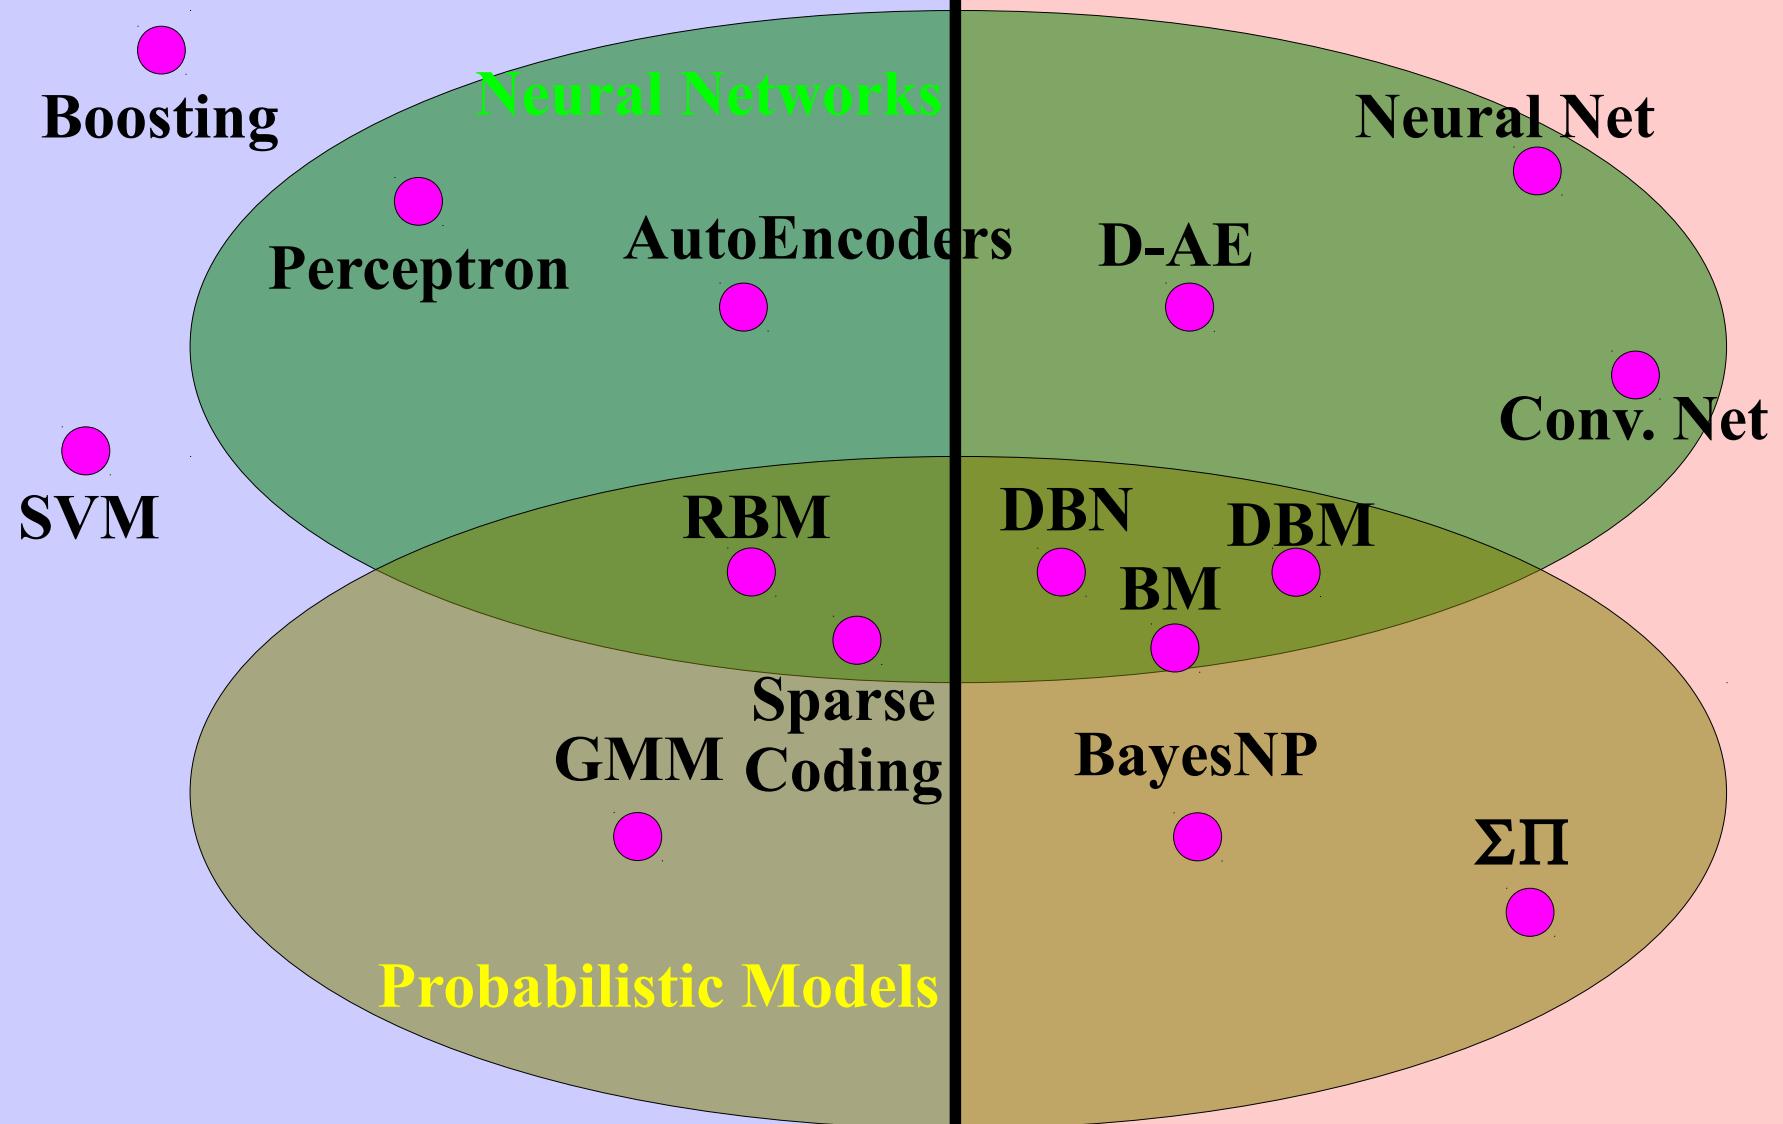
**SHALLOW**

**DEEP**



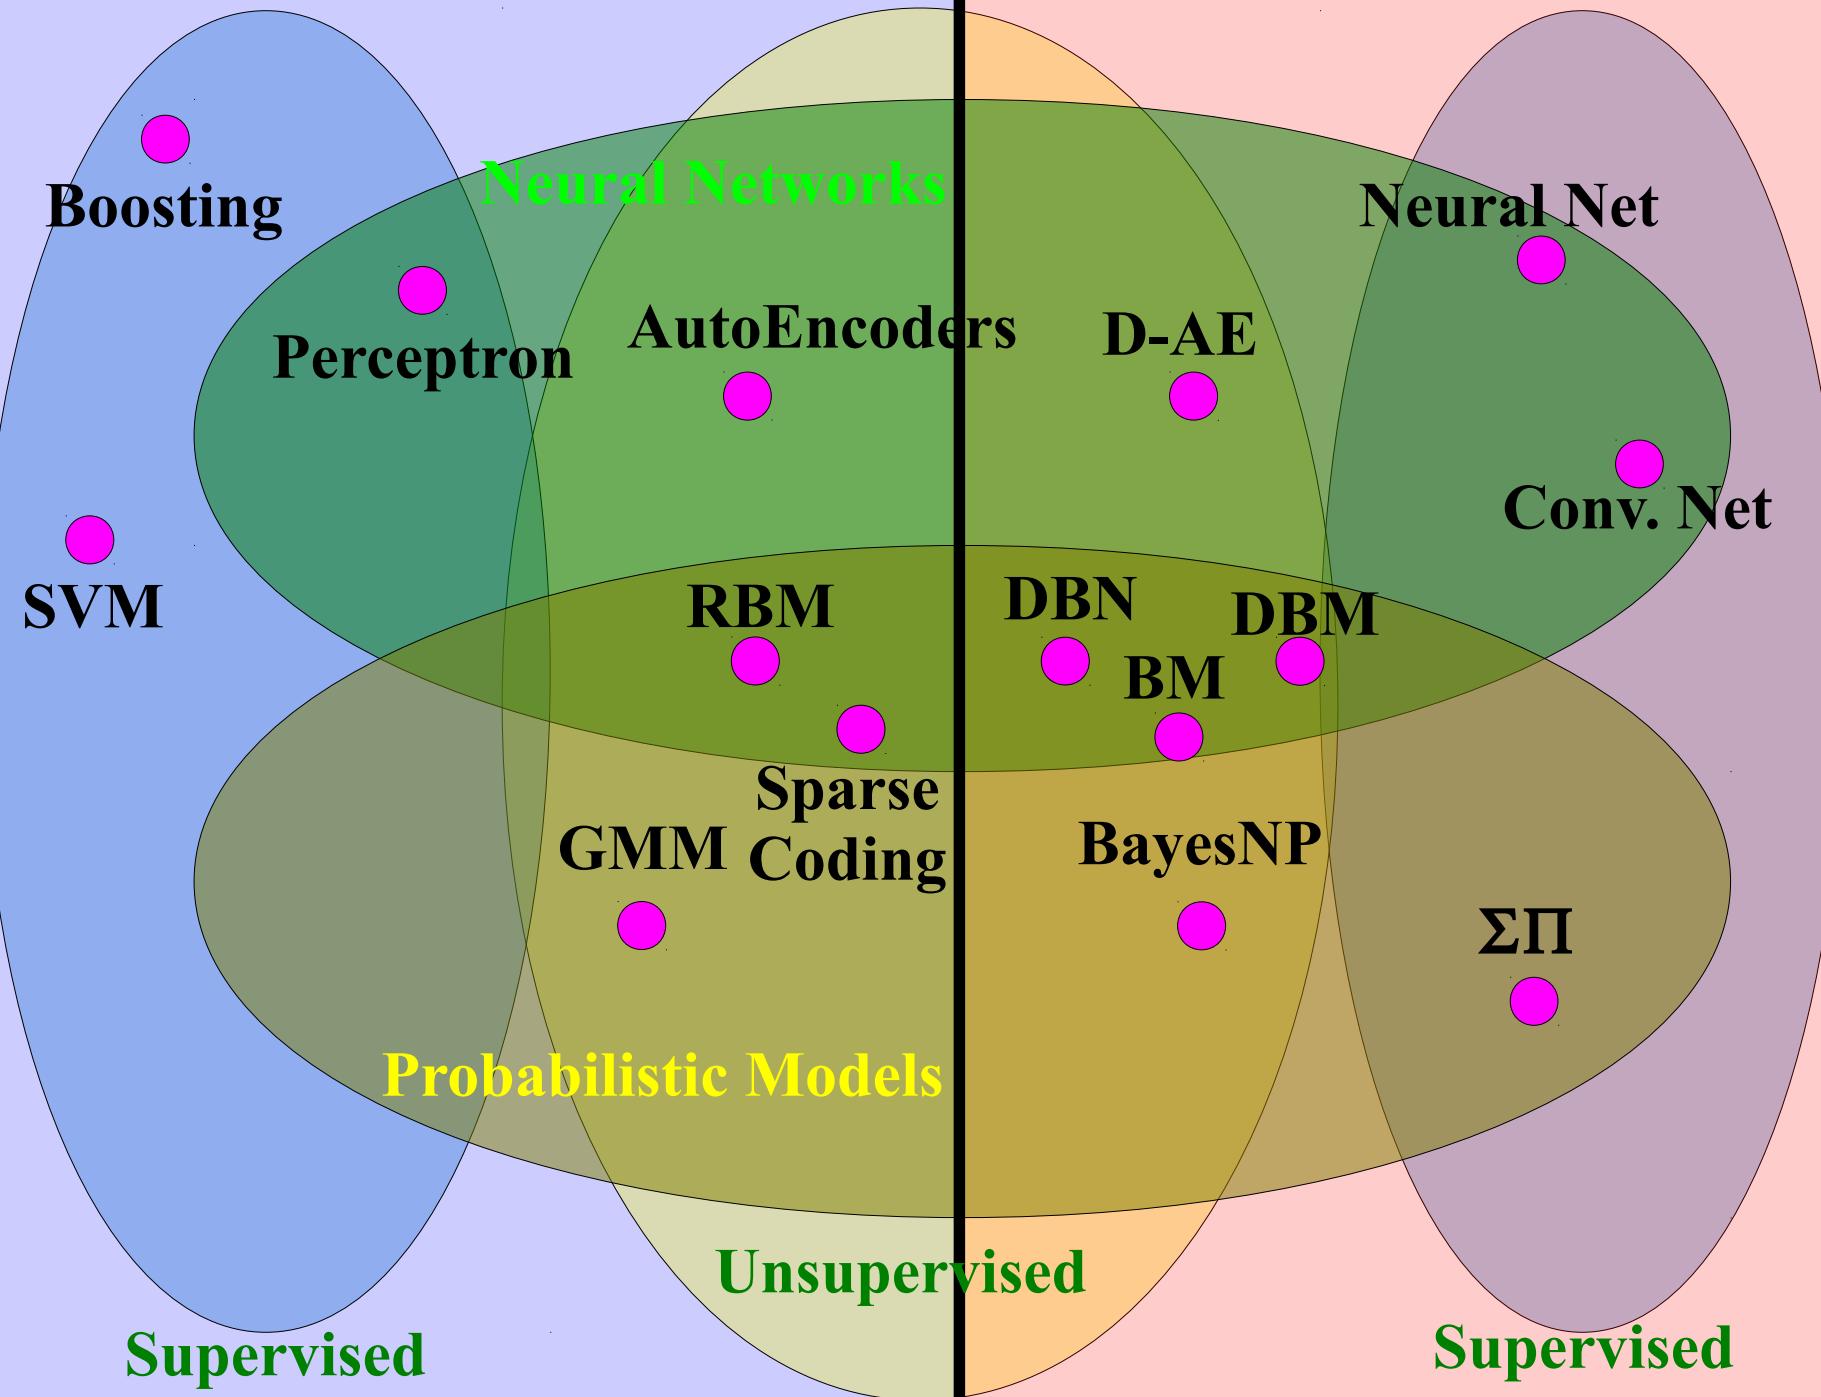
**SHALLOW**

**DEEP**



**SHALLOW**

**DEEP**



In this talk, we'll focus on **convolutional networks**.

# Outline

- Motivation
- Deep Learning: The Big Picture
- From neural nets to convolutional nets
- Applications
- A practical guide

# Linear Classifier: SVM

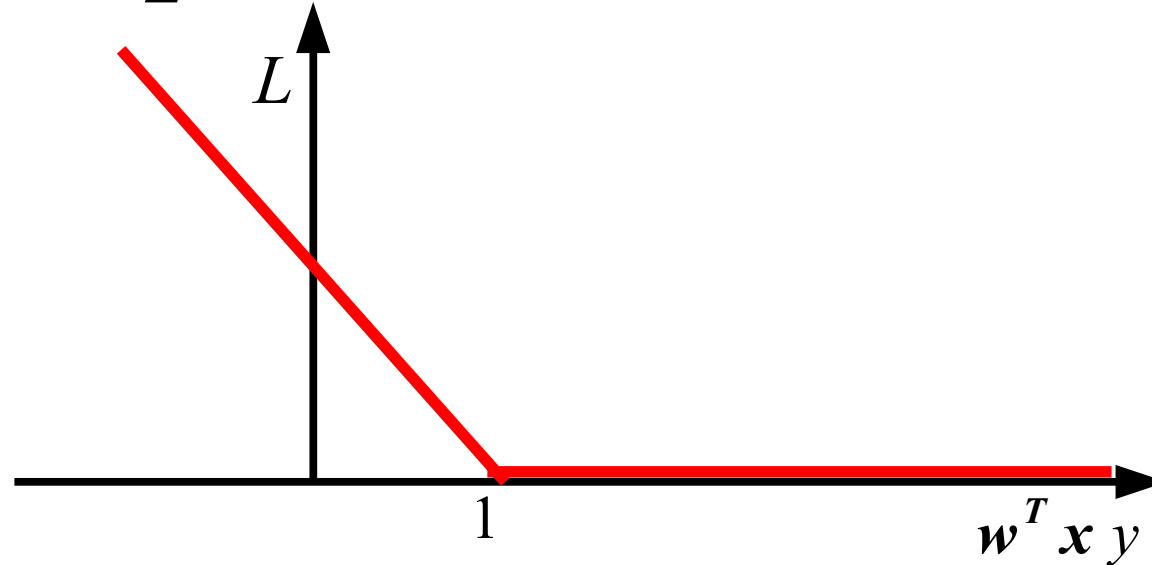
Input:  $x \in R^D$

Binary label:  $y \in \{-1, +1\}$

Parameters:  $w \in R^D$

Output prediction:  $w^T x$

Loss:  $L = \frac{1}{2} \|w\|^2 + \lambda \max[0, 1 - w^T x y]$



Hinge Loss

# Linear Classifier: Logistic Regression

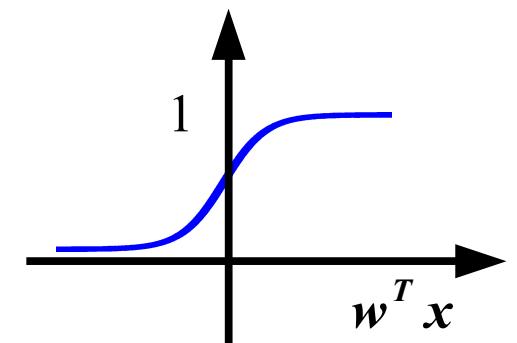
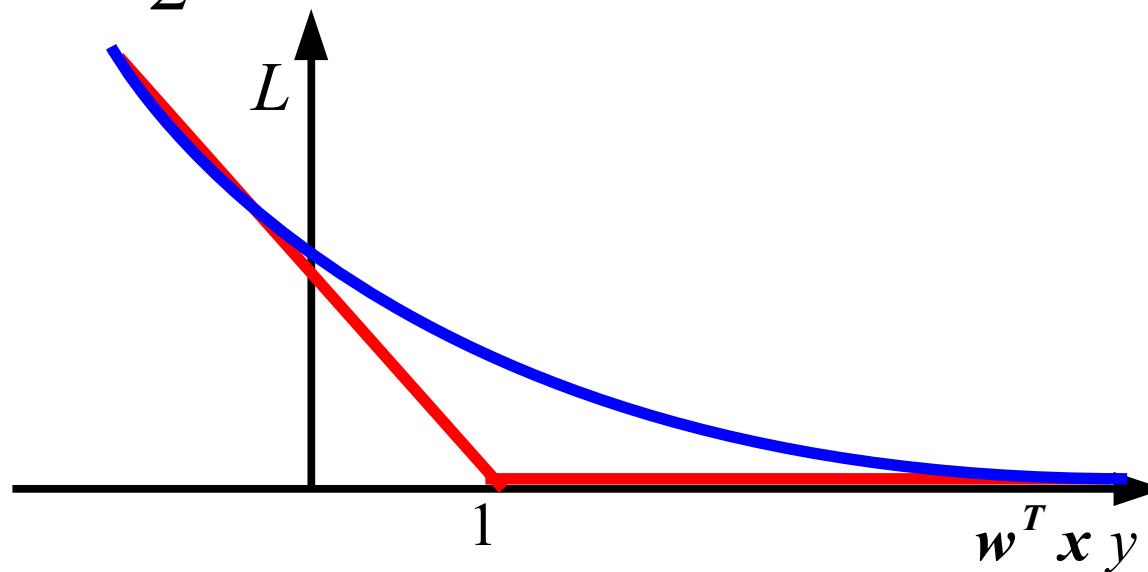
Input:  $x \in R^D$

Binary label:  $y \in \{-1, +1\}$

Parameters:  $w \in R^D$

Output prediction:  $p(y=1|x) = \frac{1}{1+e^{-w^T x}}$

Loss:  $L = \frac{1}{2} \|w\|^2 + \lambda \log(1 + \exp(-w^T x y))$



Log Loss

# Linear Classifier: Logistic Regression

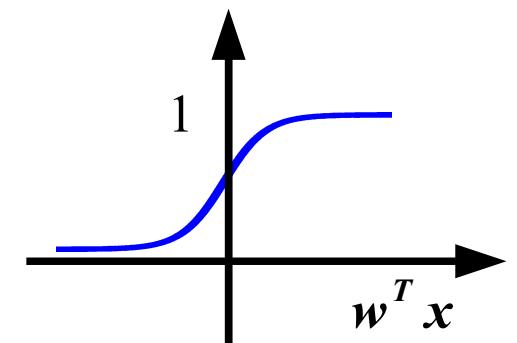
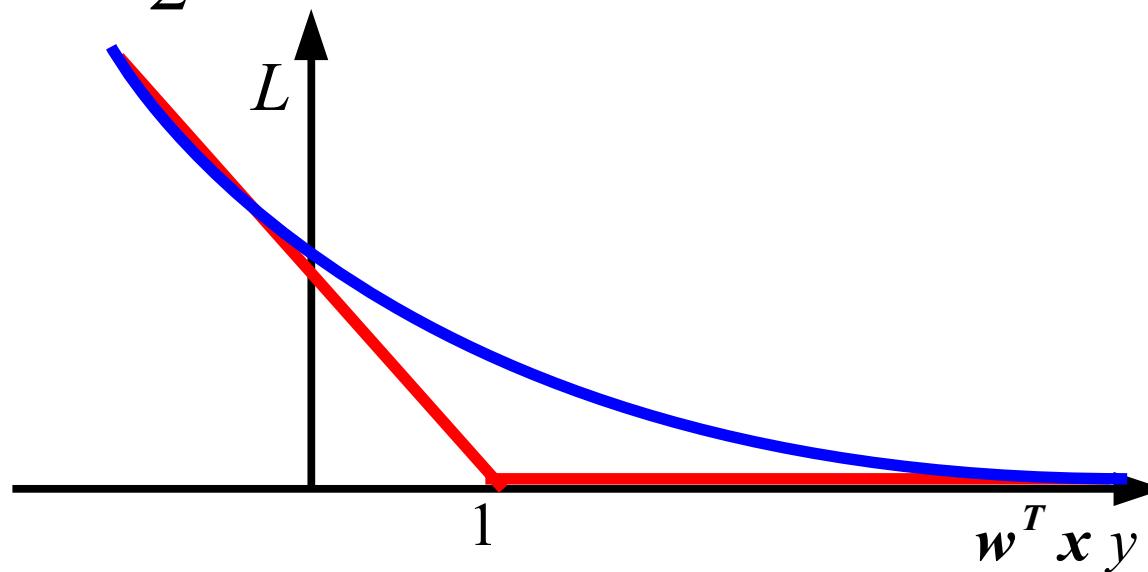
Input:  $x \in R^D$

Binary label:  $y \in \{-1, +1\}$

Parameters:  $w \in R^D$

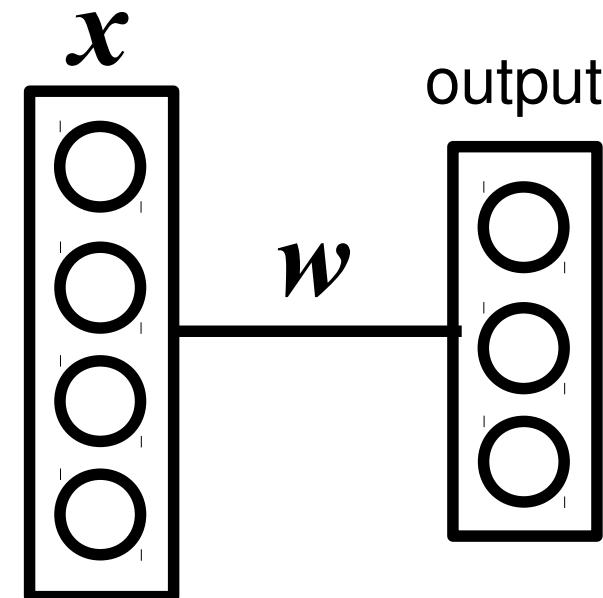
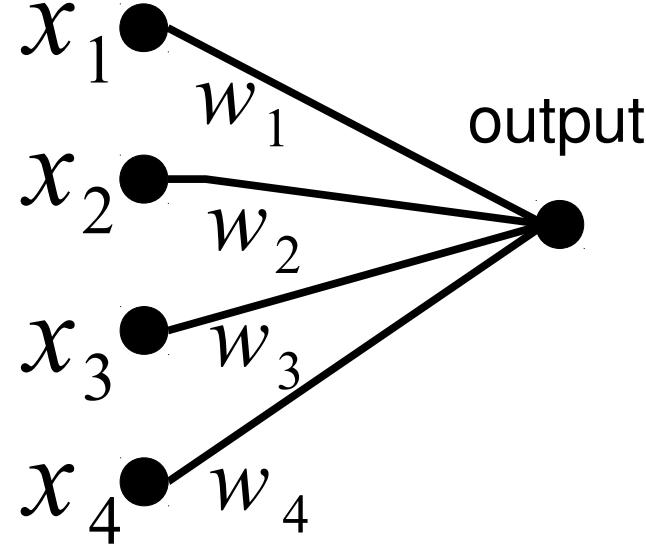
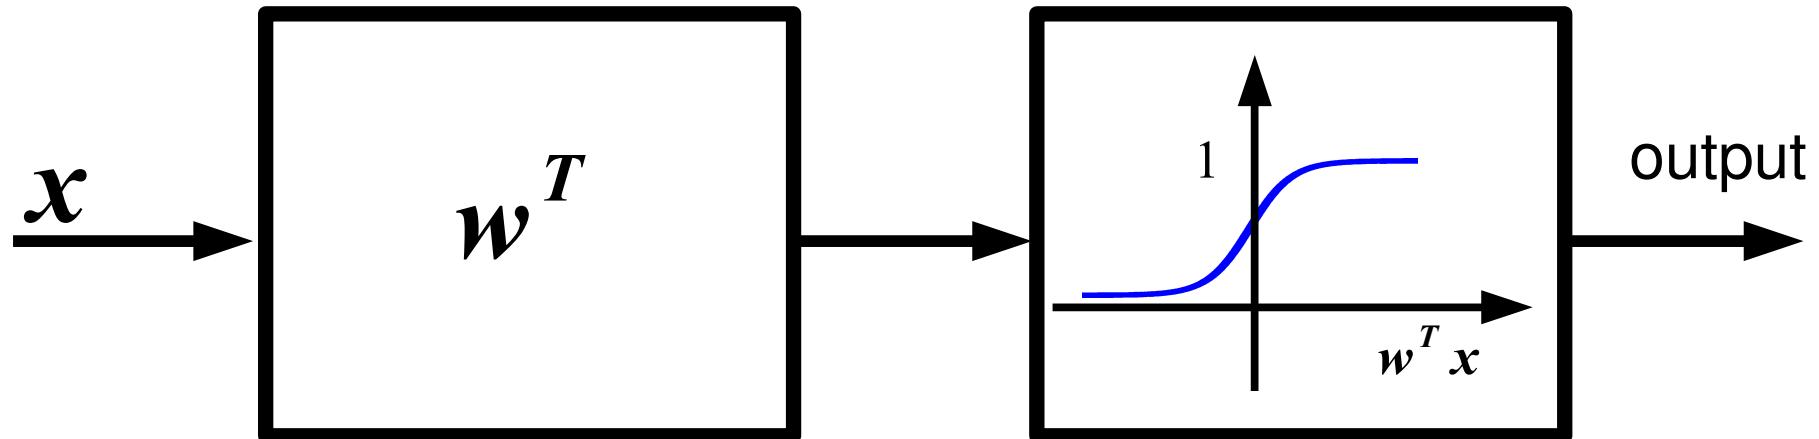
Output prediction:  $p(y=1|x) = \frac{1}{1+e^{-w^T x}}$

Loss:  $L = \frac{1}{2} \|w\|^2 - \lambda \log(p(y|x))$

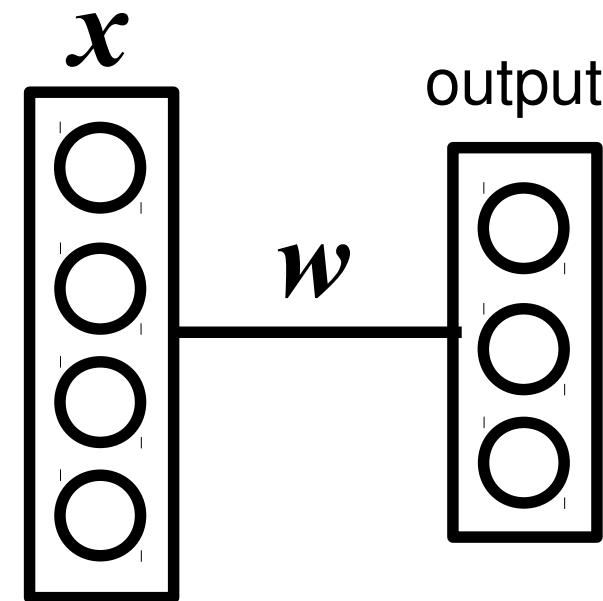
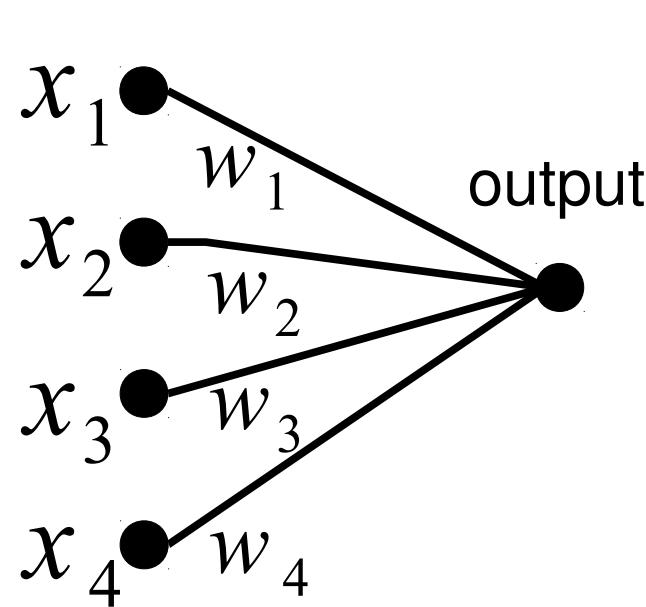
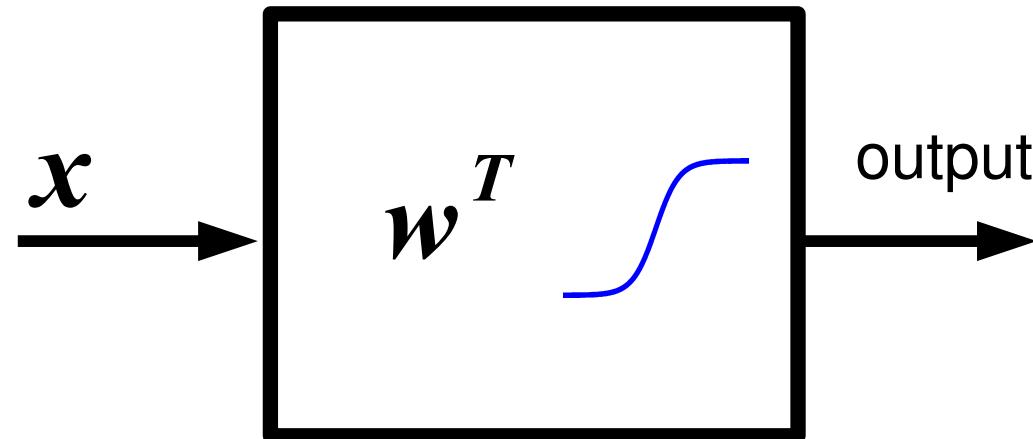


Log Loss

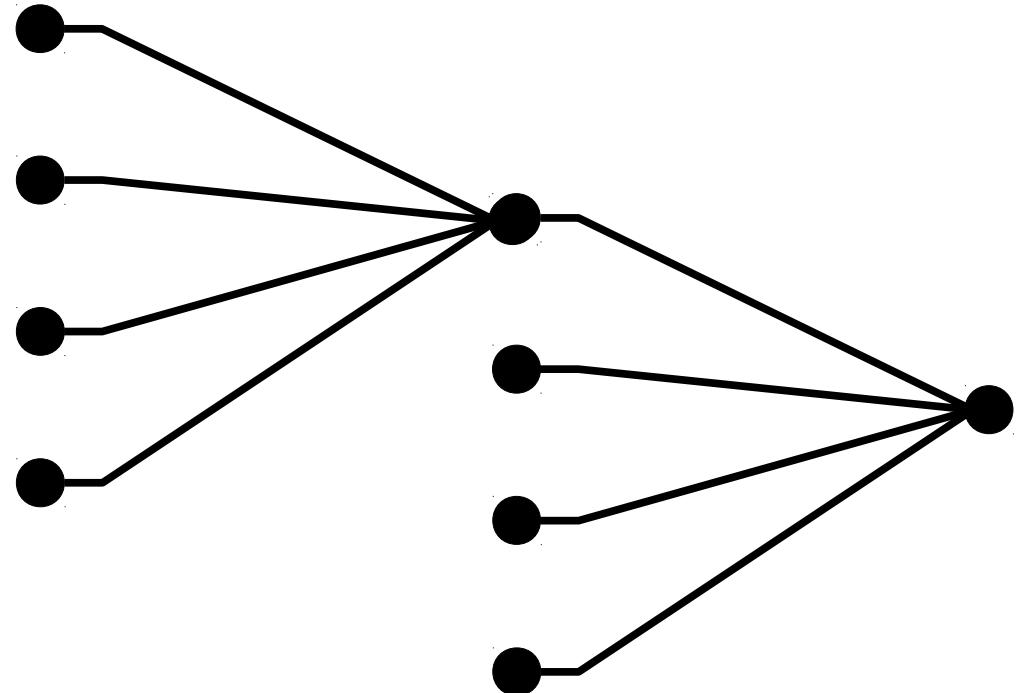
# Graphical Representation



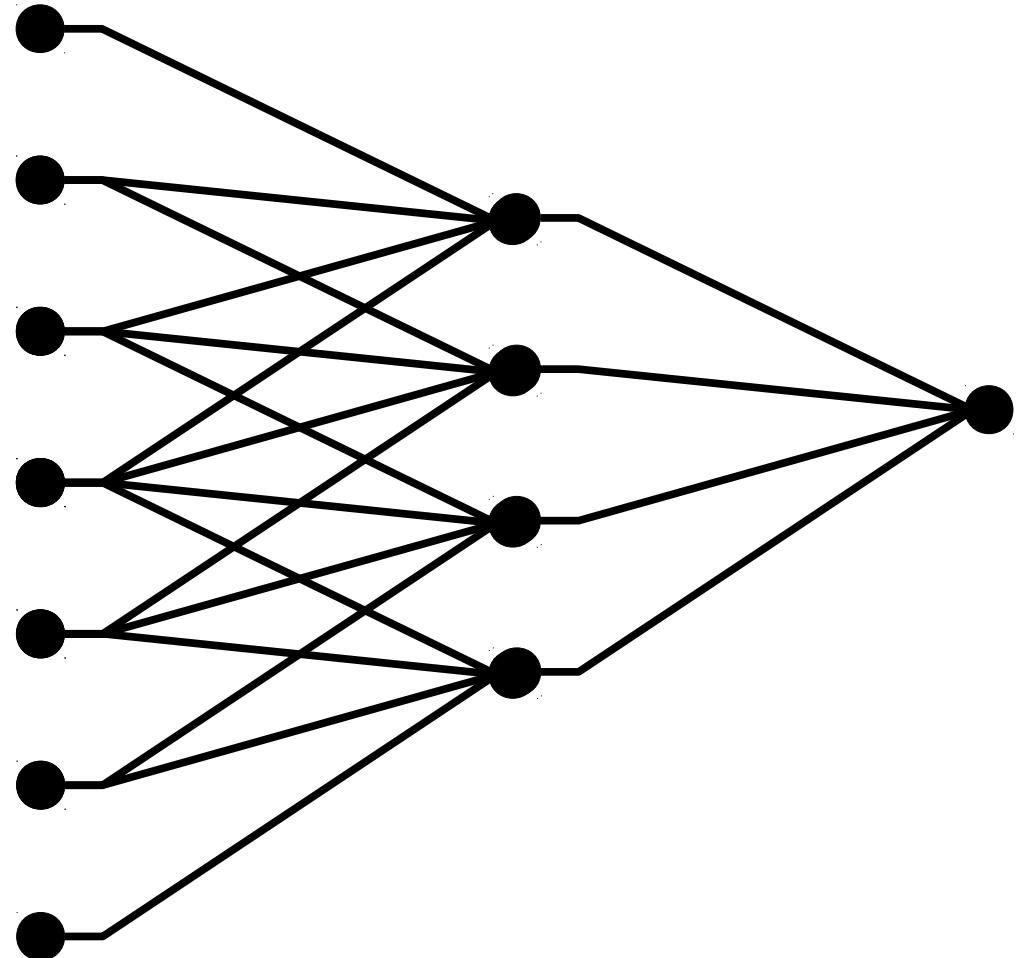
# Graphical Representation



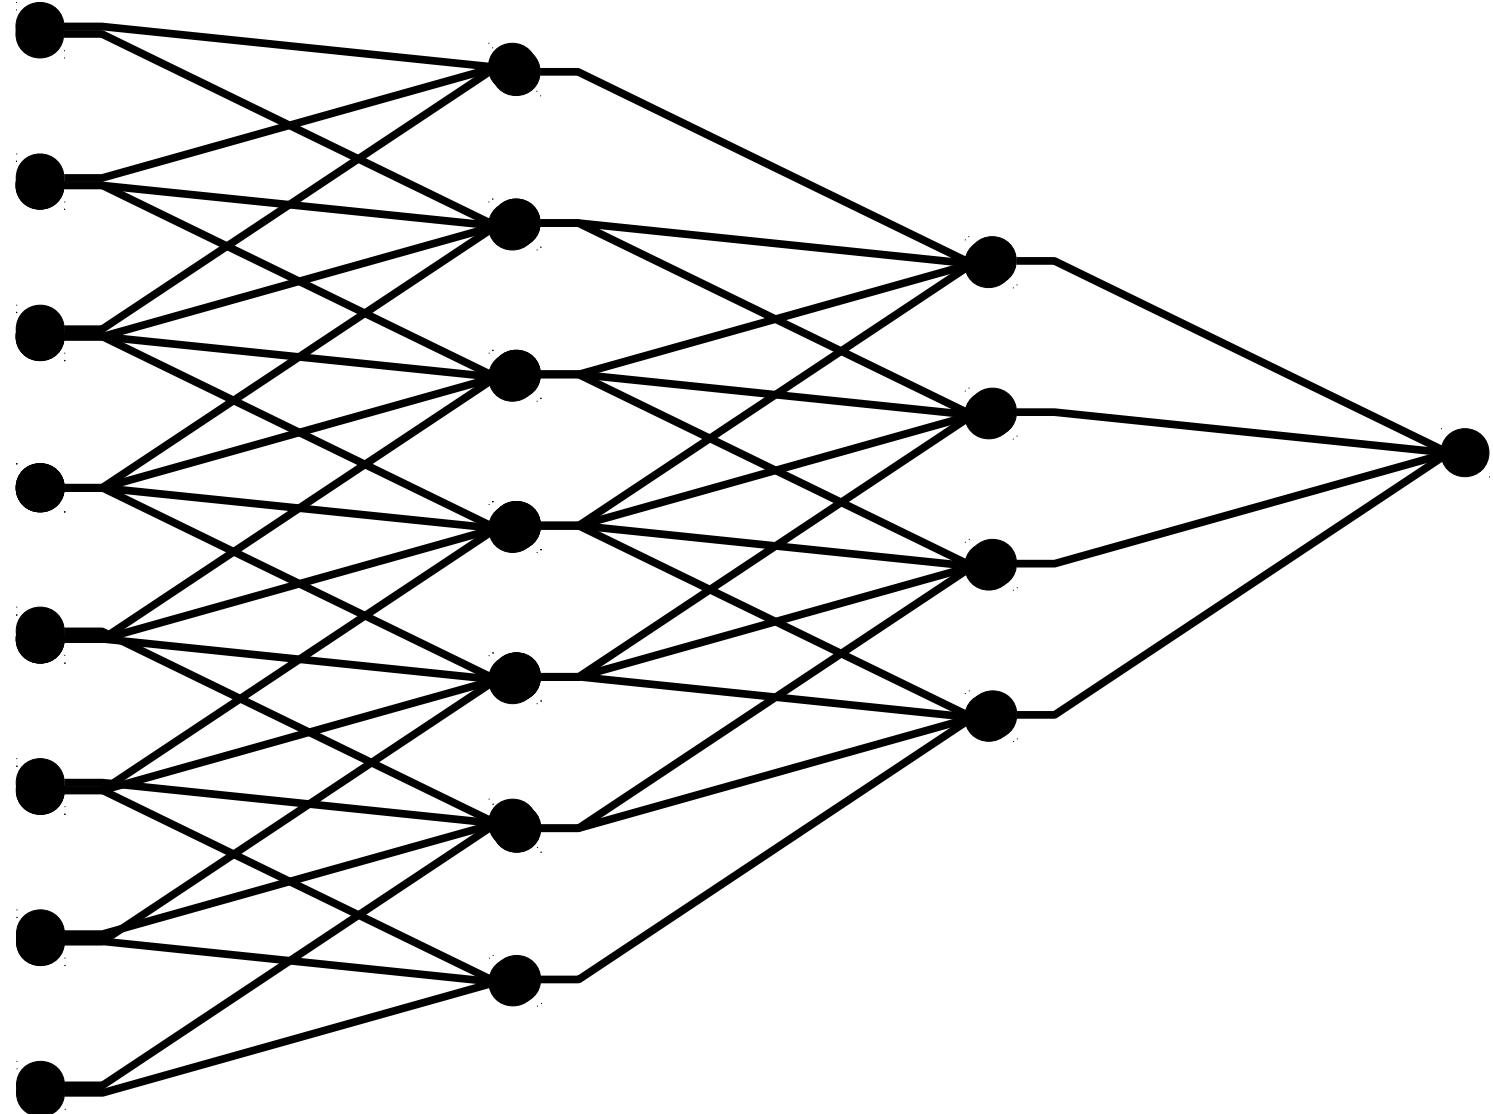
# From Logistic Regression To Neural Nets



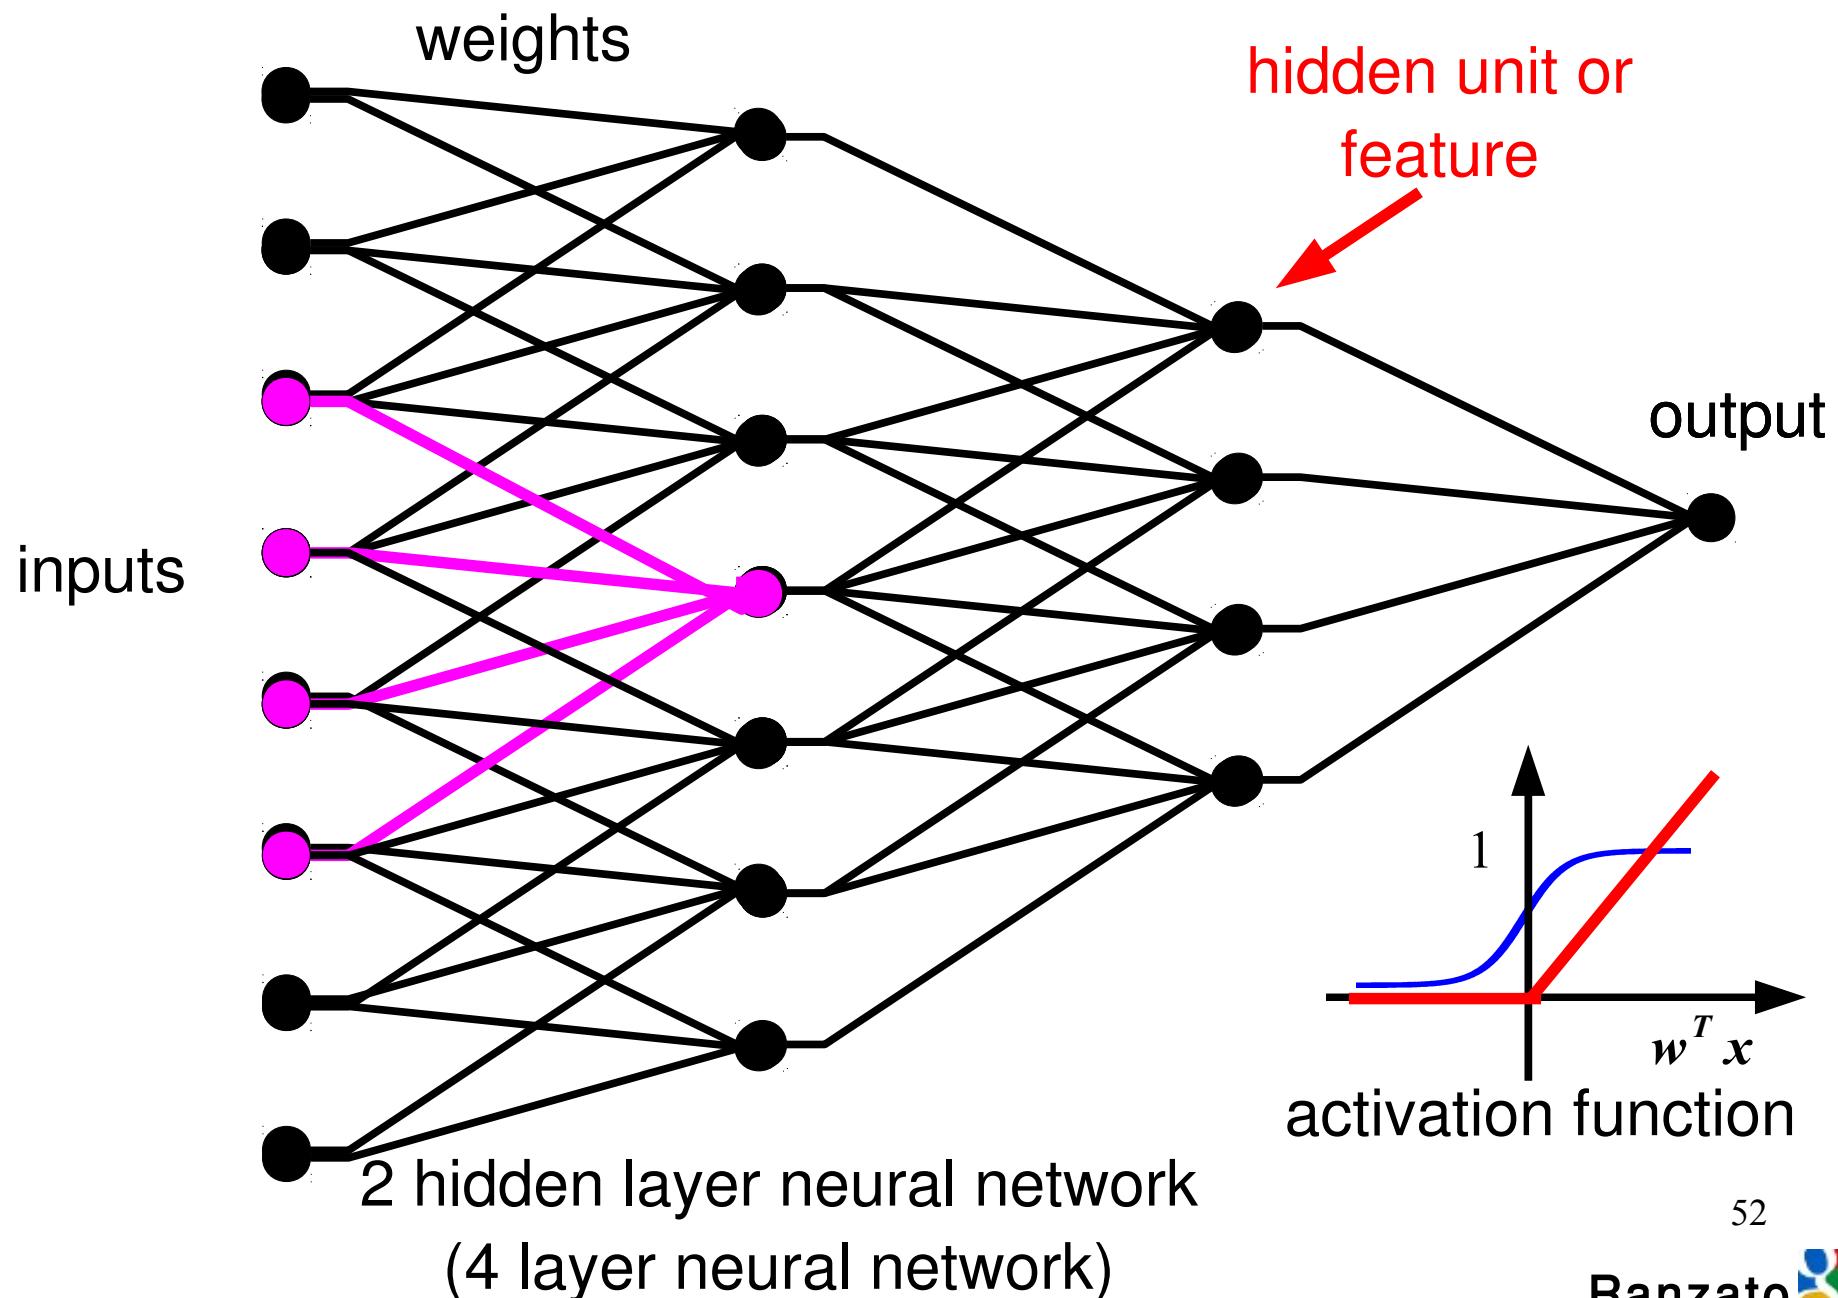
# From Logistic Regression To Neural Nets



# From Logistic Regression To Neural Nets

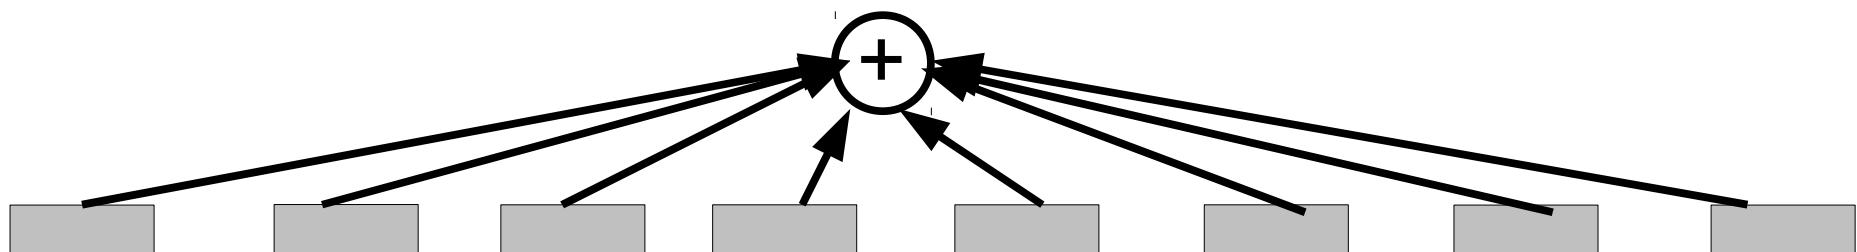


# Neural Network



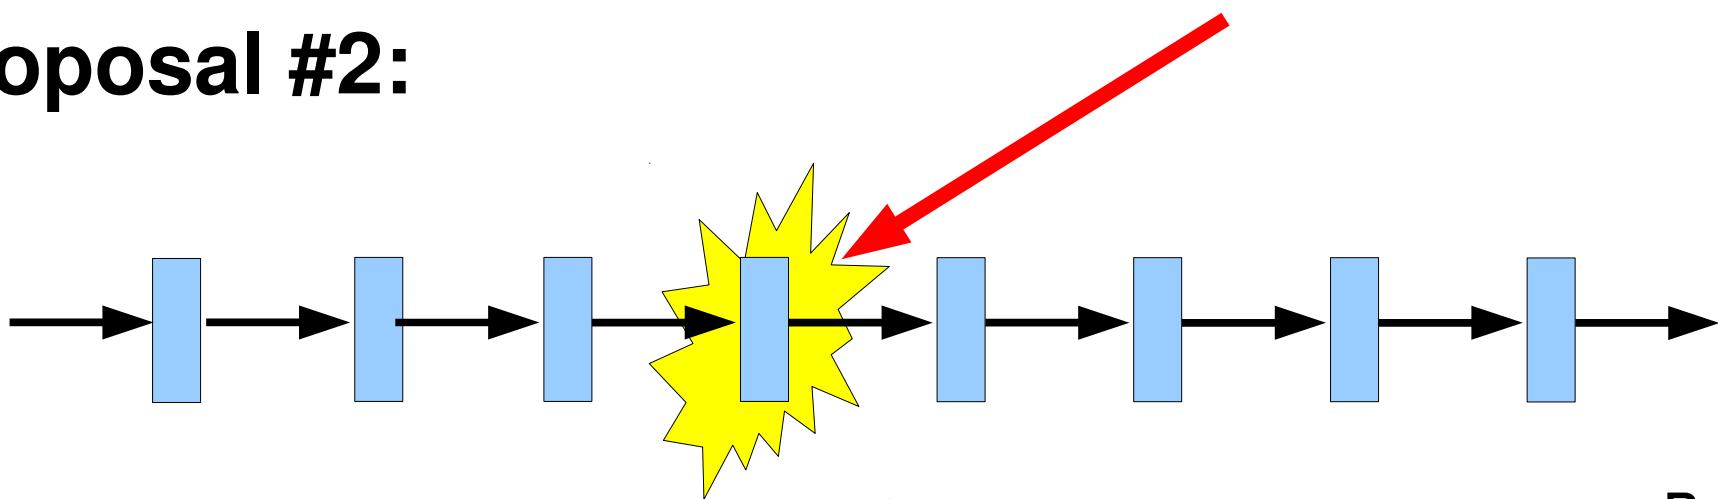
# Learning Non-Linear Features

Proposal #1:

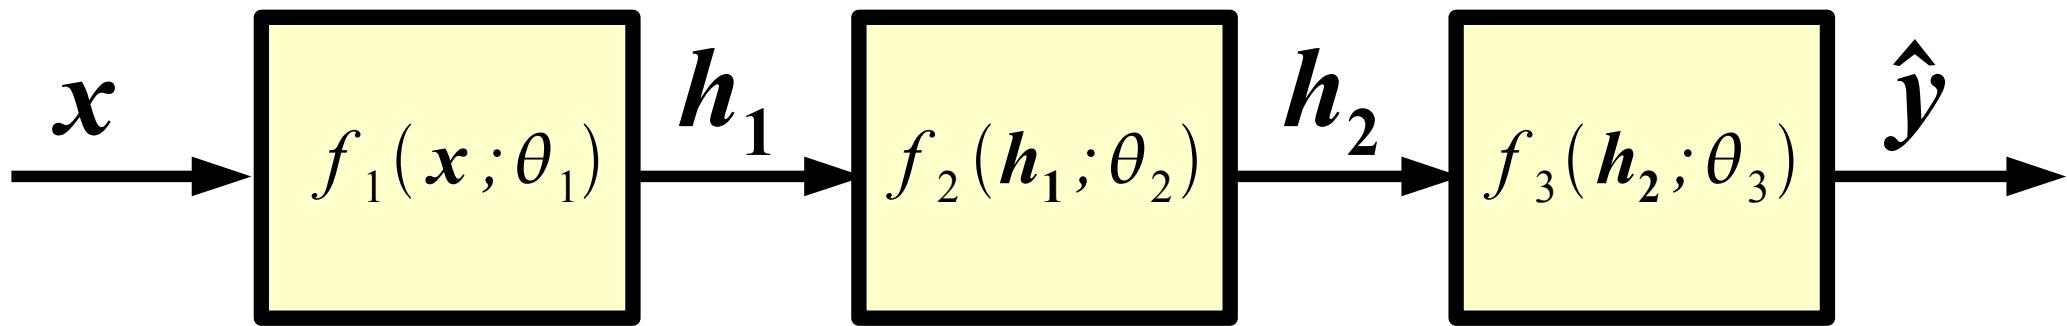


Each of box is a feature detector

Proposal #2:



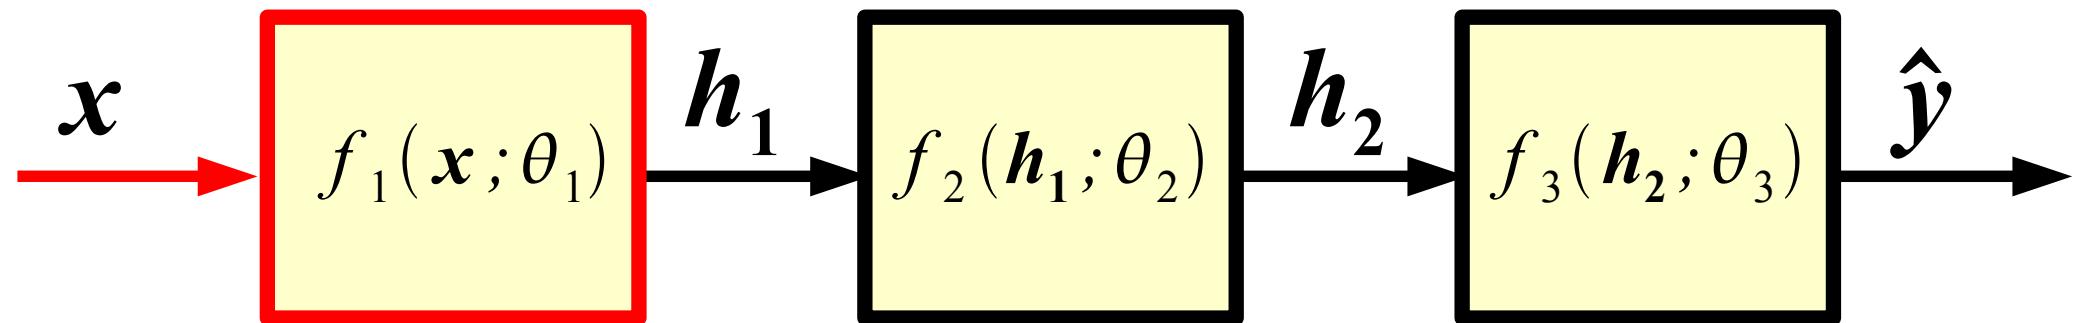
# Neural Nets



**NOTE:** In practice, each module does NOT need to be a logistic regression classifier.

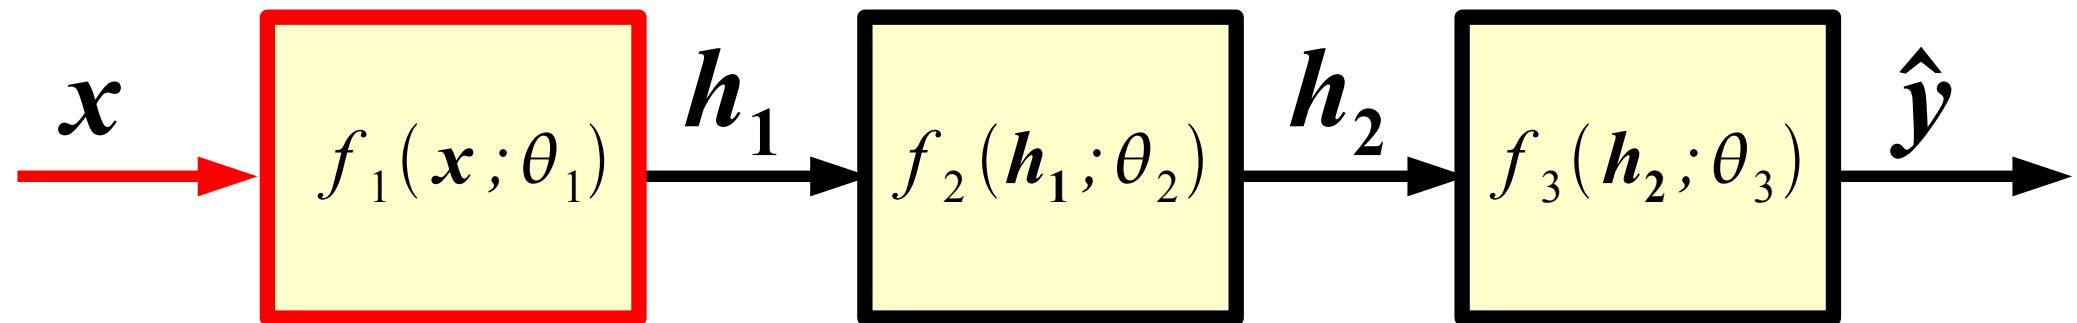
Any (a.e. differentiable) non-linear transformation is potentially good.

# Forward Propagation (FPROP)



1) Given  $x$  compute:  $h_1 = f_1(x; \theta_1)$

# Forward Propagation (FPROP)

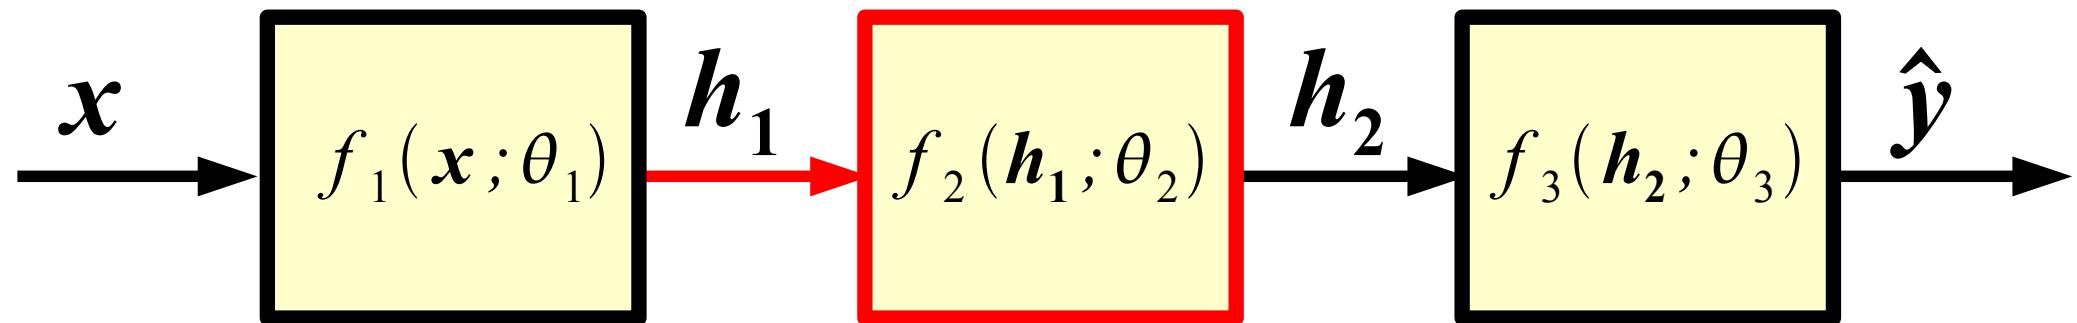


1) Given  $x$  compute:  $h_1 = f_1(x; \theta_1)$

For instance,

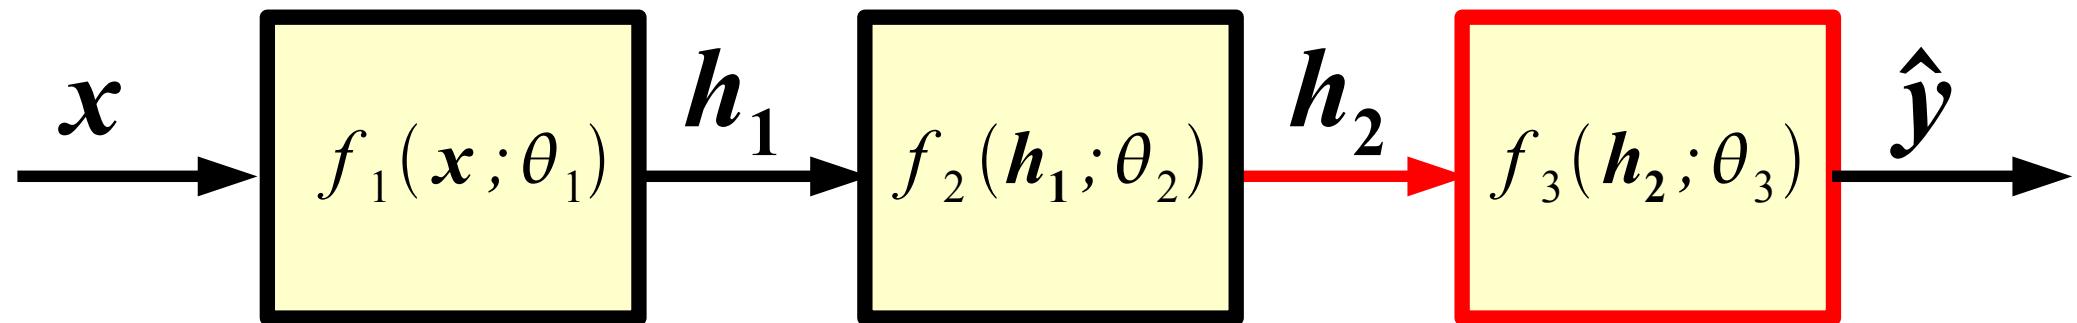
$$h_1 = \max(0, W_1 x + b_1)$$

# Forward Propagation (FPROP)



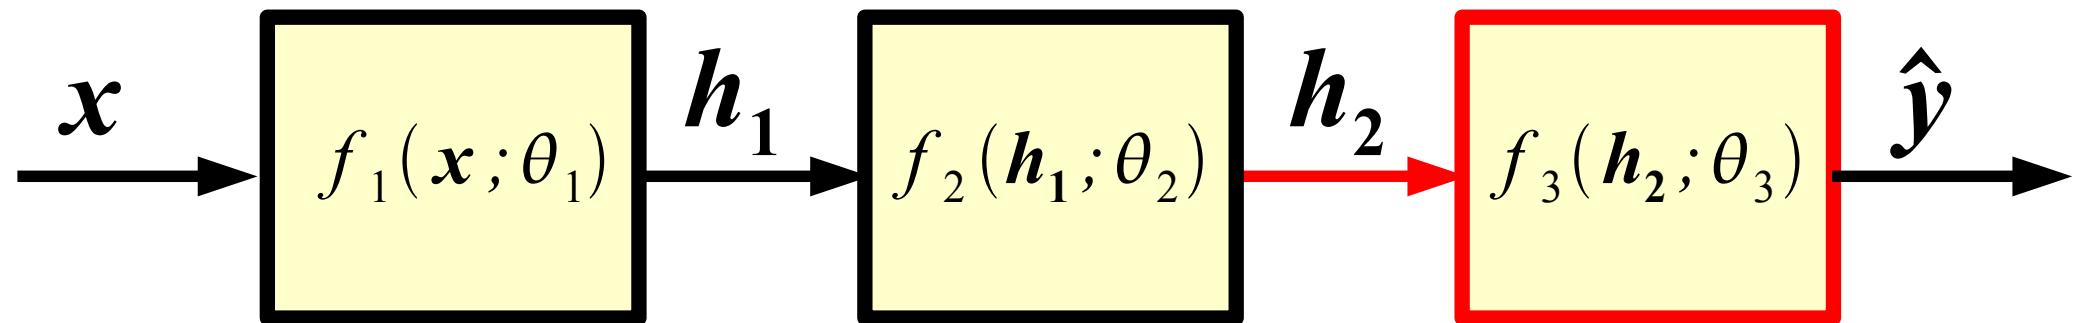
- 1) Given  $x$  compute:  $h_1 = f_1(x; \theta_1)$
- 2) Given  $h_1$  compute:  $h_2 = f_2(h_1; \theta_2)$

# Forward Propagation (FPROP)



- 1) Given  $x$  compute:  $h_1 = f_1(x; \theta_1)$
- 2) Given  $h_1$  compute:  $h_2 = f_2(h_1; \theta_2)$
- 3)** Given  $h_2$  compute:  $\hat{y} = f_3(h_2; \theta_3)$

# Forward Propagation (FPROP)

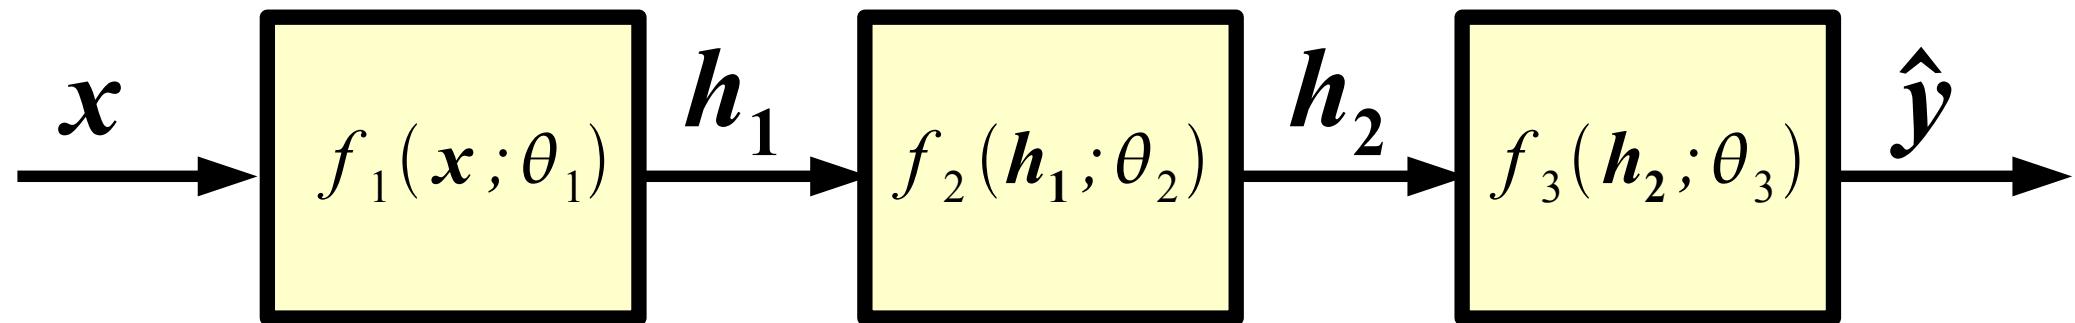


- 1) Given  $x$  compute:  $h_1 = f_1(x; \theta_1)$
- 2) Given  $h_1$  compute:  $h_2 = f_2(h_1; \theta_2)$
- 3)** Given  $h_2$  compute:  $\hat{y} = f_3(h_2; \theta_3)$

For instance,

$$\hat{y}_i = p(\text{class} = i | x) = \frac{e^{W_{3i} h_2 + b_{3i}}}{\sum_k e^{W_{3k} h_2 + b_{3k}}}$$

# Forward Propagation (FPROP)

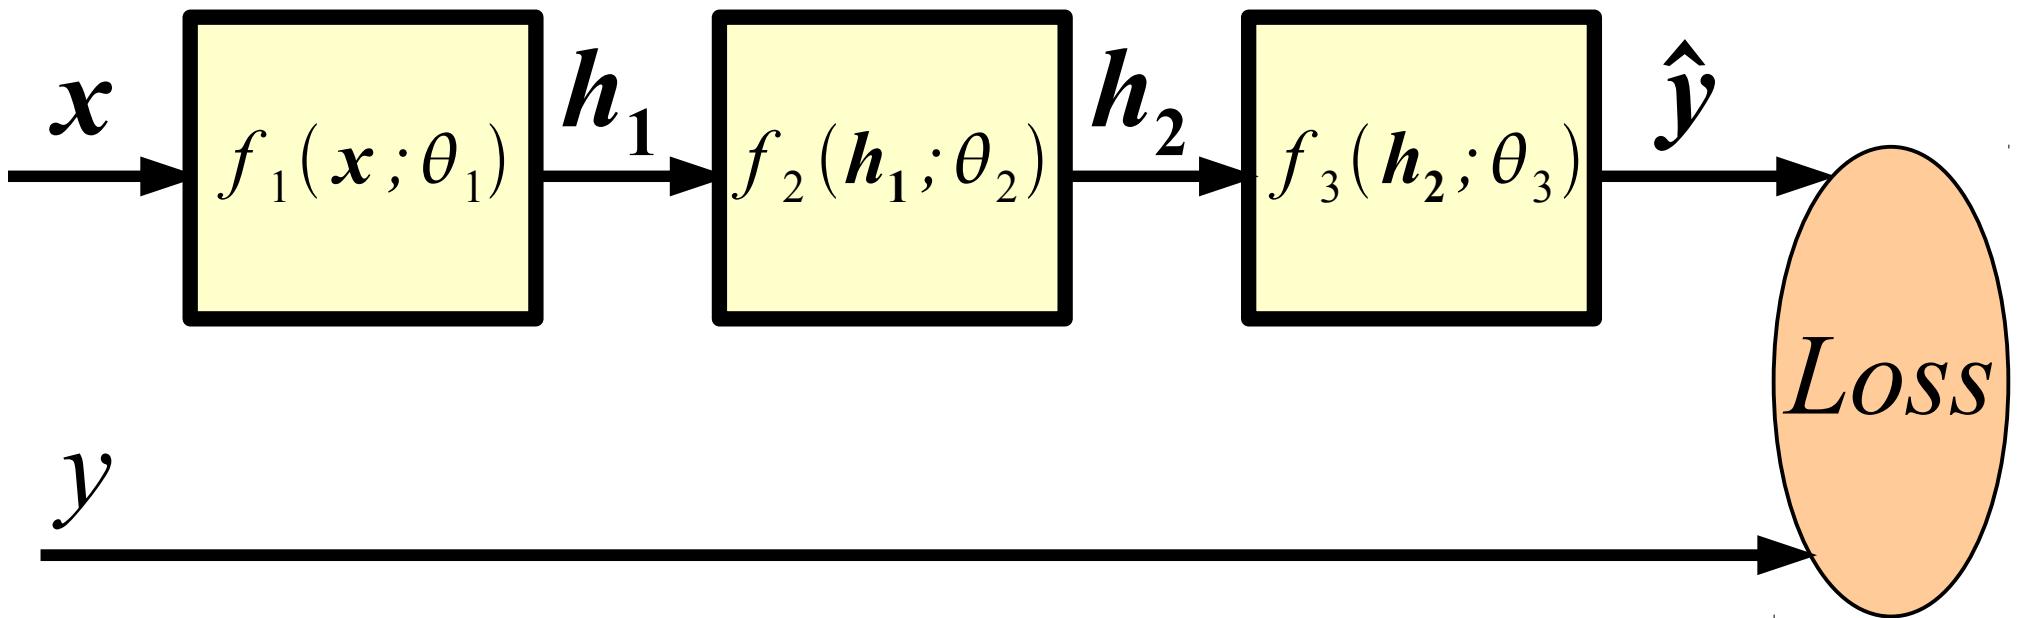


- 1) Given  $x$  compute:  $h_1 = f_1(x; \theta_1)$
- 2) Given  $h_1$  compute:  $h_2 = f_2(h_1; \theta_2)$
- 3) Given  $h_2$  compute:  $\hat{y} = f_3(h_2; \theta_3)$

This is the typical processing at test time.

At training time, we need to compute an error measure and tune the parameters to decrease the error.

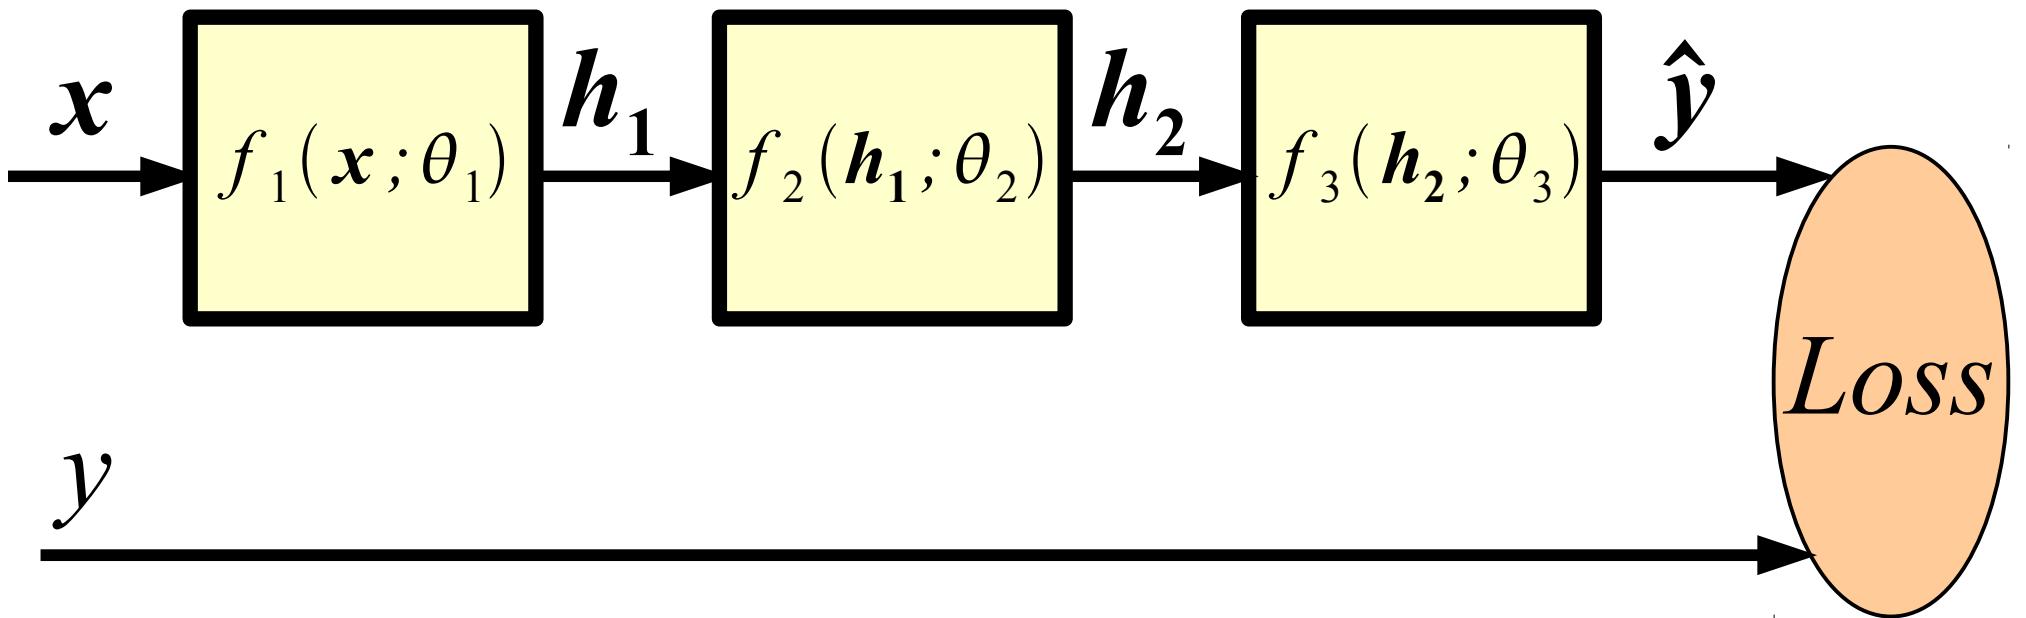
# Loss



The measure of how well the model fits the training set is given by a suitable loss function:  $L(x, y; \theta)$

The loss depends on the input  $x$ , the target label  $y$ , and the parameters  $\theta$ .

# Loss

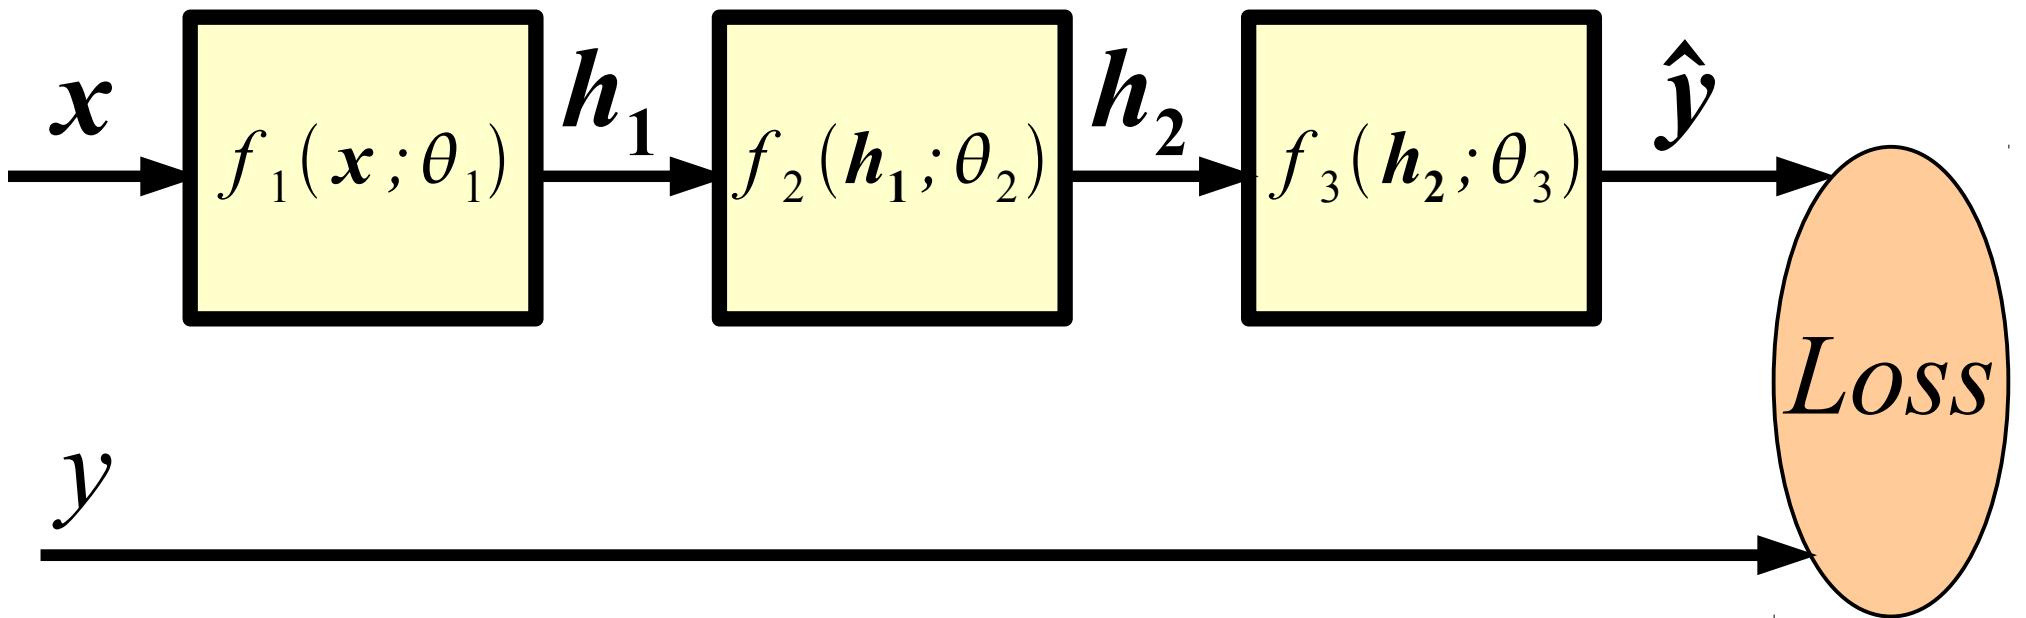


The measure of how well the model fits the training set is given by a suitable loss function:  $L(x, y; \theta)$

For instance,

$$L(x, y=k; \theta) = -\log(p(\text{class}=k|x))$$

# Loss

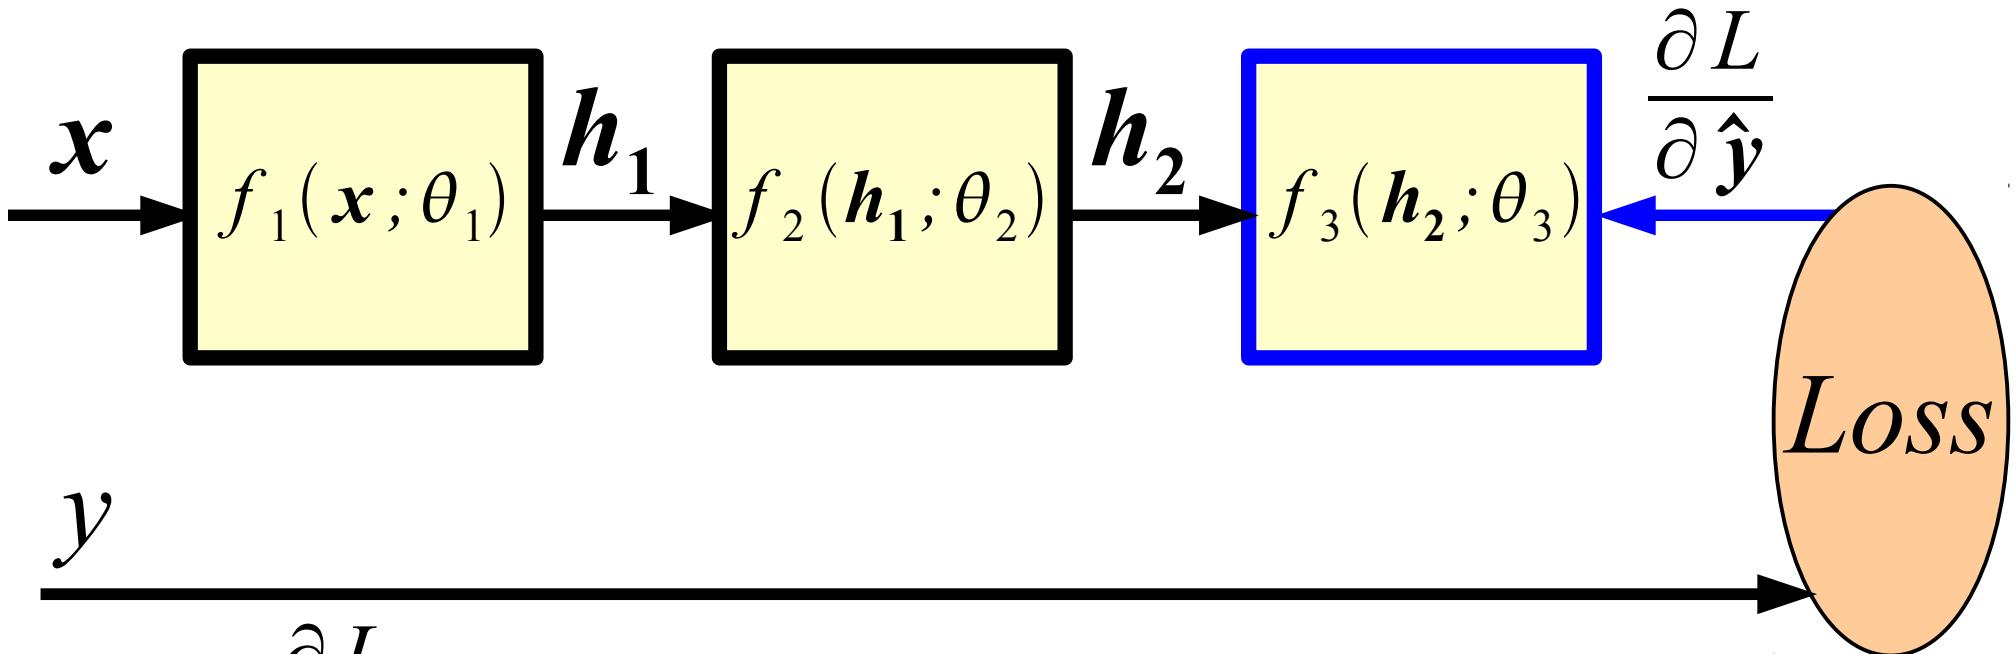


**Q.:** how to tune the parameters to decrease the loss?

If loss is (a.e.) differentiable we can compute gradients.

We can use chain-rule, a.k.a. **back-propagation**, to compute the gradients w.r.t. parameters at the lower layers.

# Backward Propagation (BPROP)

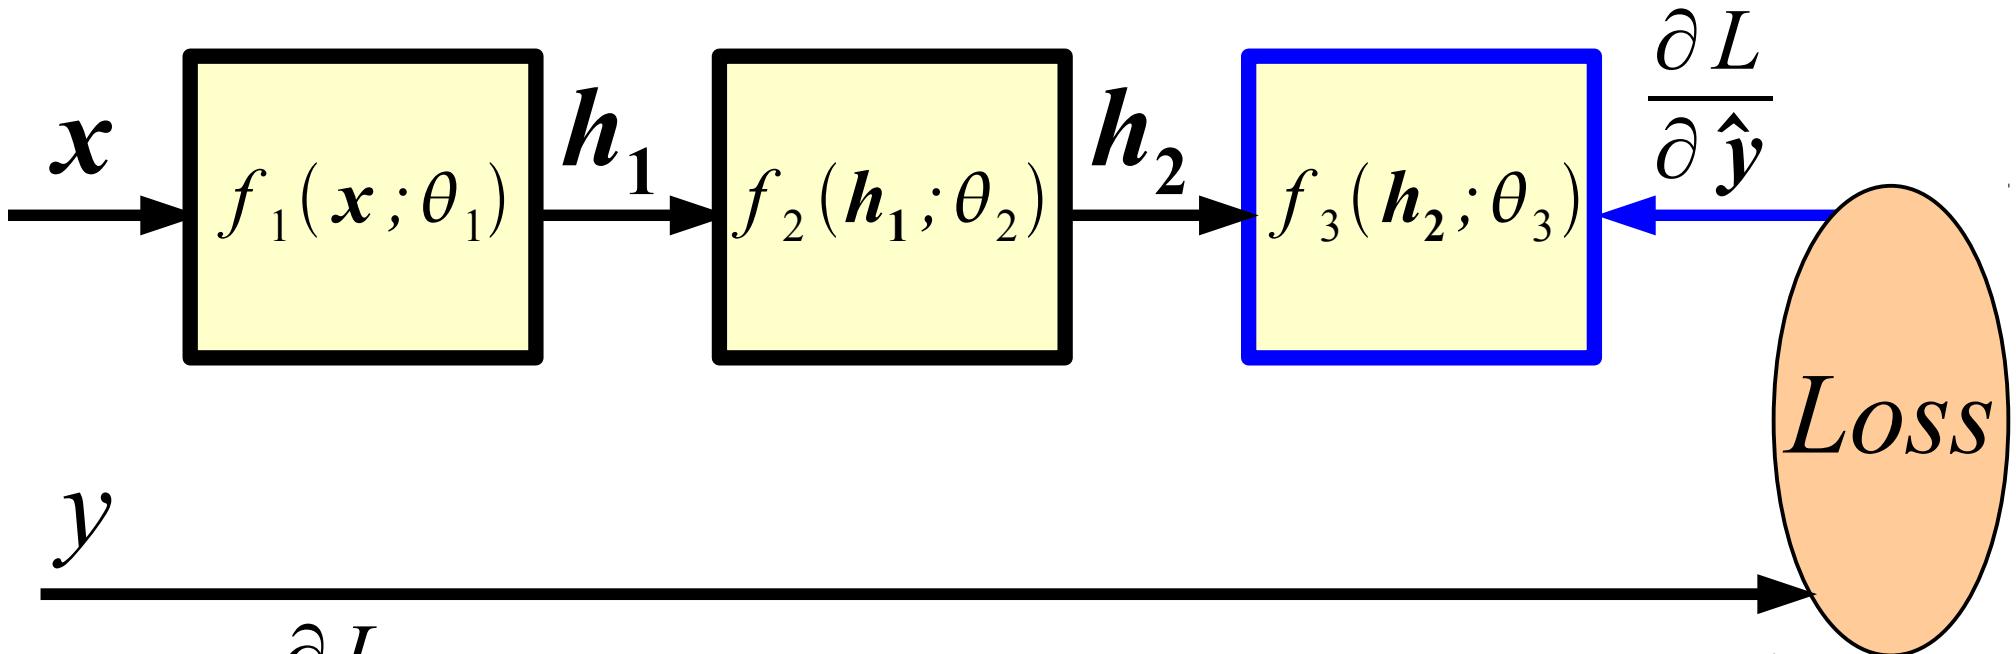


Given  $\frac{\partial L}{\partial \hat{y}}$  and assuming the Jacobian of each module is easy to compute, then we have:

$$\frac{\partial L}{\partial \theta_3} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial \theta_3}$$

$$\frac{\partial L}{\partial \mathbf{h}_2} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial \mathbf{h}_2}$$

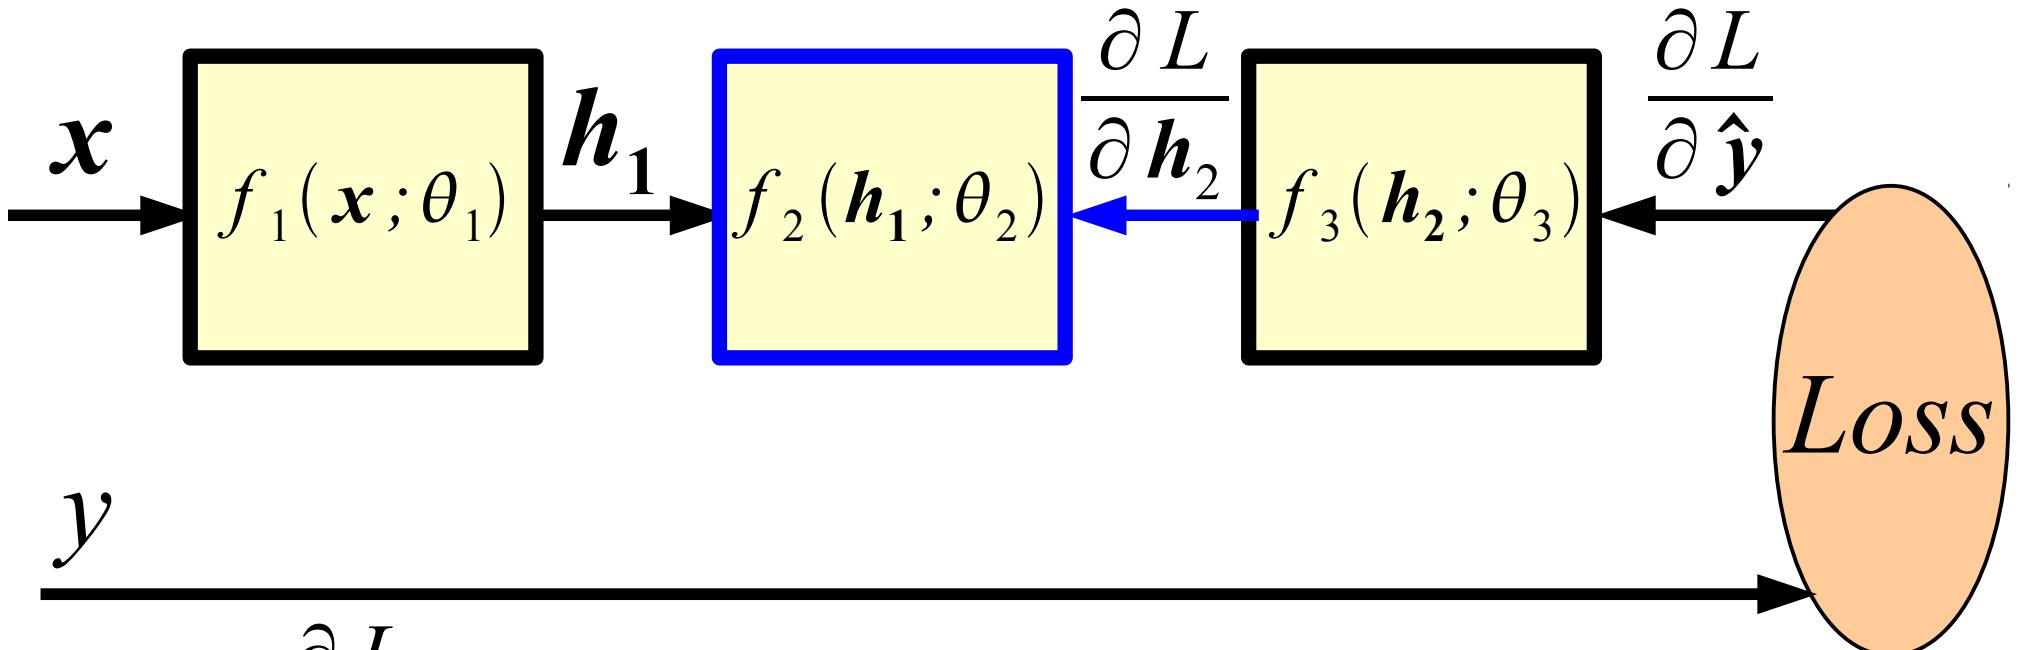
# Backward Propagation (BPROP)



Given  $\frac{\partial L}{\partial \hat{y}}$  and assuming the Jacobian of each module is easy to compute, then we have:

$$\frac{\partial L}{\partial \theta_3} = (\hat{y} - y) \ h_2' \quad \frac{\partial L}{\partial h_2} = (\hat{y} - y) \ \theta_3'$$

# Backward Propagation (BPROP)

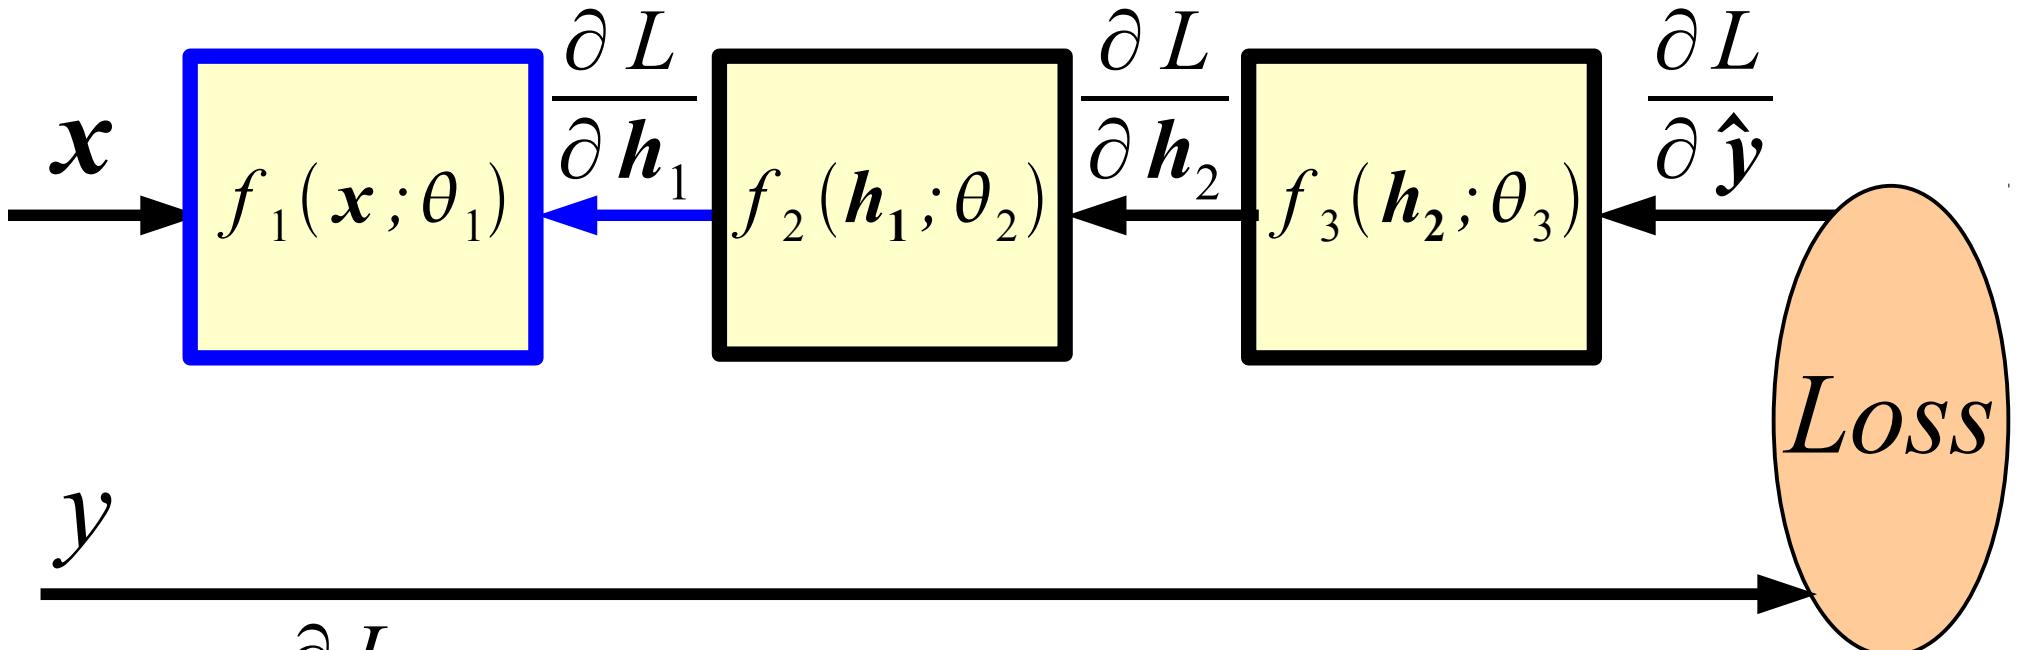


Given  $\frac{\partial L}{\partial h_2}$  we can compute now:

$$\frac{\partial L}{\partial \theta_2} = \frac{\partial L}{\partial h_2} \frac{\partial h_2}{\partial \theta_2}$$

$$\frac{\partial L}{\partial h_1} = \frac{\partial L}{\partial h_2} \frac{\partial h_2}{\partial h_1}$$

# Backward Propagation (BPROP)



Given  $\frac{\partial L}{\partial h_1}$  we can compute now:

$$\frac{\partial L}{\partial \theta_1} = \frac{\partial L}{\partial h_1} \frac{\partial h_1}{\partial \theta_1}$$

# Optimization

**Stochastic Gradient Descent (on mini-batches):**

$$\theta \leftarrow \theta - \eta \frac{\partial L}{\partial \theta}, \eta \in R$$

**Stochastic Gradient Descent with Momentum:**

$$\begin{aligned}\theta &\leftarrow \theta - \eta \Delta \\ \Delta &\leftarrow 0.9 \Delta + \frac{\partial L}{\partial \theta}\end{aligned}$$

LeCun et al. “Efficient BackProp” Neural Networks: Tricks of the trade 1998

Schaul et al. “No more pesky learning rates” ICML 2013

Sutskever et al. “On the importance of initialization and momentum...” ICML 2013

# Toy Code: Neural Net Trainer

```
% F-PROP
for i = 1 : nr_layers - 1
    [h{i} jac{i}] = nonlinearity(W{i} * h{i-1} + b{i});
end
h{nr_layers-1} = W{nr_layers-1} * h{nr_layers-2} + b{nr_layers-1};
prediction = softmax(h{l-1});

% CROSS ENTROPY LOSS
loss = - sum(sum(log(prediction) .* target)) / batch_size;

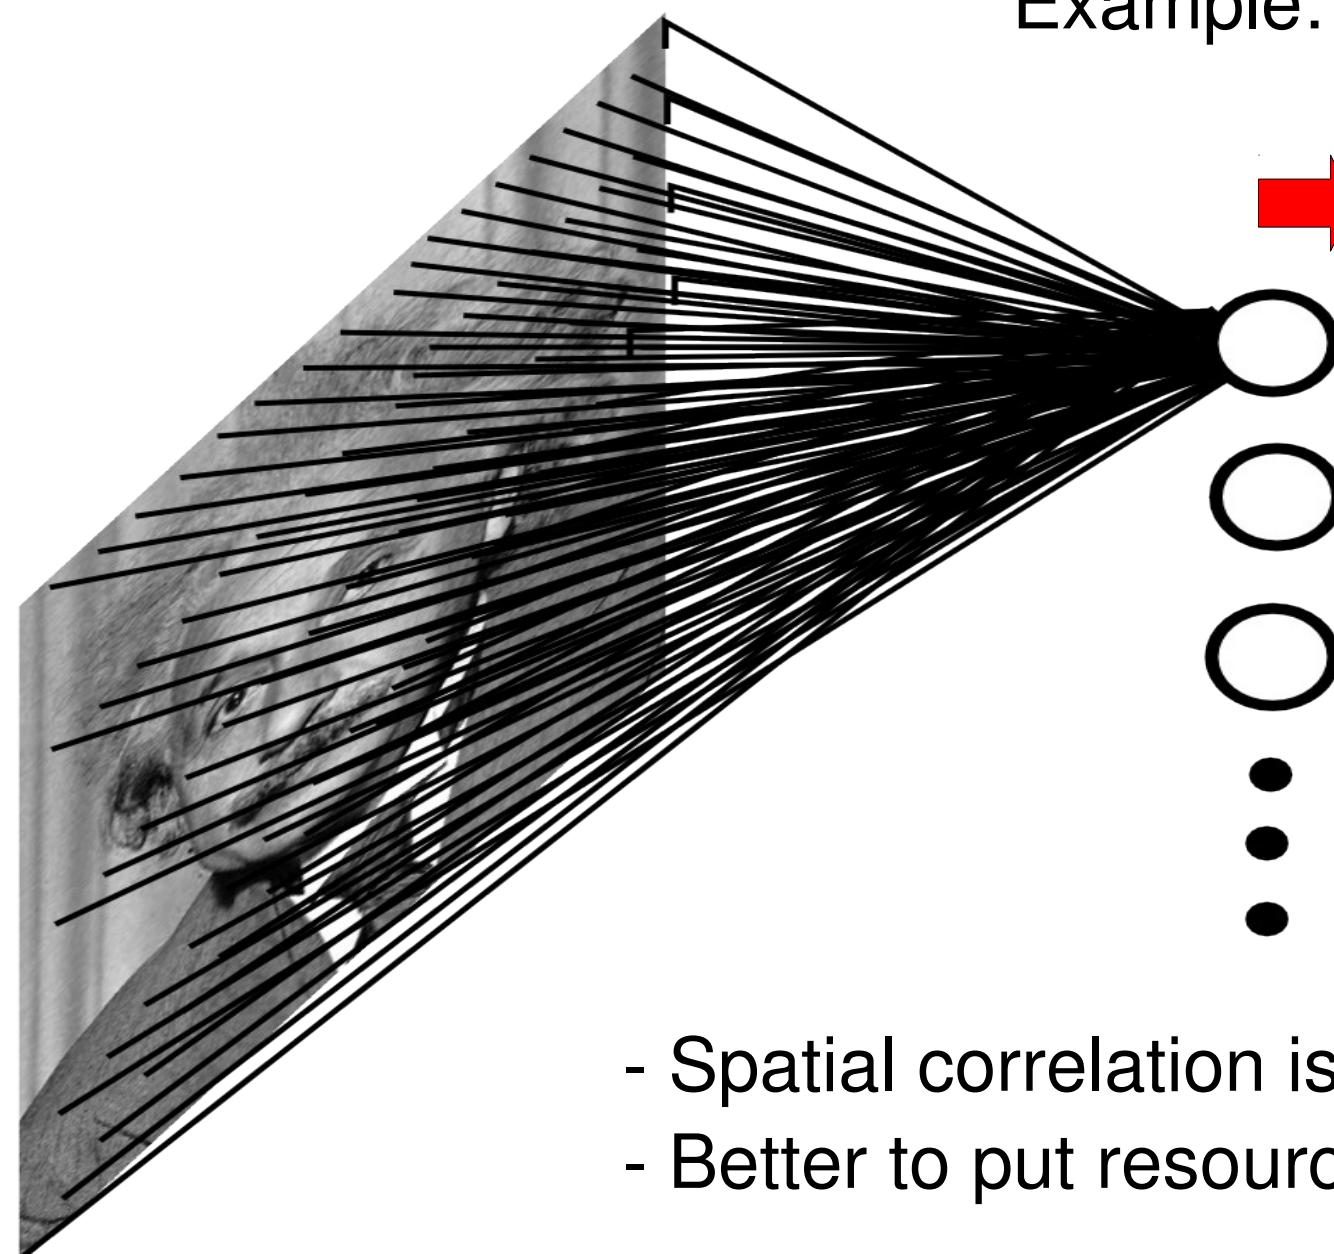
% B-PROP
dh{l-1} = prediction - target;
for i = nr_layers - 1 : -1 : 1
    Wgrad{i} = dh{i} * h{i-1}';
    bgrad{i} = sum(dh{i}, 2);
    dh{i-1} = (W{i}' * dh{i}) .* jac{i-1};
end

% UPDATE
for i = 1 : nr_layers - 1
    W{i} = W{i} - (lr / batch_size) * Wgrad{i};
    b{i} = b{i} - (lr / batch_size) * bgrad{i};
end
```

# KEY IDEAS: Training NNets

- Neural Net = stack of feature detectors
- F-Prop / B-Prop
- Learning by SGD

# FULLY CONNECTED NEURAL NET



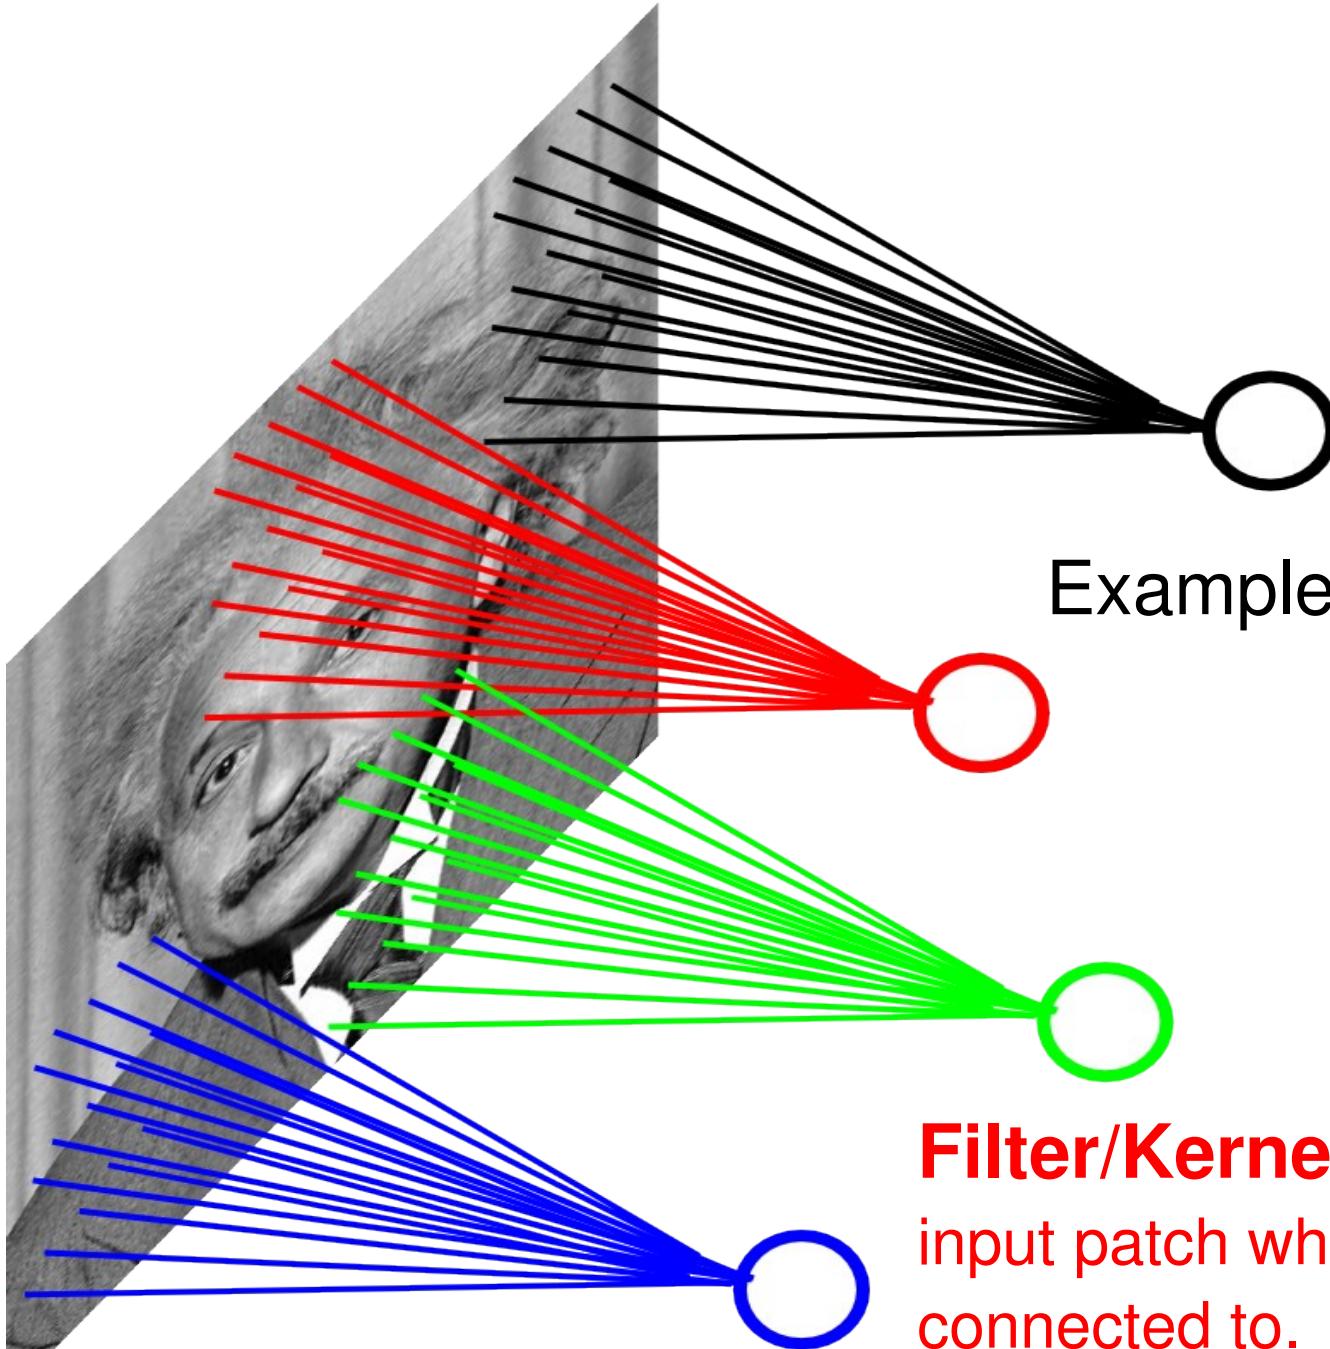
Example: 1000x1000 image

1M hidden units

**10<sup>12</sup> parameters!!!**

- Spatial correlation is local
- Better to put resources elsewhere!

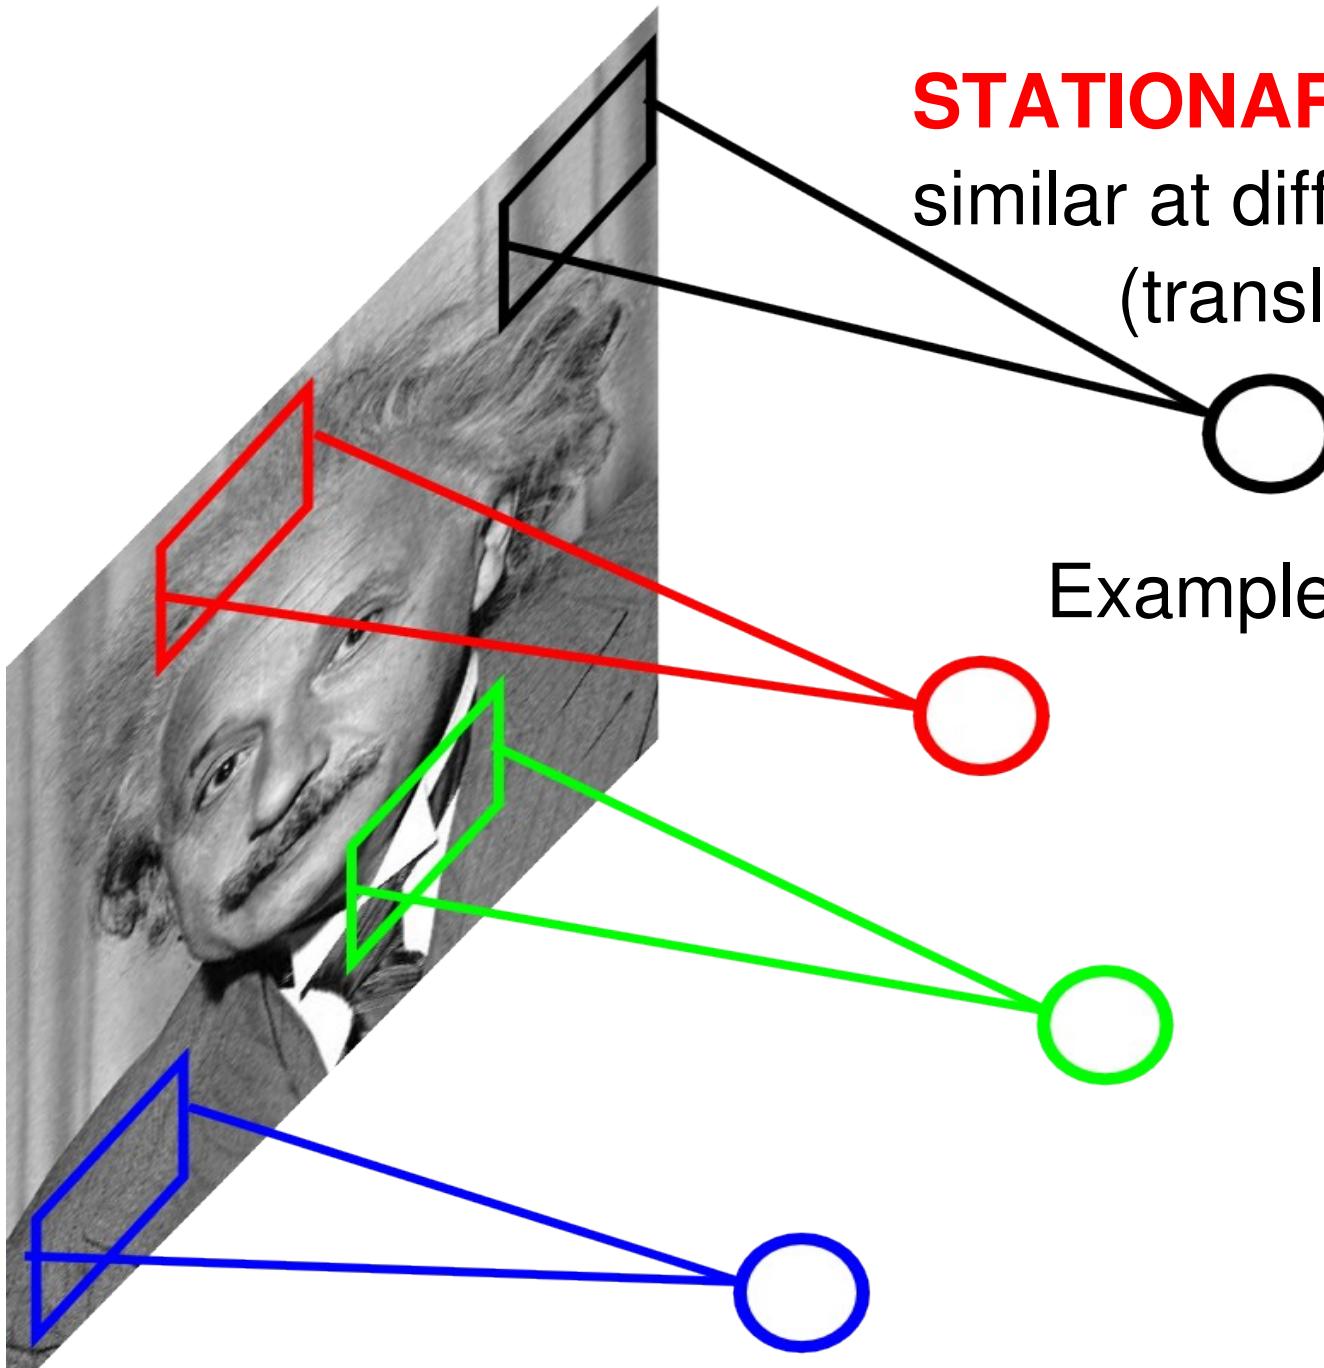
# LOCALLY CONNECTED NEURAL NET



Example: 1000x1000 image  
1M hidden units  
Filter size: 10x10  
100M parameters

**Filter/Kernel/Receptive field:**  
input patch which the hidden unit is  
connected to.

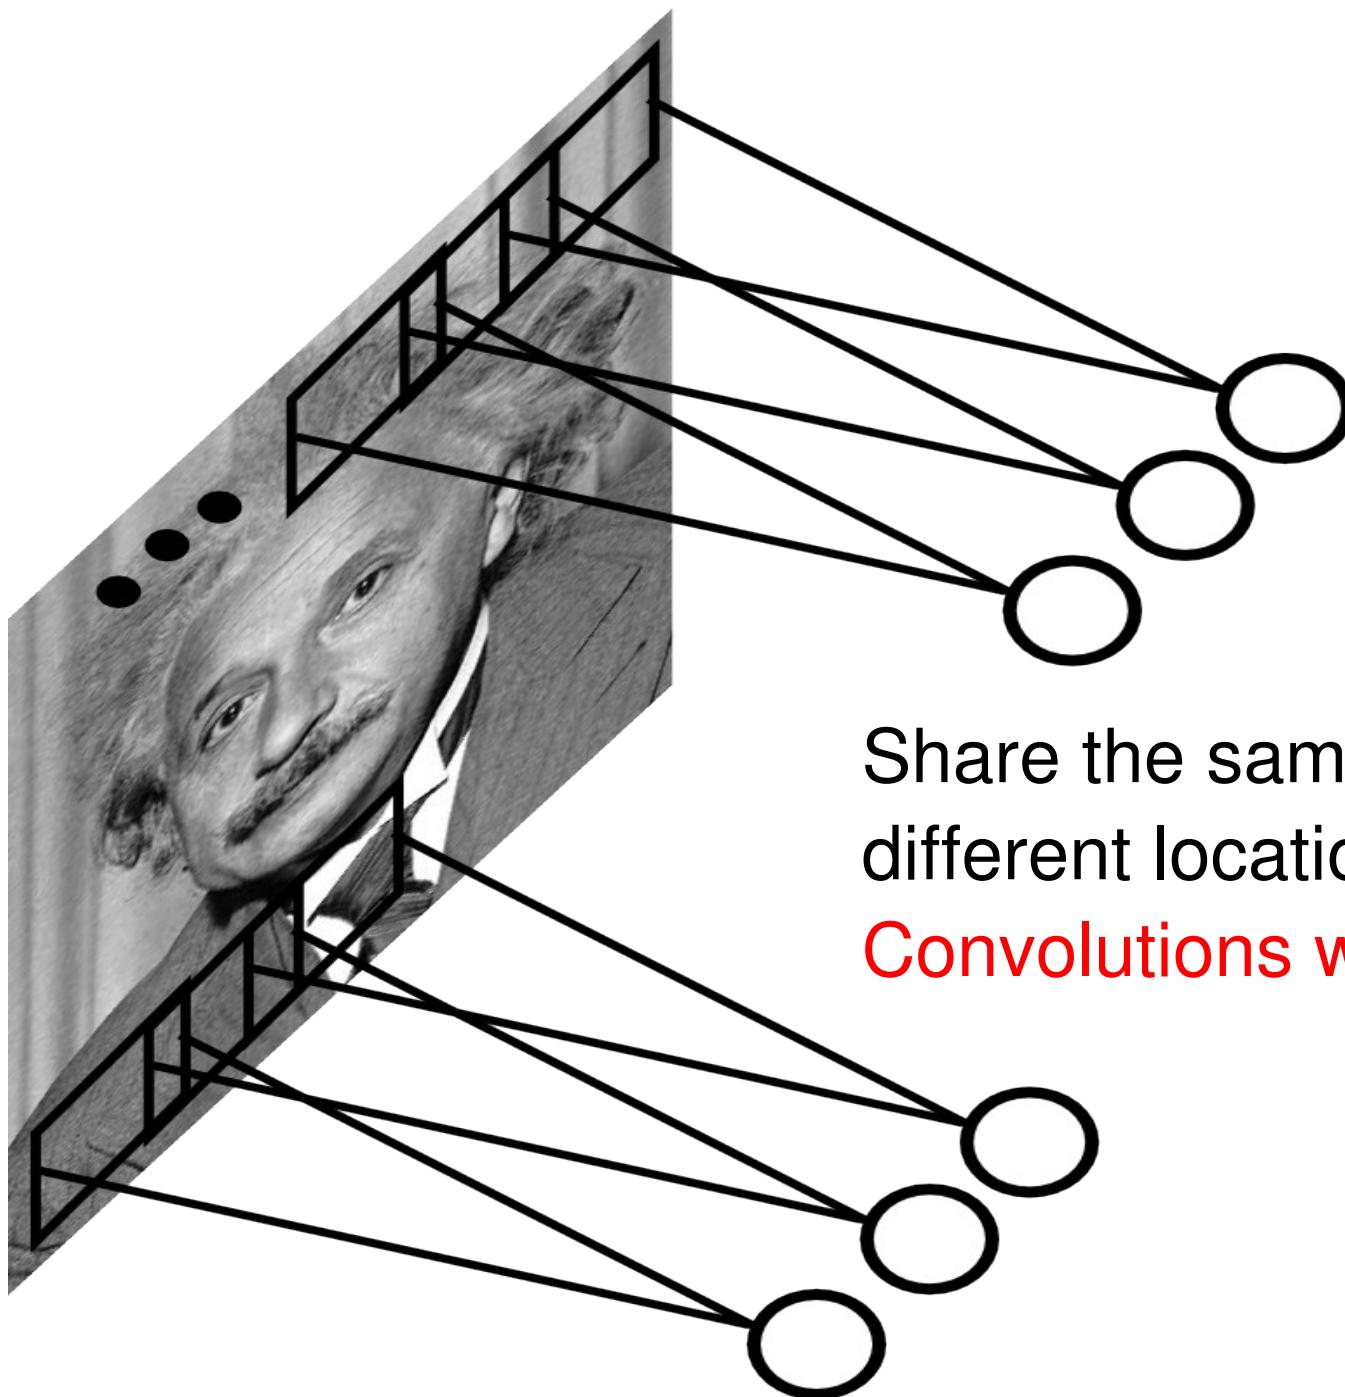
# LOCALLY CONNECTED NEURAL NET



**STATIONARITY?** Statistics are similar at different locations  
(translation invariance)

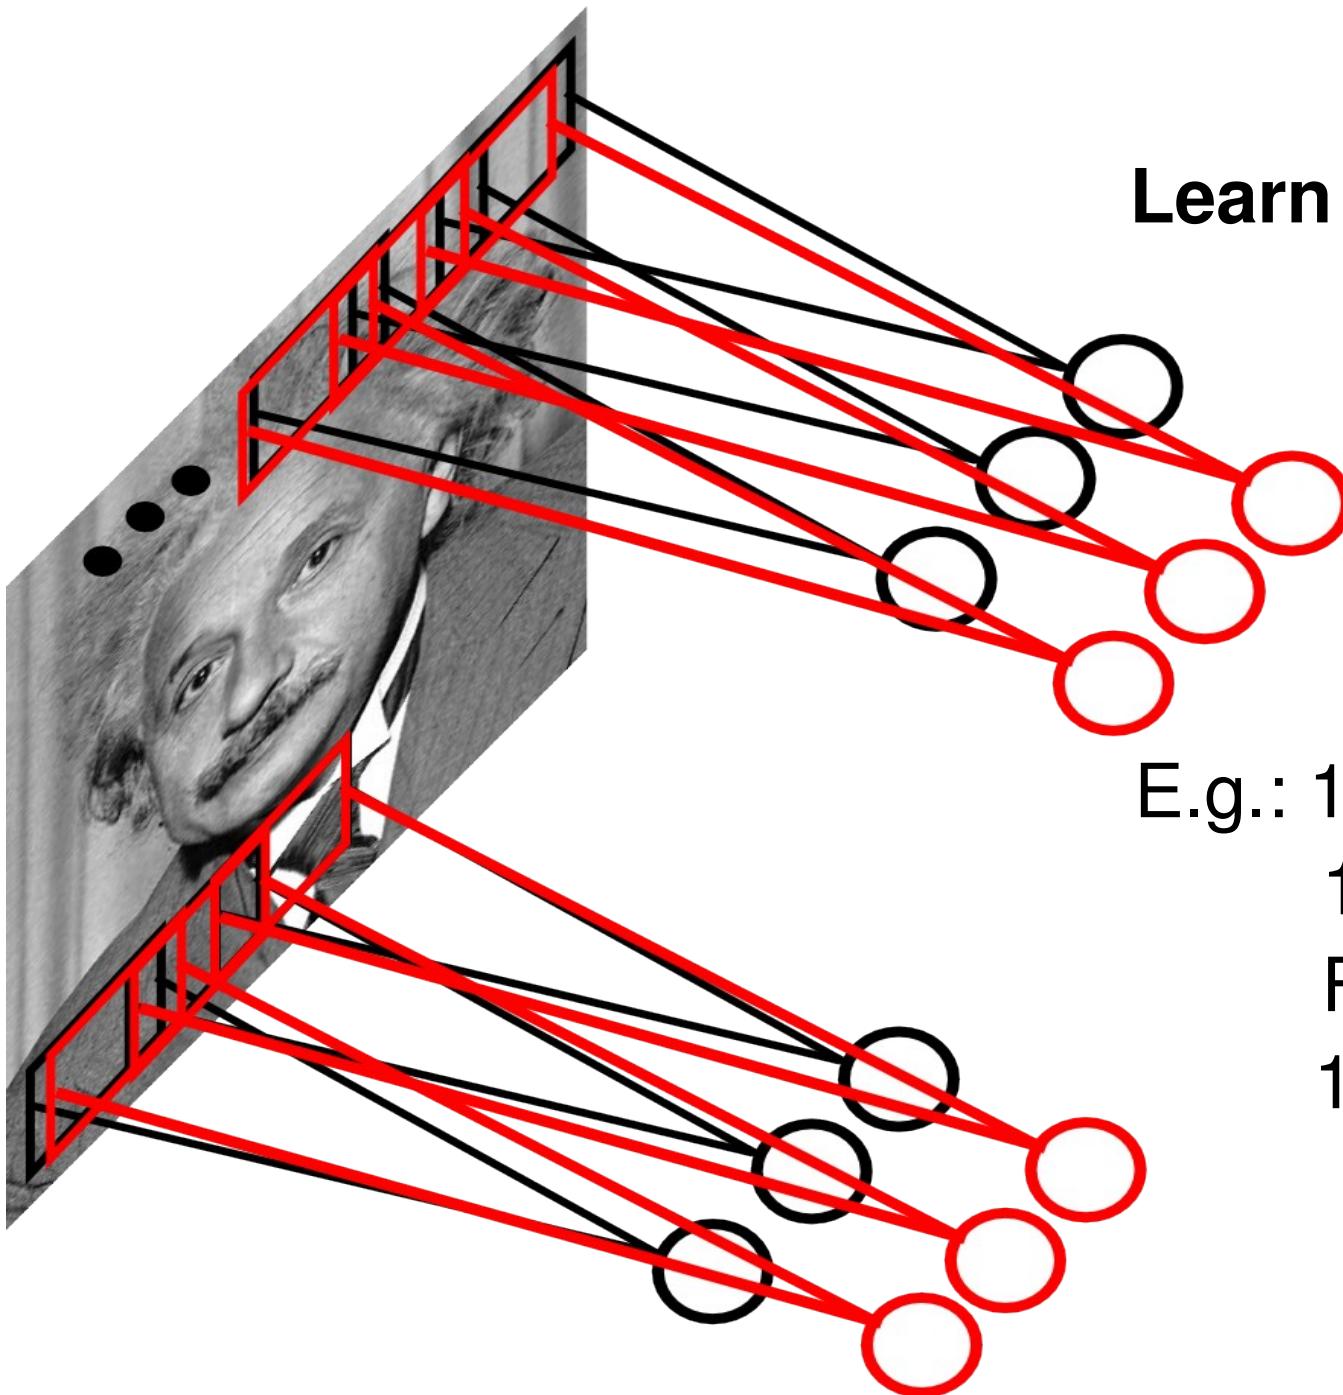
Example: 1000x1000 image  
1M hidden units  
Filter size: 10x10  
100M parameters

# CONVOLUTIONAL NET



Share the same parameters across  
different locations:  
**Convolutions with learned kernels**

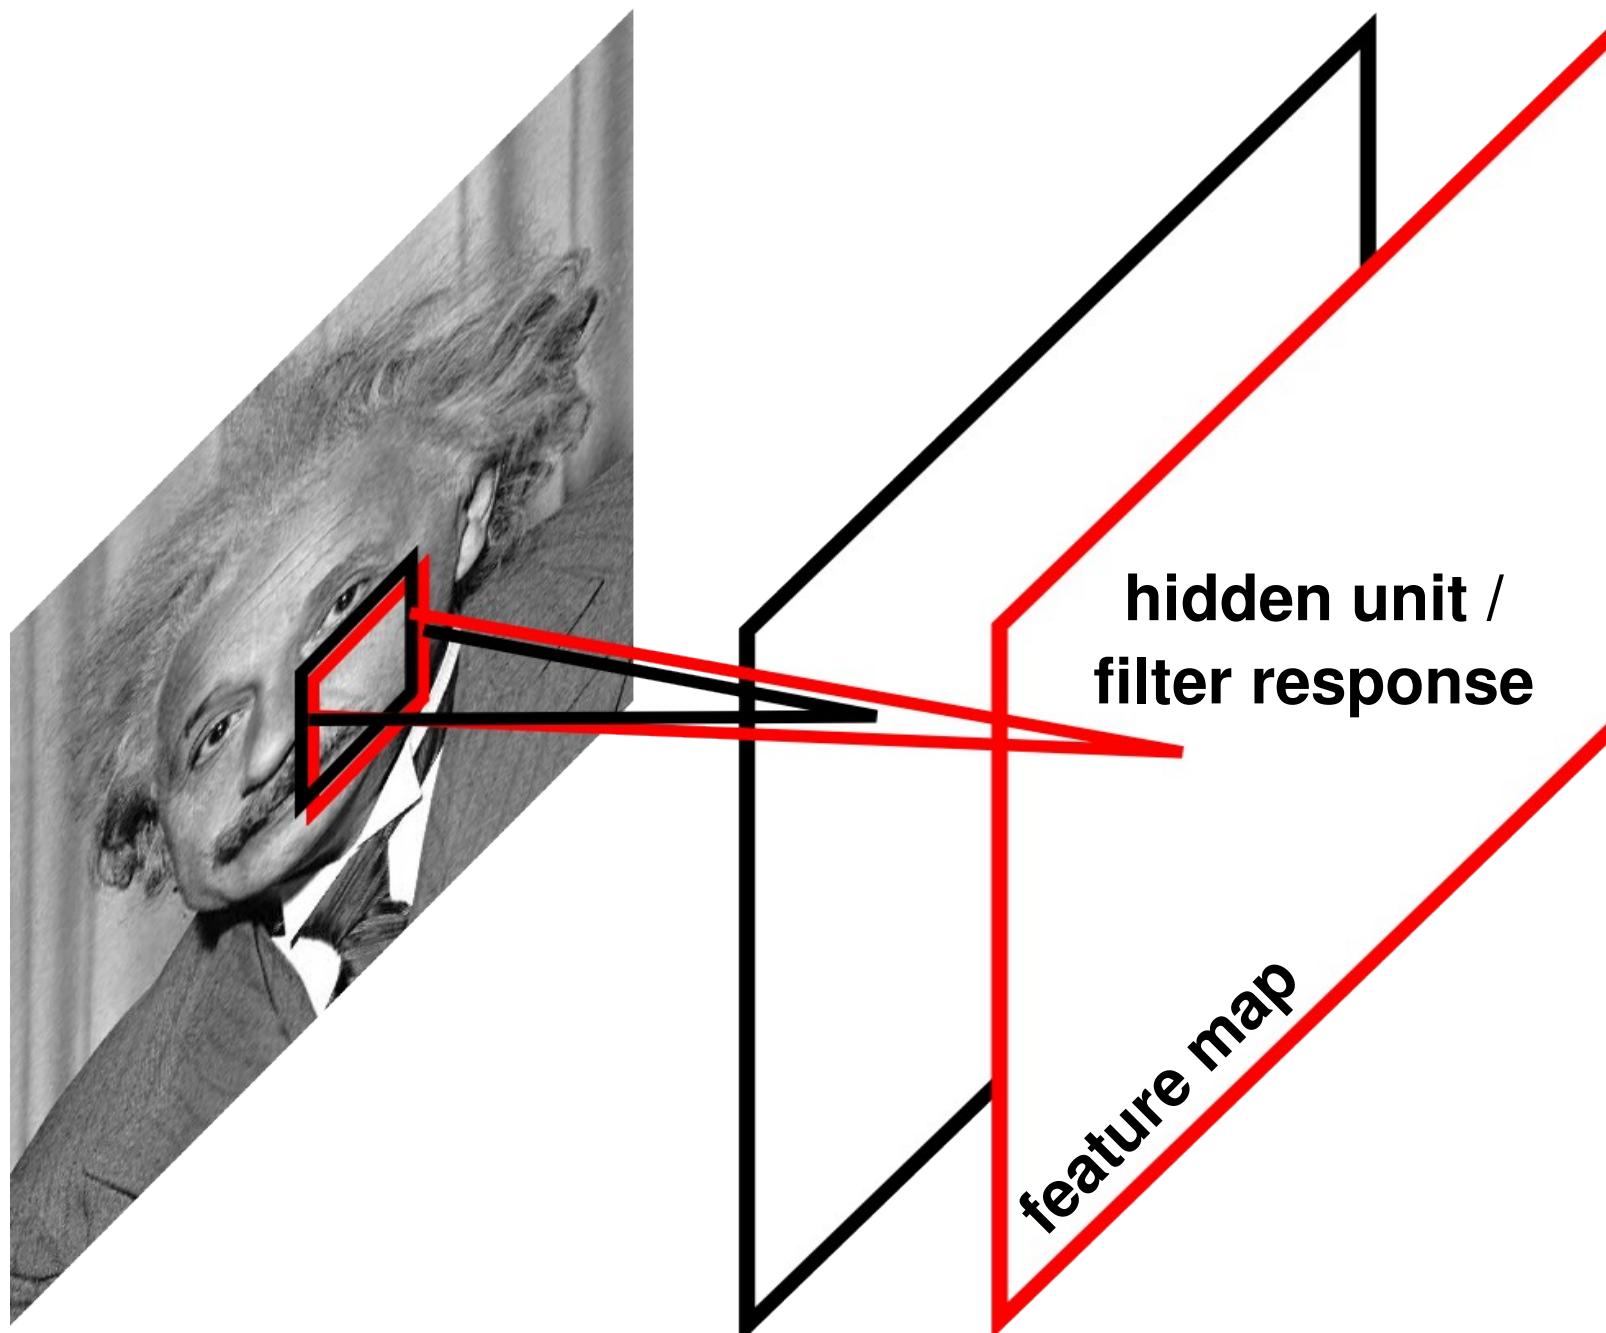
# CONVOLUTIONAL NET



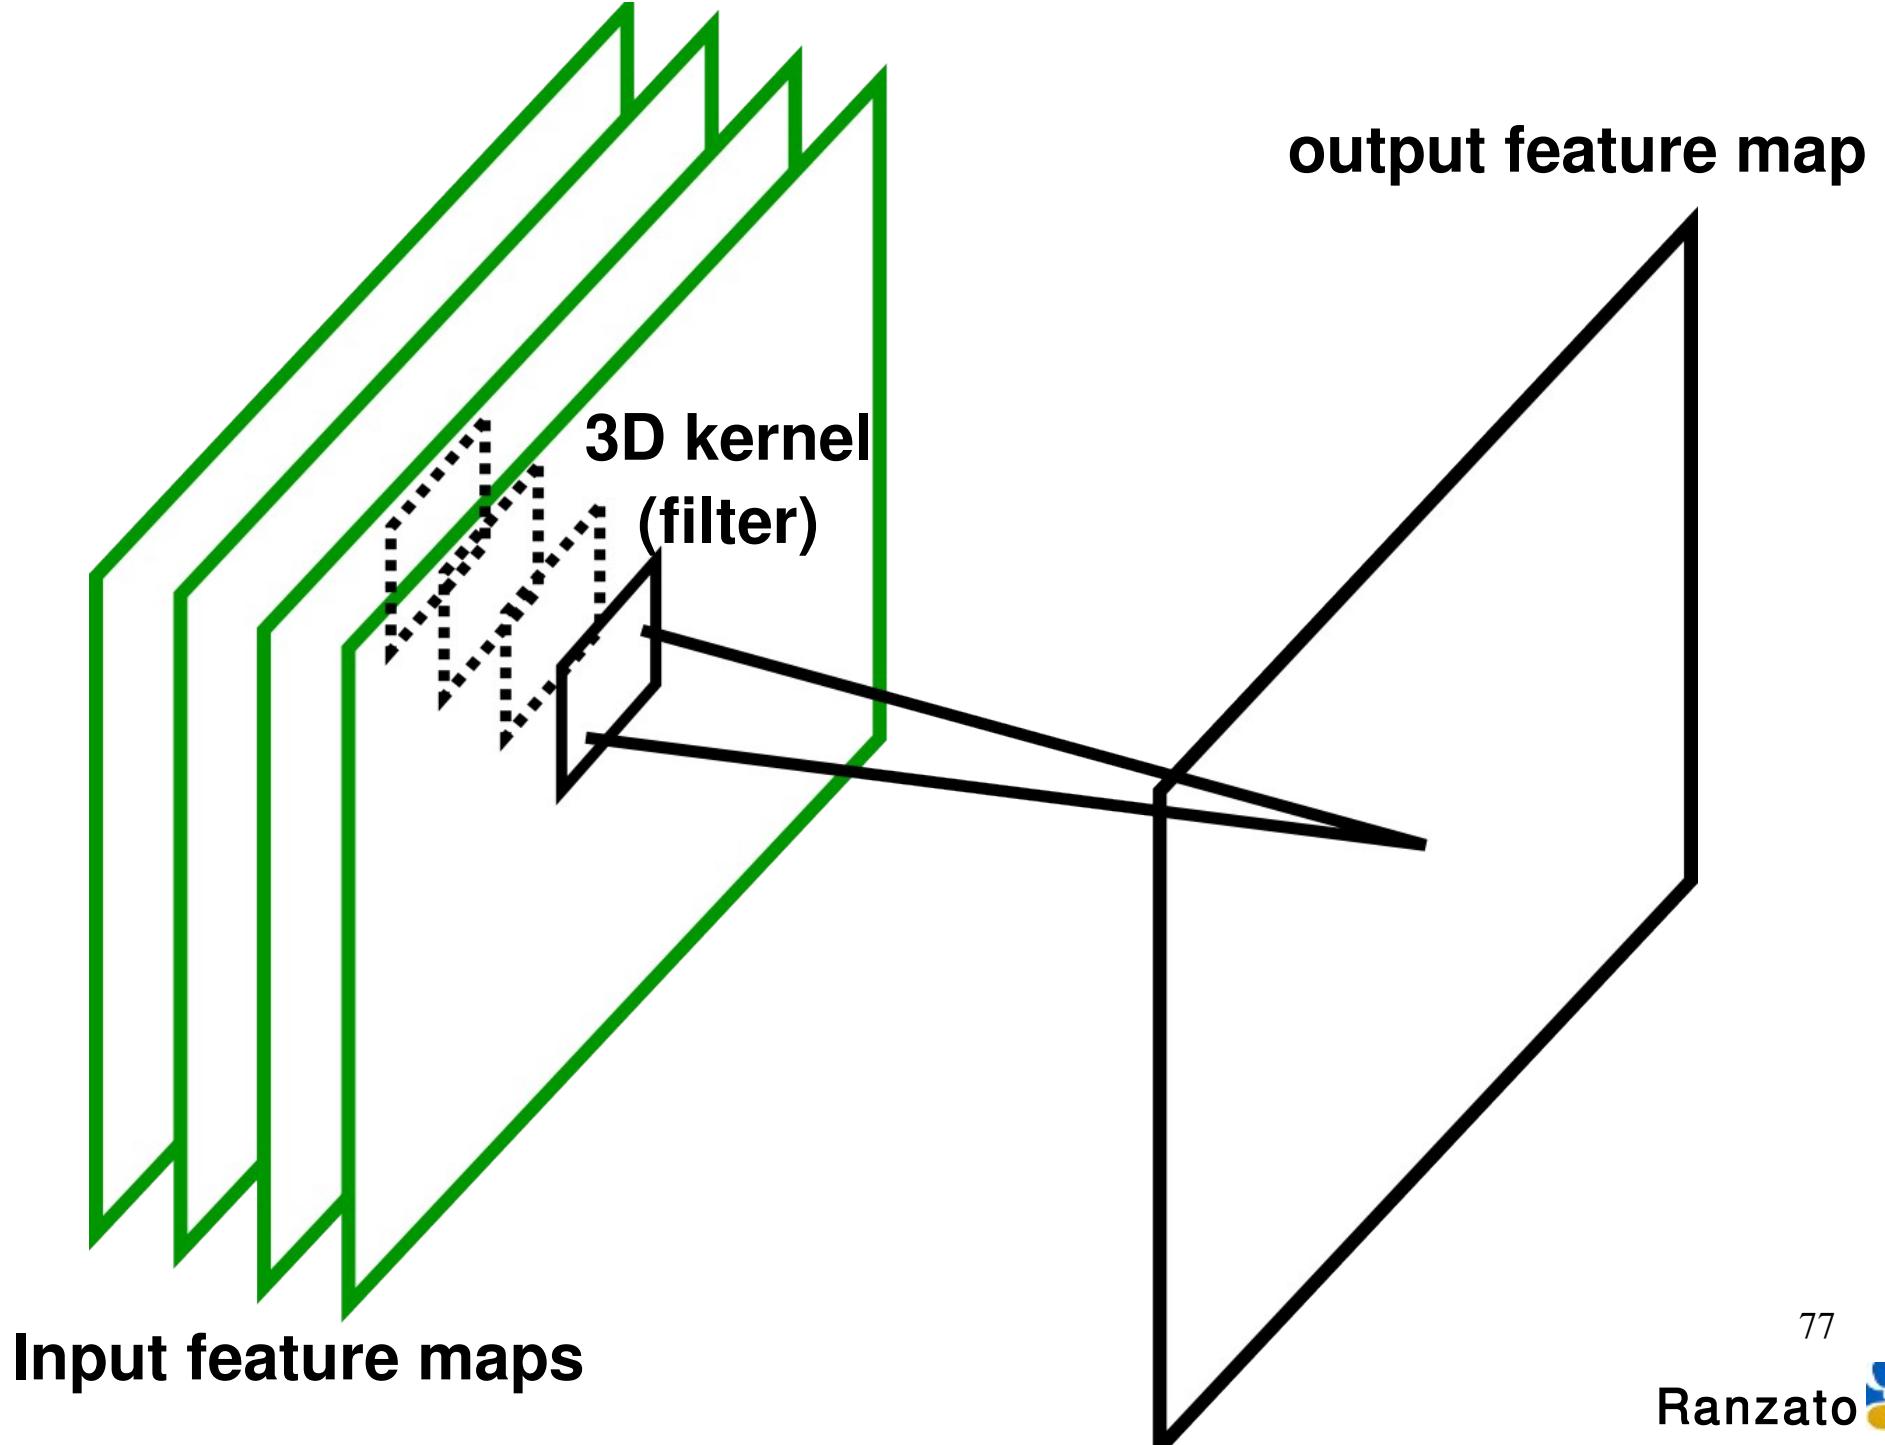
**Learn multiple filters.**

E.g.: 1000x1000 image  
100 Filters  
Filter size: 10x10  
10K parameters

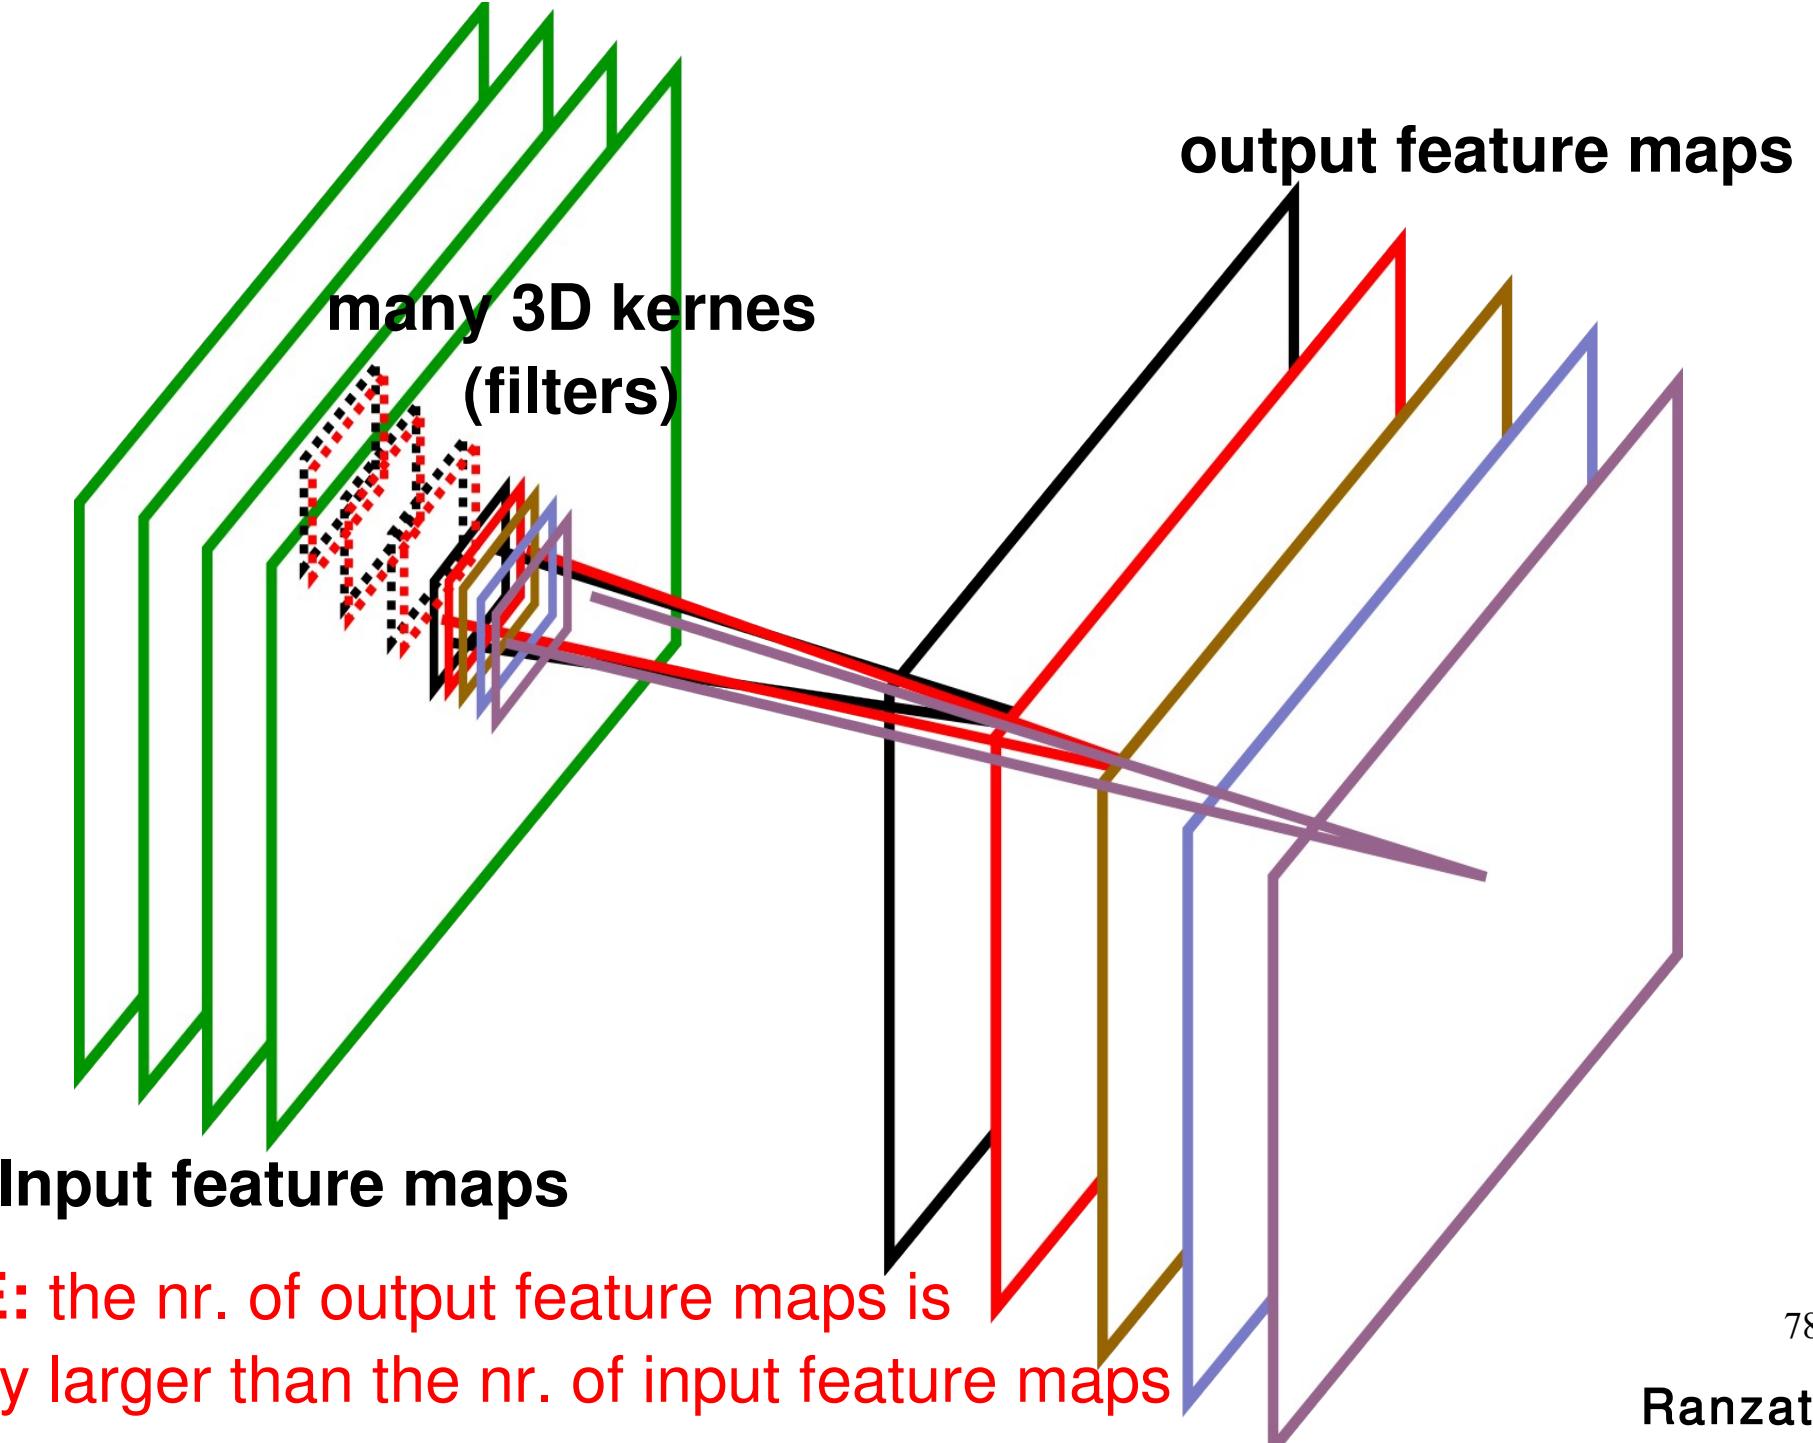
# CONVOLUTIONAL NET



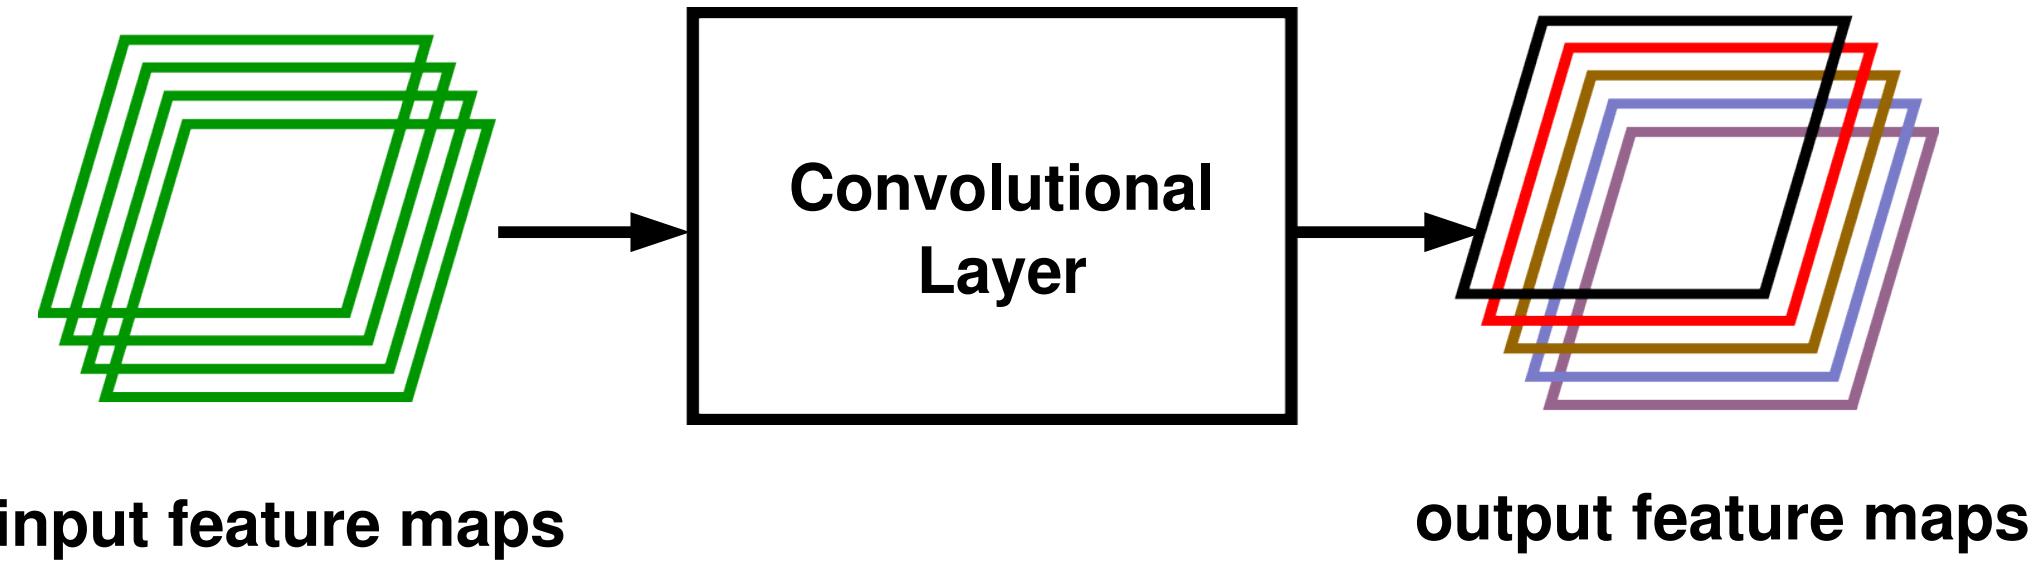
# CONVOLUTIONAL LAYER



# CONVOLUTIONAL LAYER



# CONVOLUTIONAL LAYER



**NOTE:** the nr. of output feature maps is  
usually larger than the nr. of input feature maps

# KEY IDEAS: CONV. NETS

A standard neural net applied to images:

- scales quadratically with the size of the input
- does not leverage stationarity

Solution:

- connect each hidden unit to a small patch of the input
- share the weight across hidden units

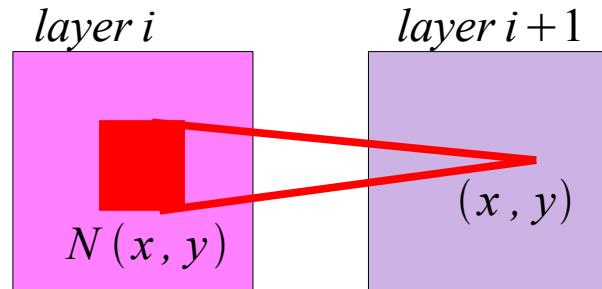
This is called: **convolutional network**.

LeCun et al. “Gradient-based learning applied to document recognition” IEEE 1998

# SPECIAL LAYERS

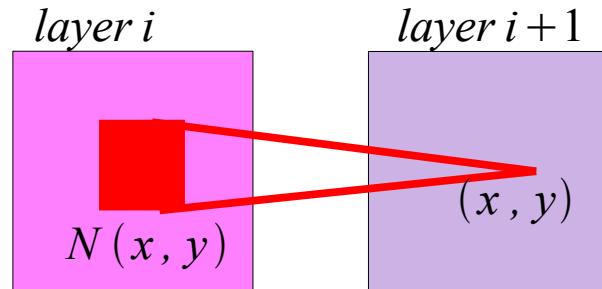
Over the years, some new modules have proven to be very effective when plugged into conv-nets:

- **Pooling (average, L2, max)**



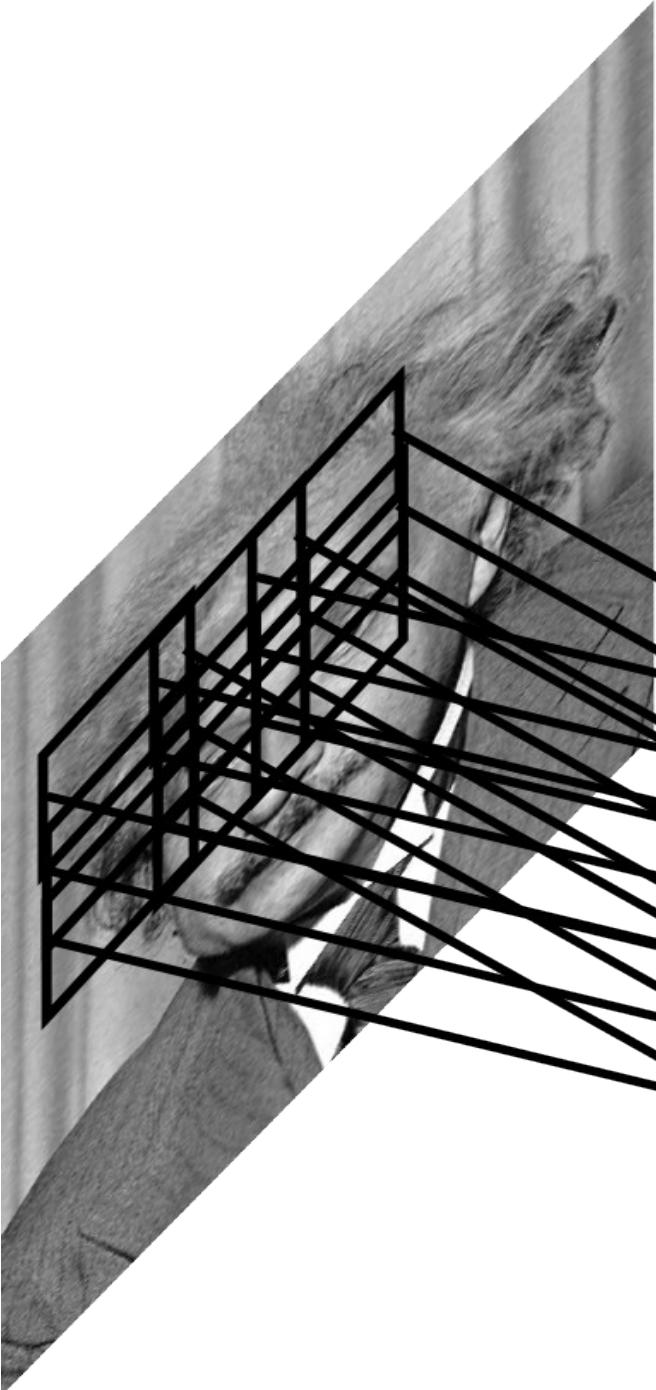
$$h_{i+1, x, y} = \max_{(j, k) \in N(x, y)} h_{i, j, k}$$

- **Local Contrast Normalization (over space / features)**



$$h_{i+1, x, y} = \frac{h_{i, x, y} - m_{i, x, y}}{\sigma_{i, x, y}}$$

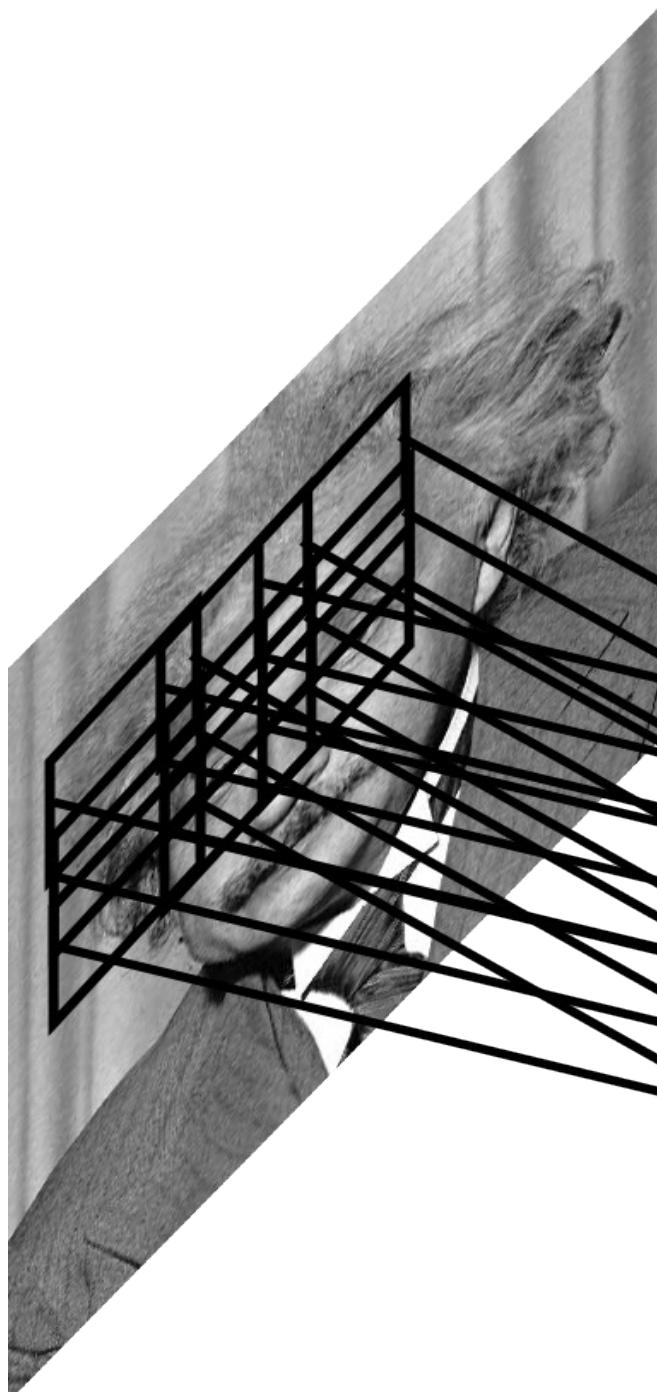
# POOLING



Let us assume filter is an “eye” detector.

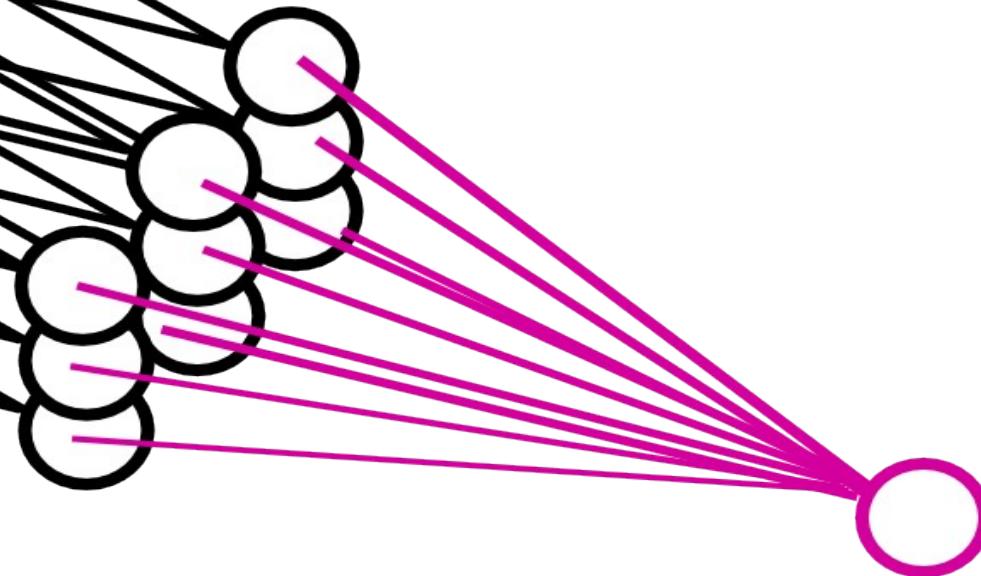
**Q.:** how can we make the detection robust to the exact location of the eye?

# POOLING

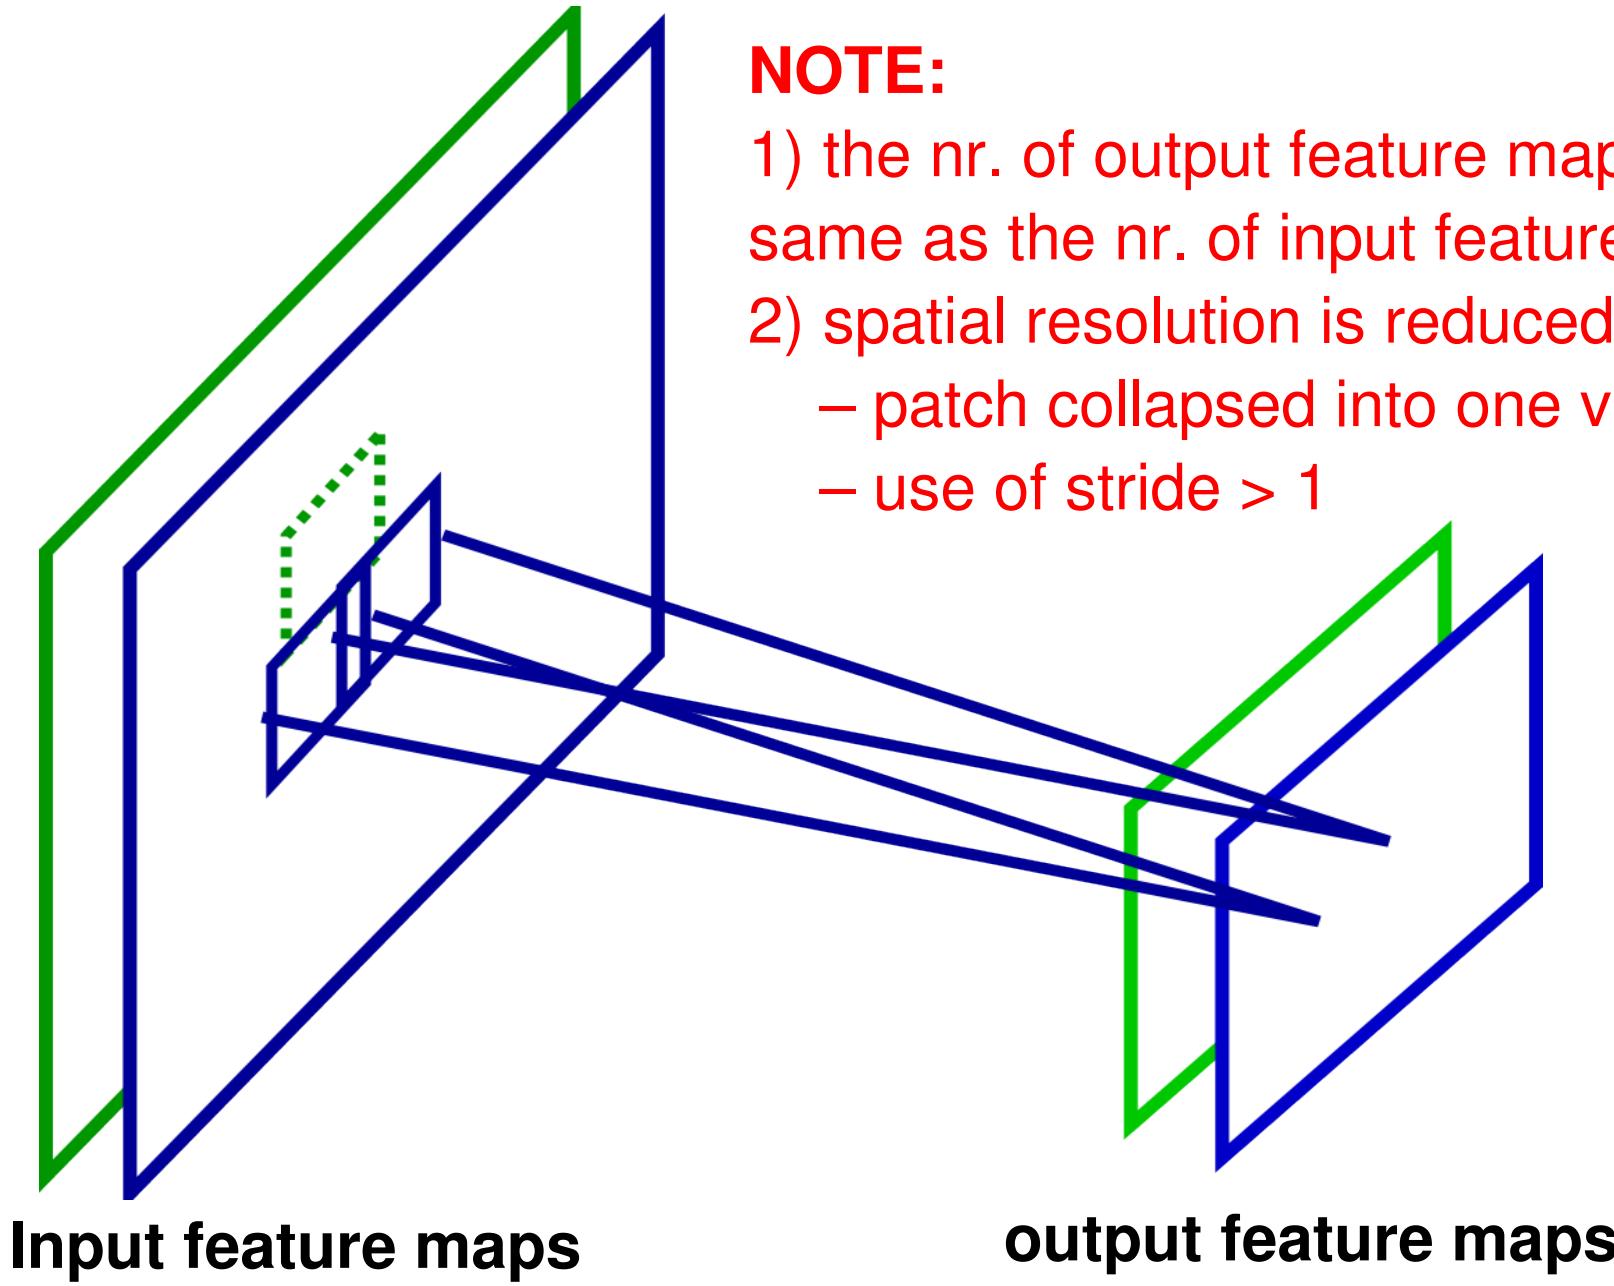


By “**pooling**” (e.g., taking max) filter responses at different locations we gain robustness to the exact spatial location of features.

$$h_{i+1, x, y} = \max_{(j, k) \in N(x, y)} h_{i, j, k}$$



# POOLING LAYER



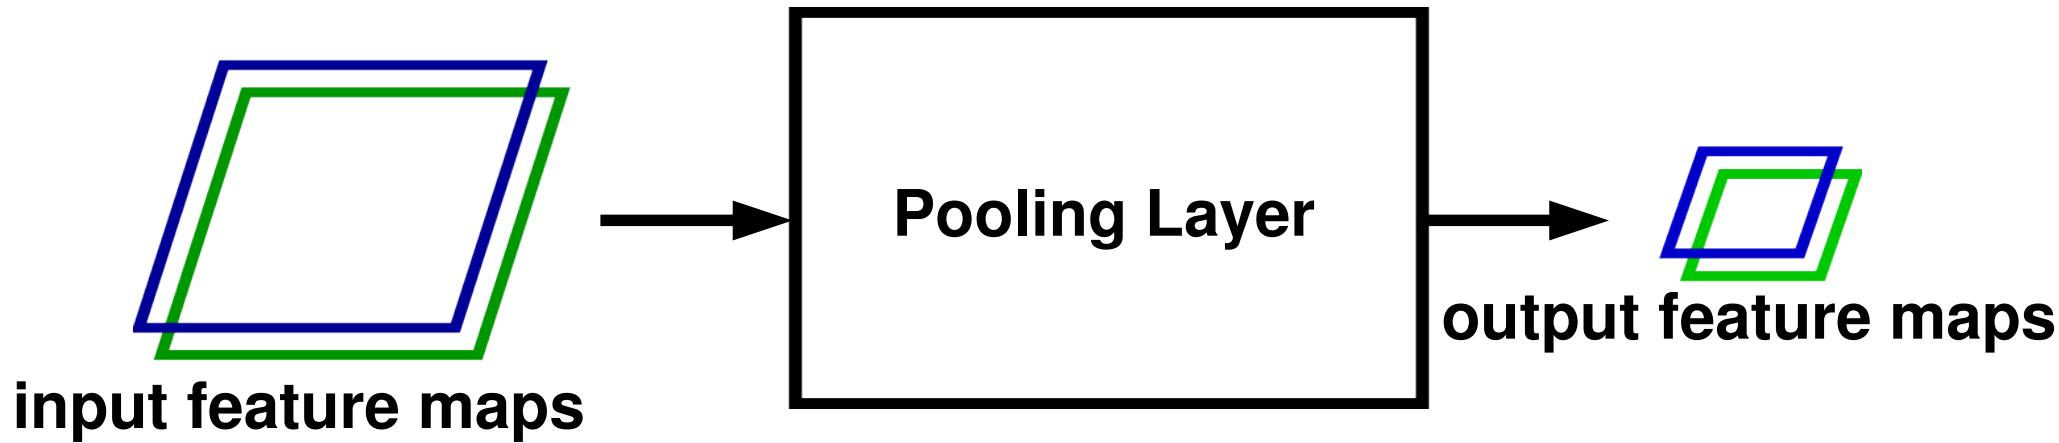
## NOTE:

- 1) the nr. of output feature maps is the same as the nr. of input feature maps
- 2) spatial resolution is reduced
  - patch collapsed into one value
  - use of stride > 1

# POOLING LAYER

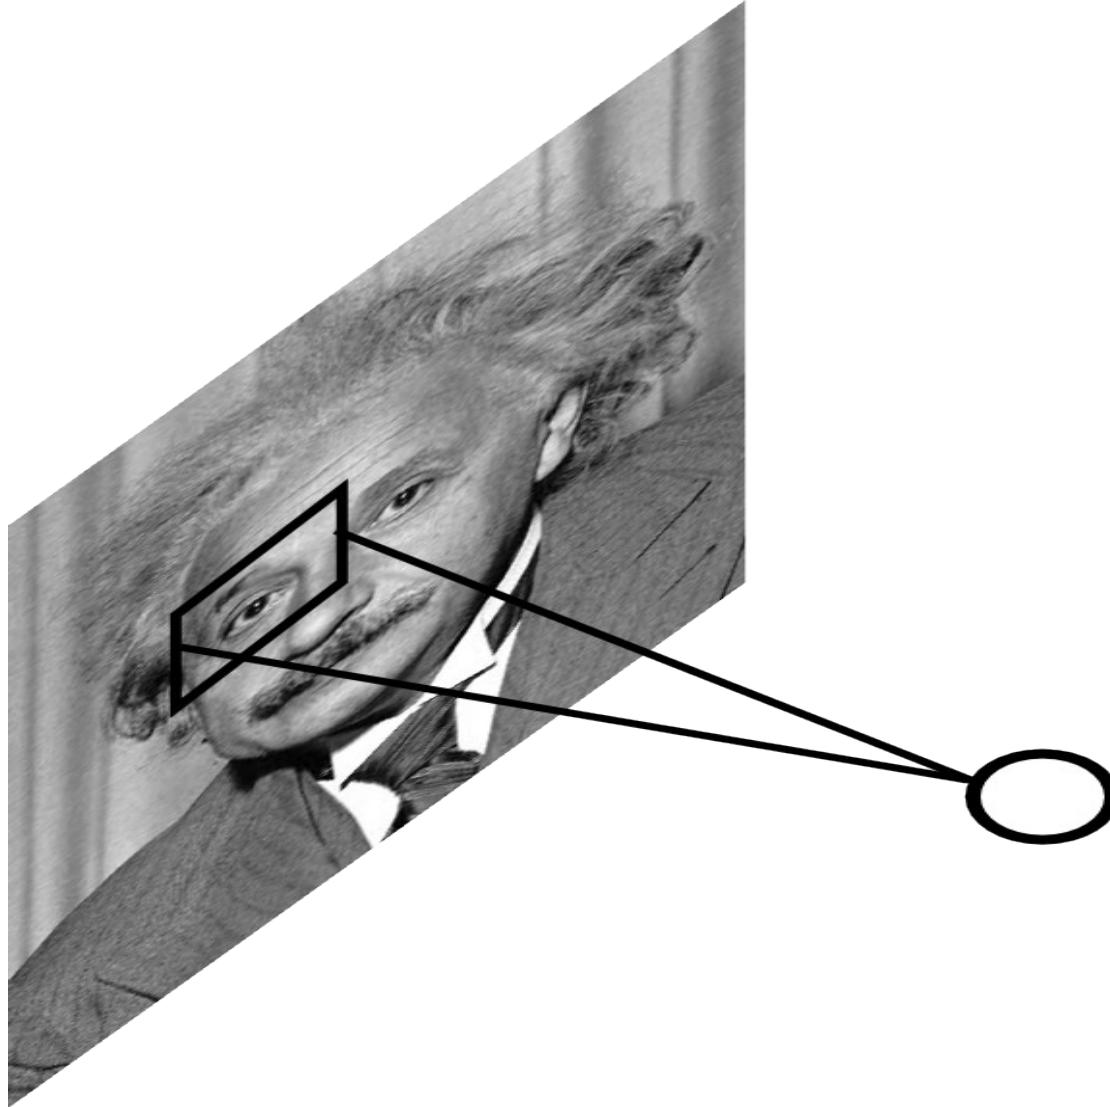
## NOTE:

- 1) the nr. of output feature maps is the same as the nr. of input feature maps
- 2) spatial resolution is reduced
  - patch collapsed into one value
  - use of stride > 1



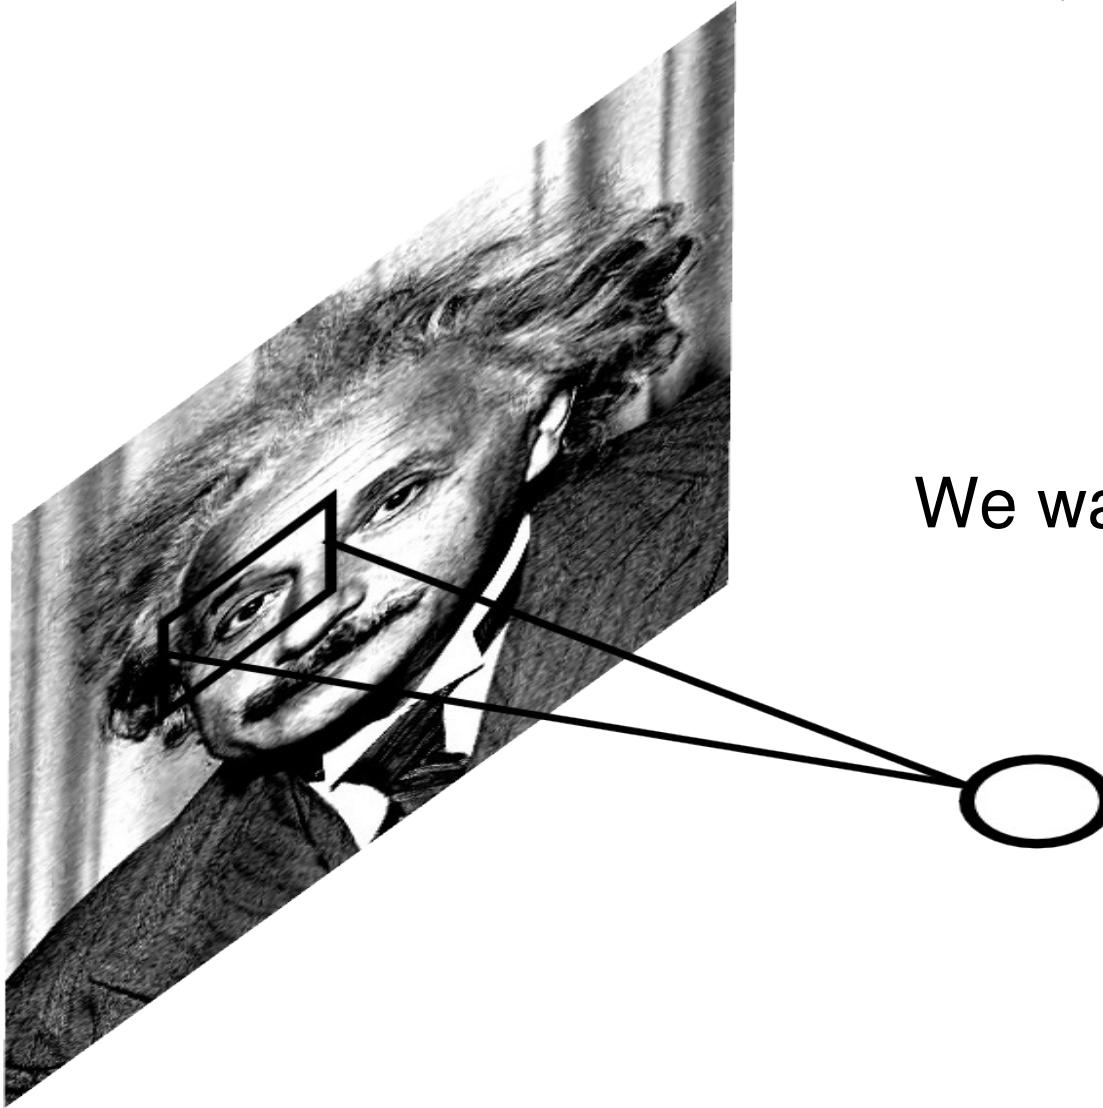
# LOCAL CONTRAST NORMALIZATION

$$h_{i+1,x,y} = \frac{h_{i,x,y} - m_{i,N}(x,y)}{\sigma_{i,N}(x,y)}$$



# LOCAL CONTRAST NORMALIZATION

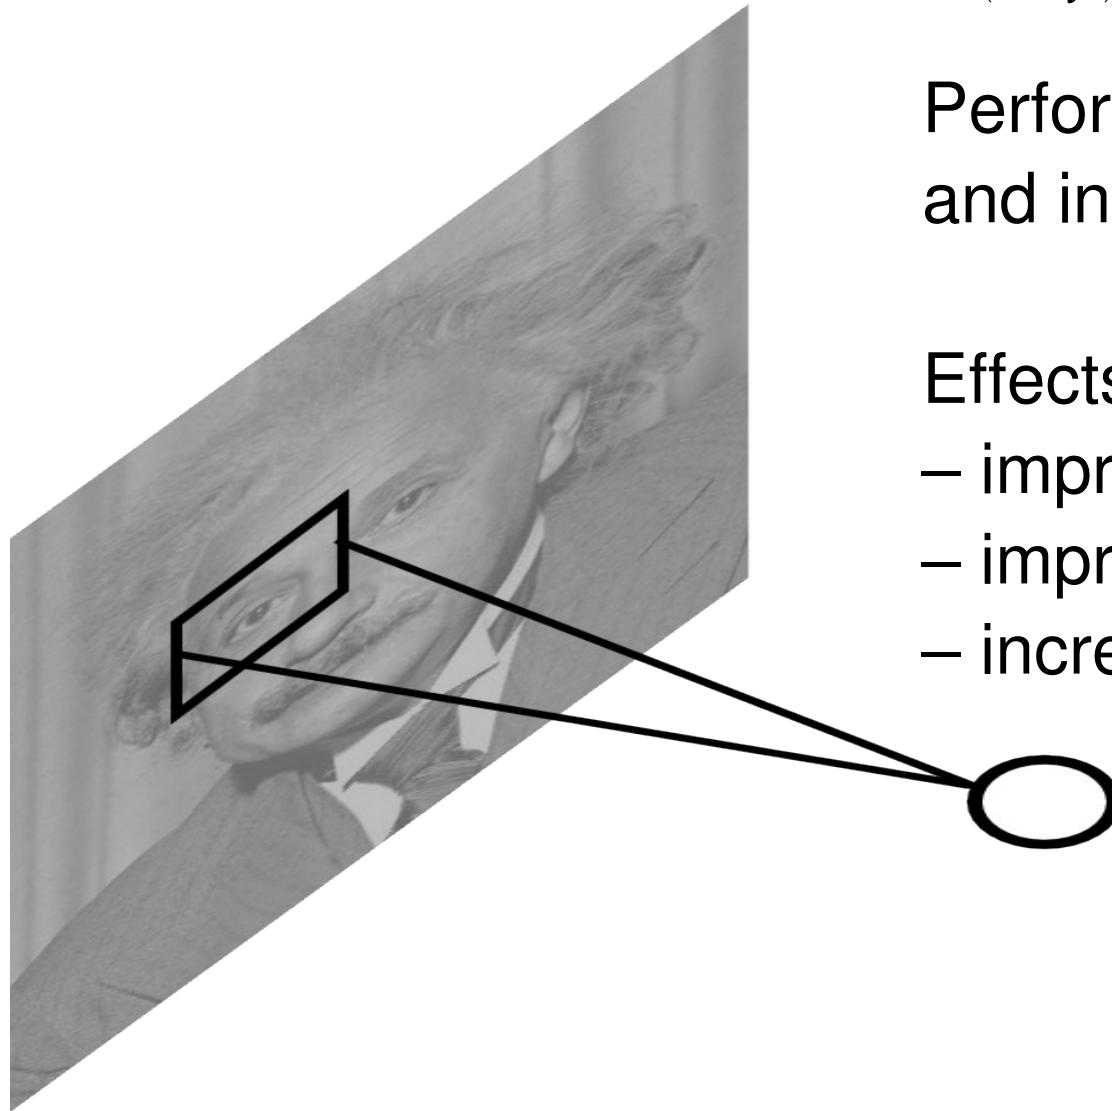
$$h_{i+1,x,y} = \frac{h_{i,x,y} - m_{i,N}(x,y)}{\sigma_{i,N}(x,y)}$$



We want the same response.

# LOCAL CONTRAST NORMALIZATION

$$h_{i+1,x,y} = \frac{h_{i,x,y} - m_{i,N}(x,y)}{\sigma_{i,N}(x,y)}$$



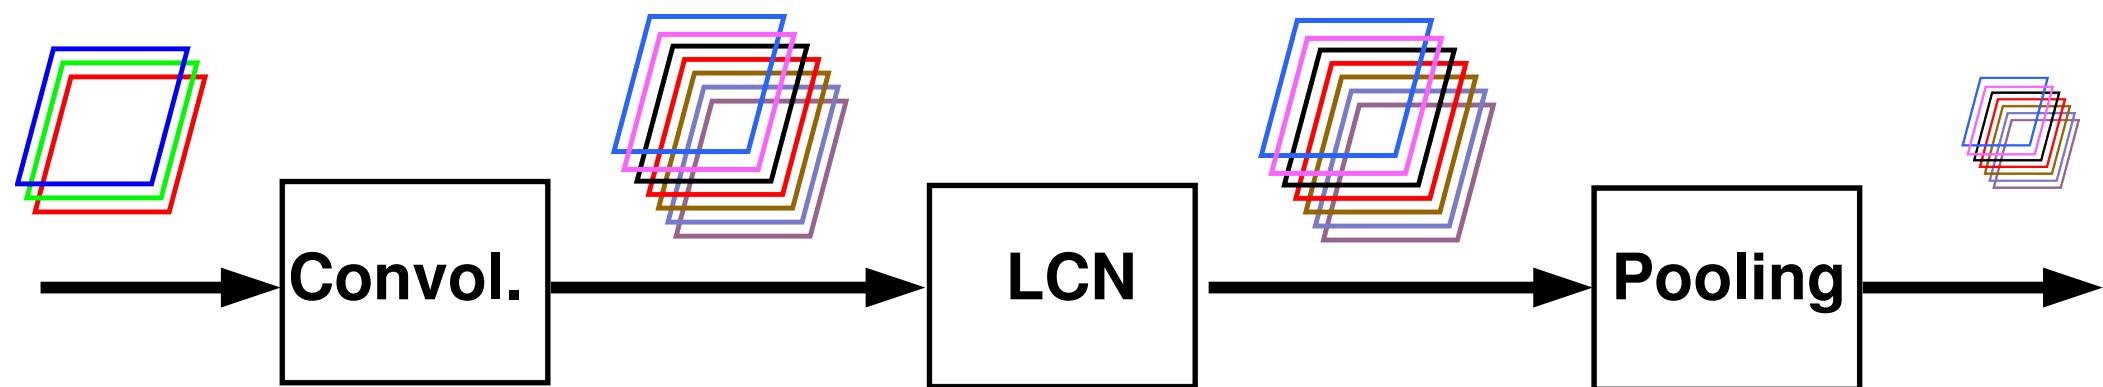
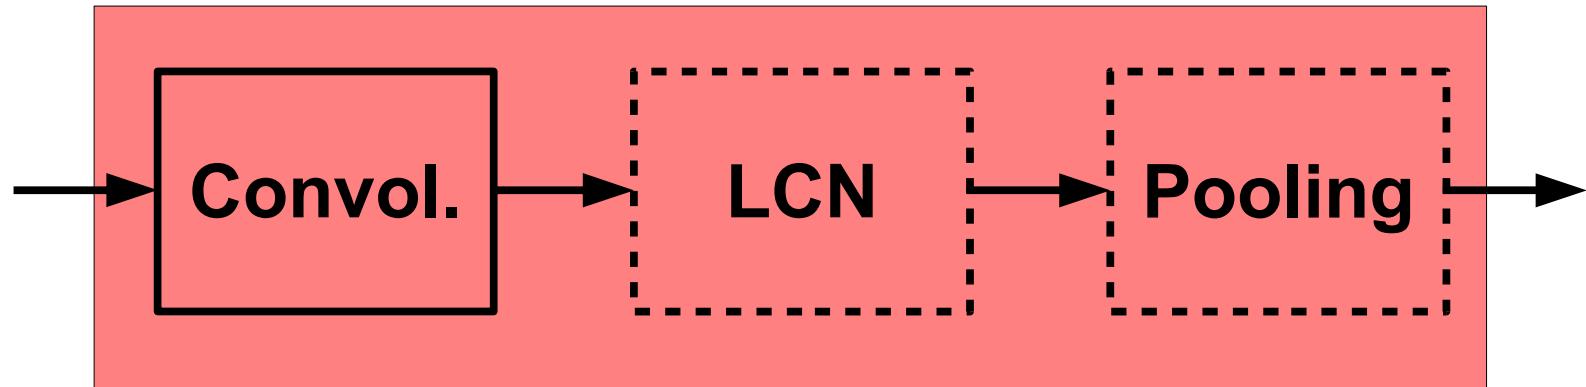
Performed also across features and in the higher layers.

Effects:

- improves invariance
- improves optimization
- increases sparsity

# CONV NETS: TYPICAL ARCHITECTURE

One stage (zoom)

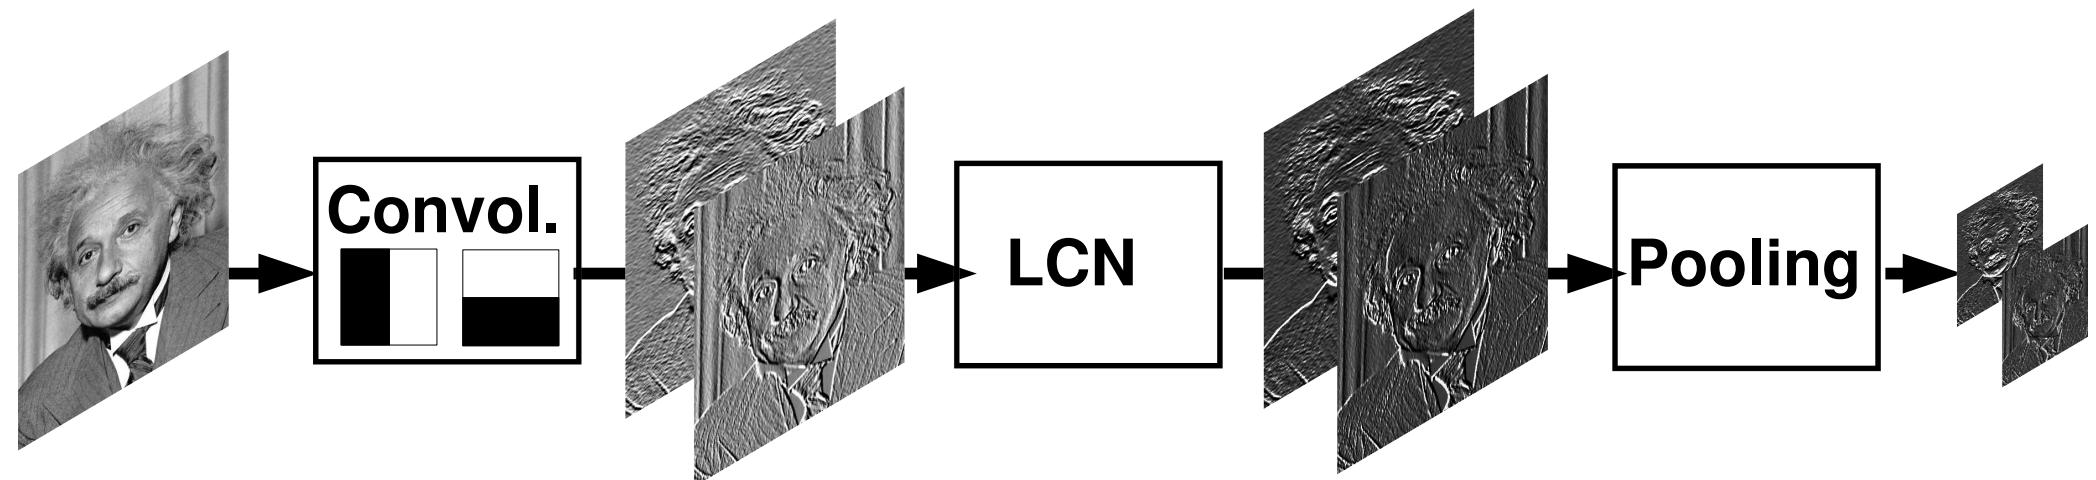
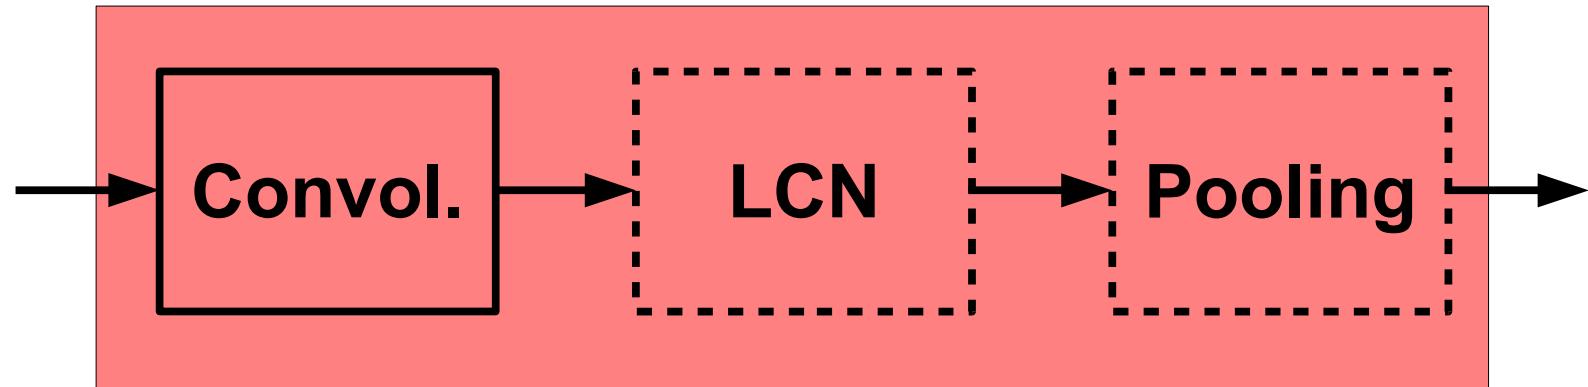


Convolutional layer increases nr. feature maps.

Pooling layer decreases spatial resolution.

# CONV NETS: TYPICAL ARCHITECTURE

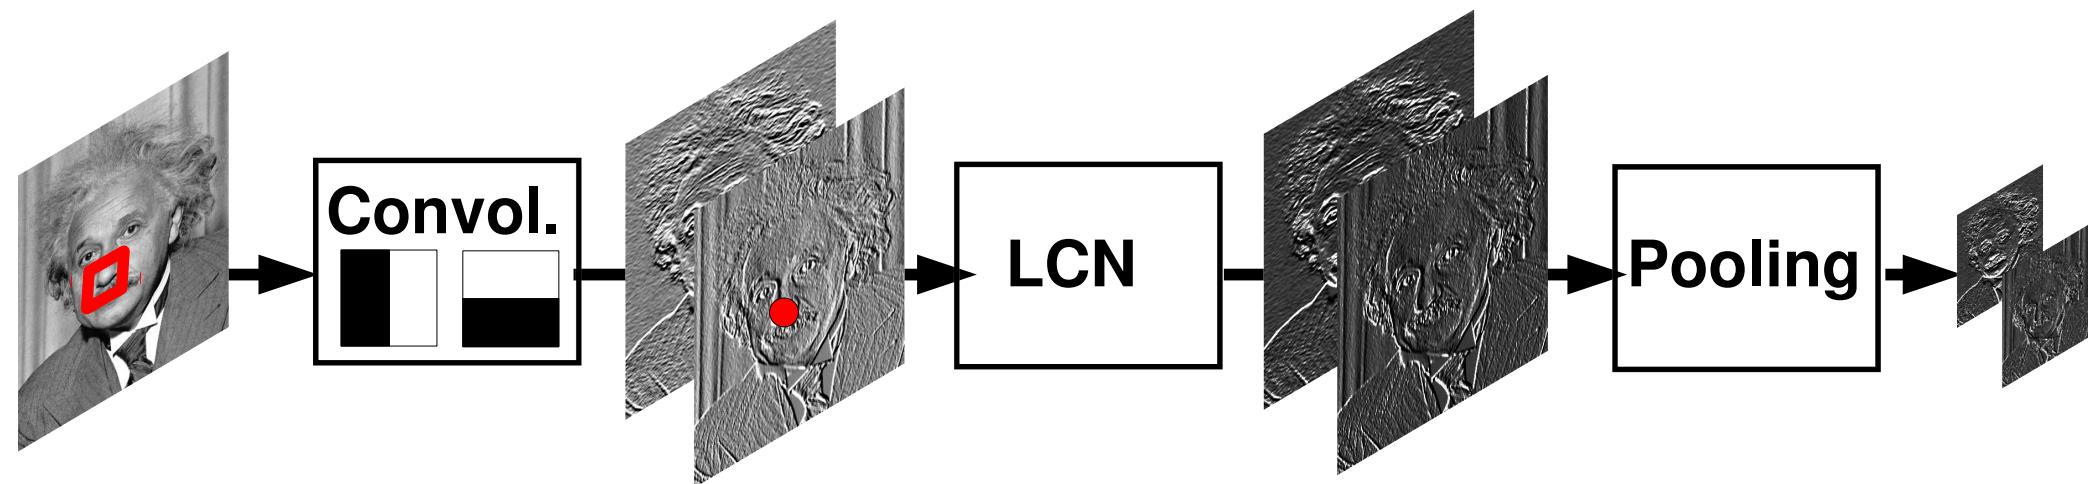
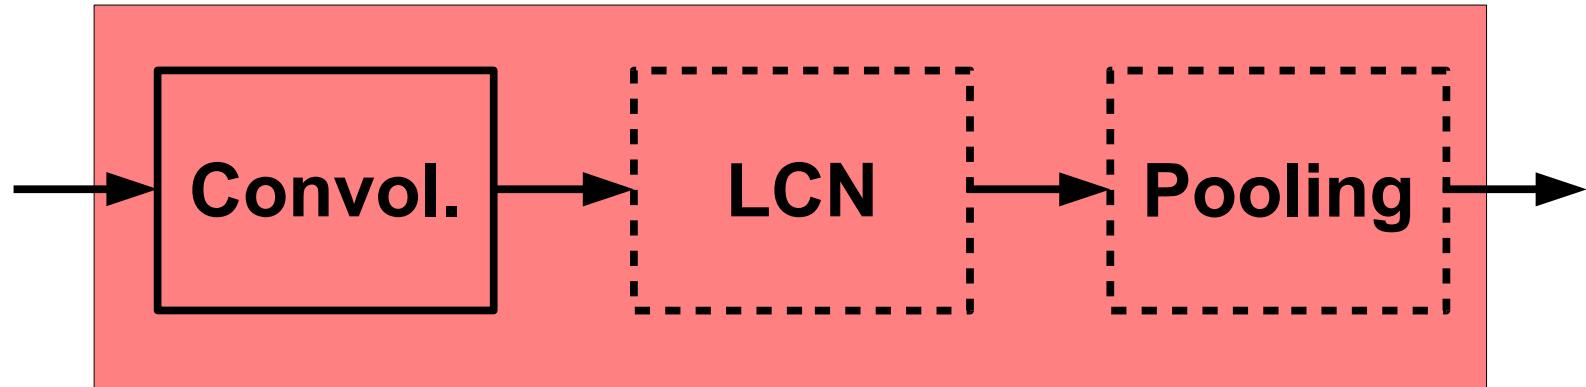
One stage (zoom)



Example with only two filters.

# CONV NETS: TYPICAL ARCHITECTURE

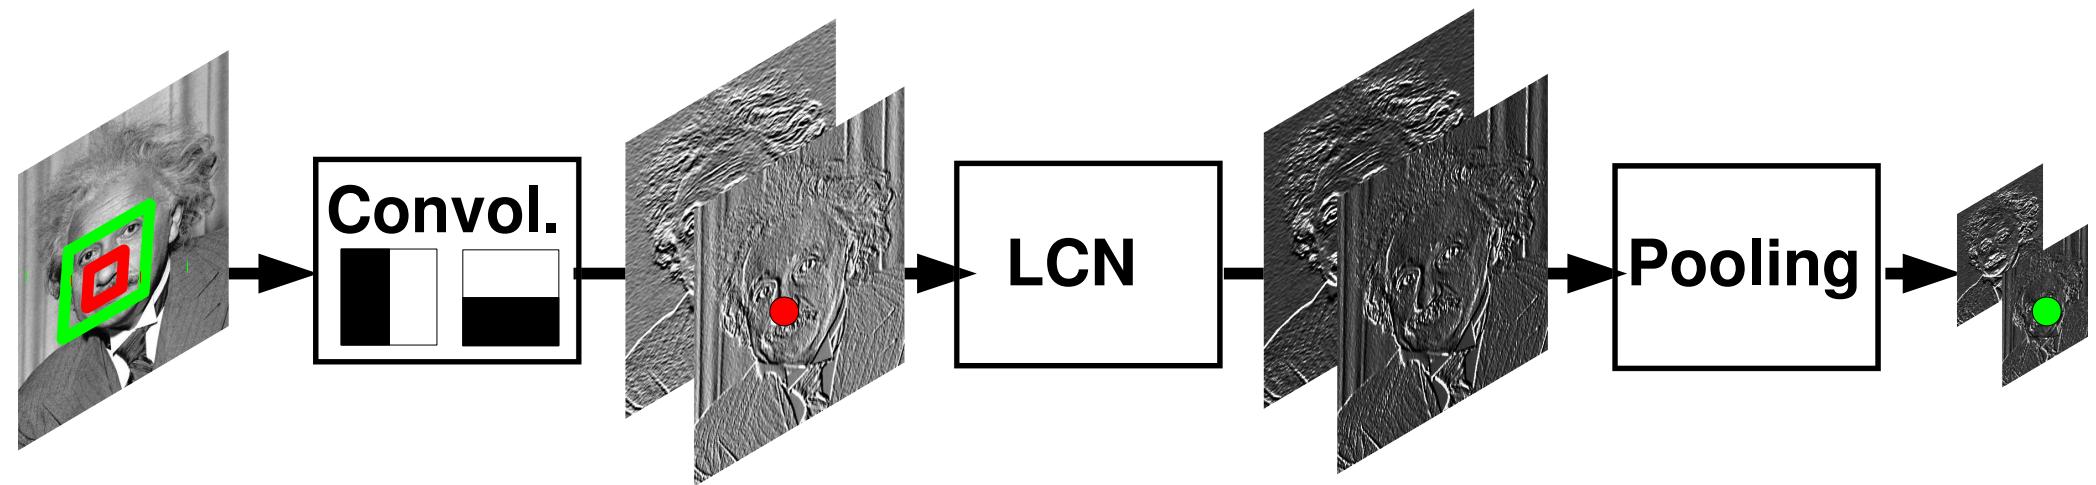
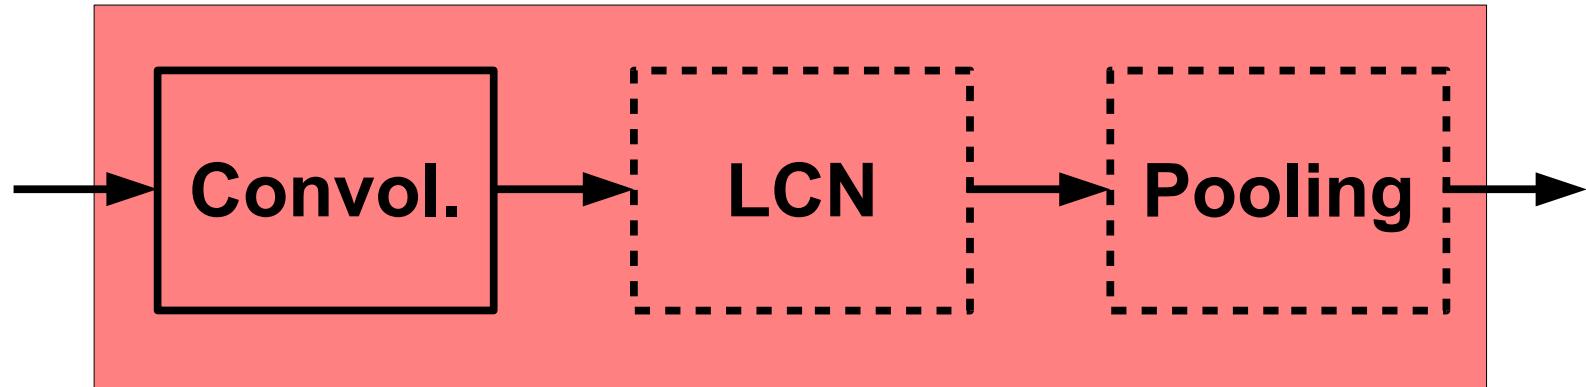
One stage (zoom)



A hidden unit in the first hidden layer is influenced by a small neighborhood (equal to size of filter).

# CONV NETS: TYPICAL ARCHITECTURE

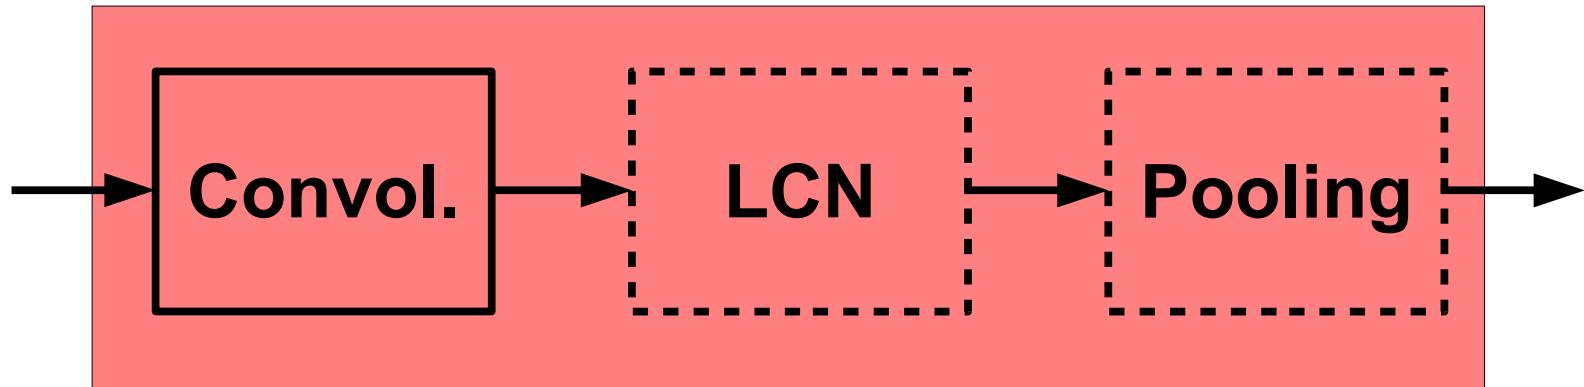
One stage (zoom)



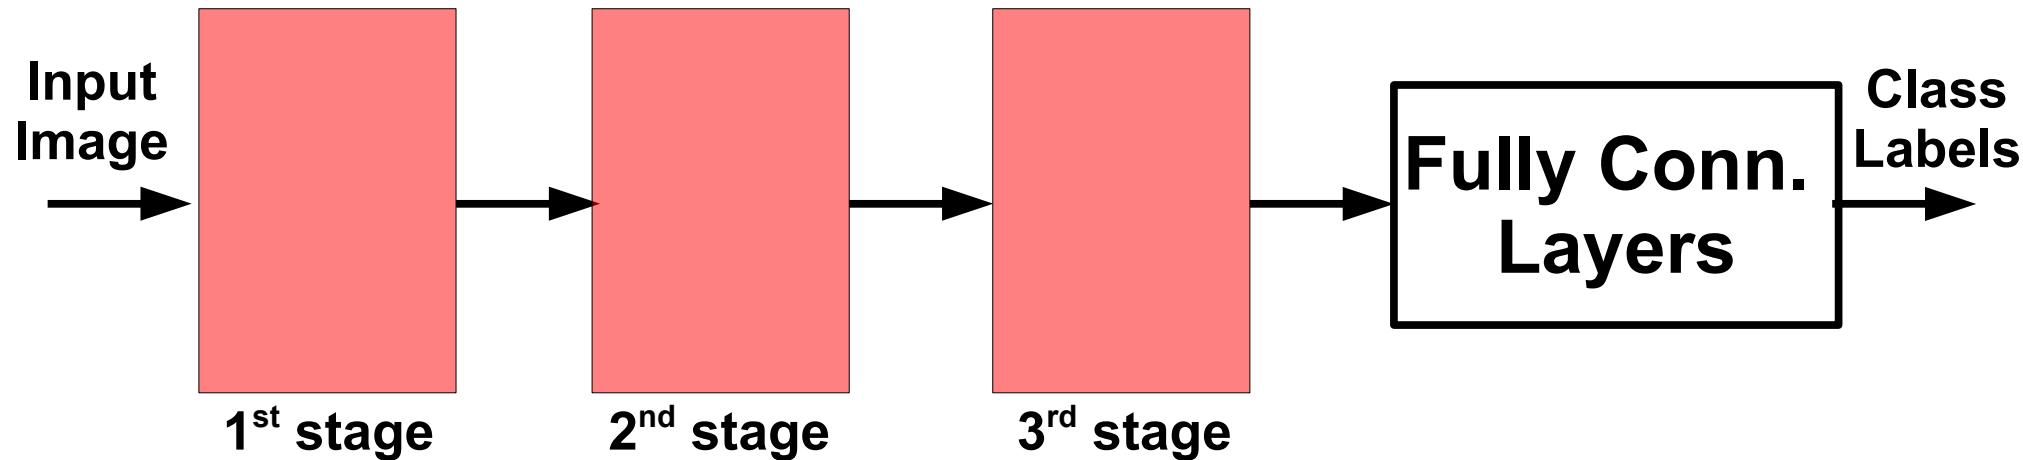
A hidden unit after the pooling layer is influenced by a larger neighborhood (it depends on filter sizes and strides).

# CONV NETS: TYPICAL ARCHITECTURE

One stage (zoom)



Whole system



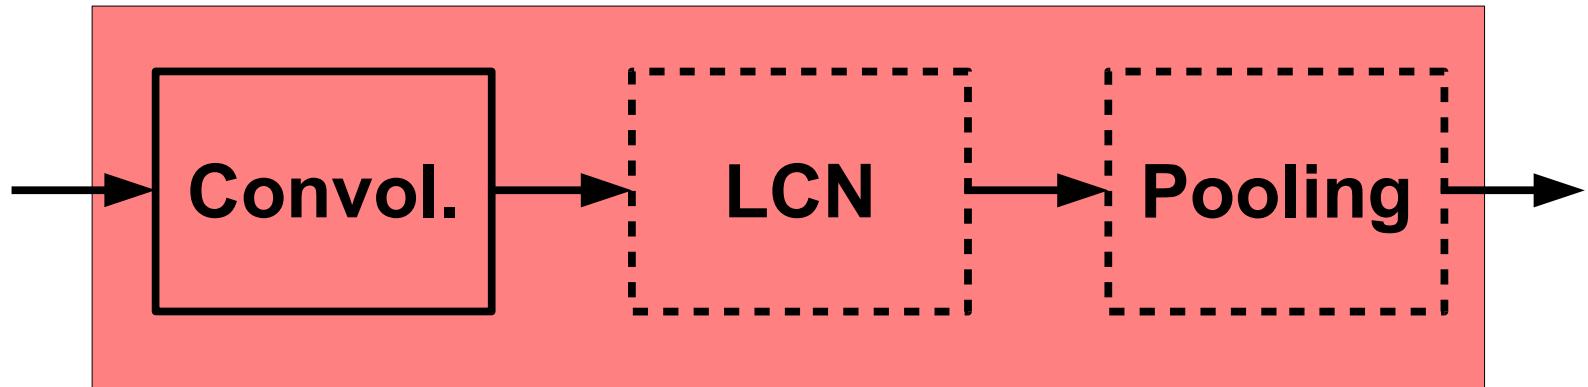
After a few stages, residual spatial resolution is very small.

93

We have learned a descriptor for the whole image.

# CONV NETS: TYPICAL ARCHITECTURE

One stage (zoom)



Conceptually similar to:

SIFT → K-Means → Pyramid Pooling → SVM

Lazebnik et al. "...Spatial Pyramid Matching..." CVPR 2006

SIFT → Fisher Vect. → Pooling → SVM

Sanchez et al. "Image classification with F.V.: Theory and practice" IJCV 2012

# CONV NETS: TRAINING

All layers are differentiable (a.e.).

We can use standard back-propagation.

## Algorithm:

**Given a small mini-batch**

- **F-PROP**
- **B-PROP**
- **PARAMETER UPDATE**

# KEY IDEAS: CONV. NETS

Conv. Nets have special layers like:

- pooling, and
- local contrast normalization

Back-propagation can still be applied.

These layers are useful to:

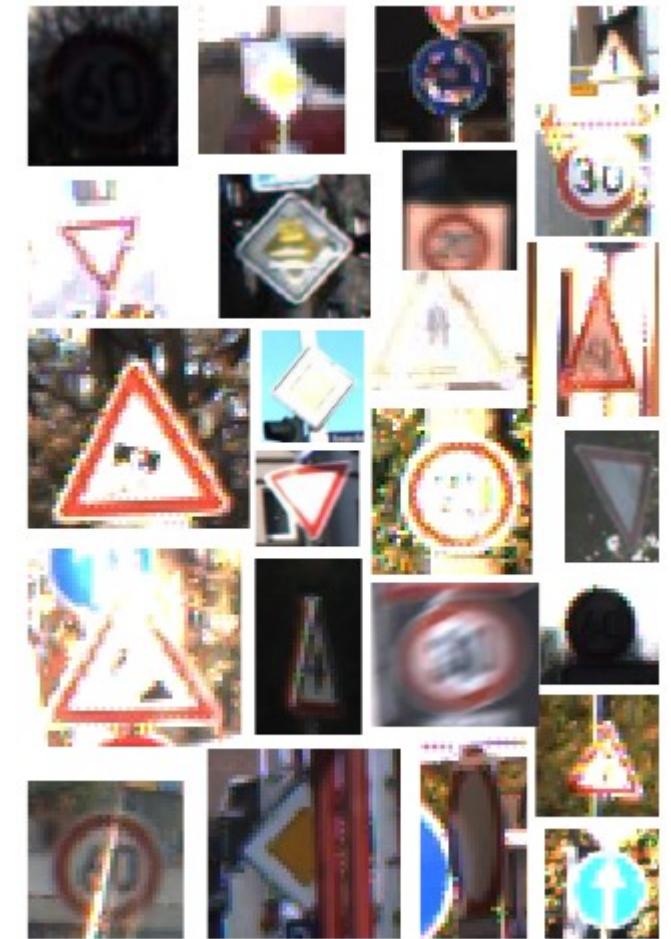
- reduce computational burden
- increase invariance
- ease the optimization

# Outline

- Motivation
- Deep Learning: The Big Picture
- From neural nets to convolutional nets
- Applications
- A practical guide

# CONV NETS: EXAMPLES

- OCR / House number & Traffic sign classification



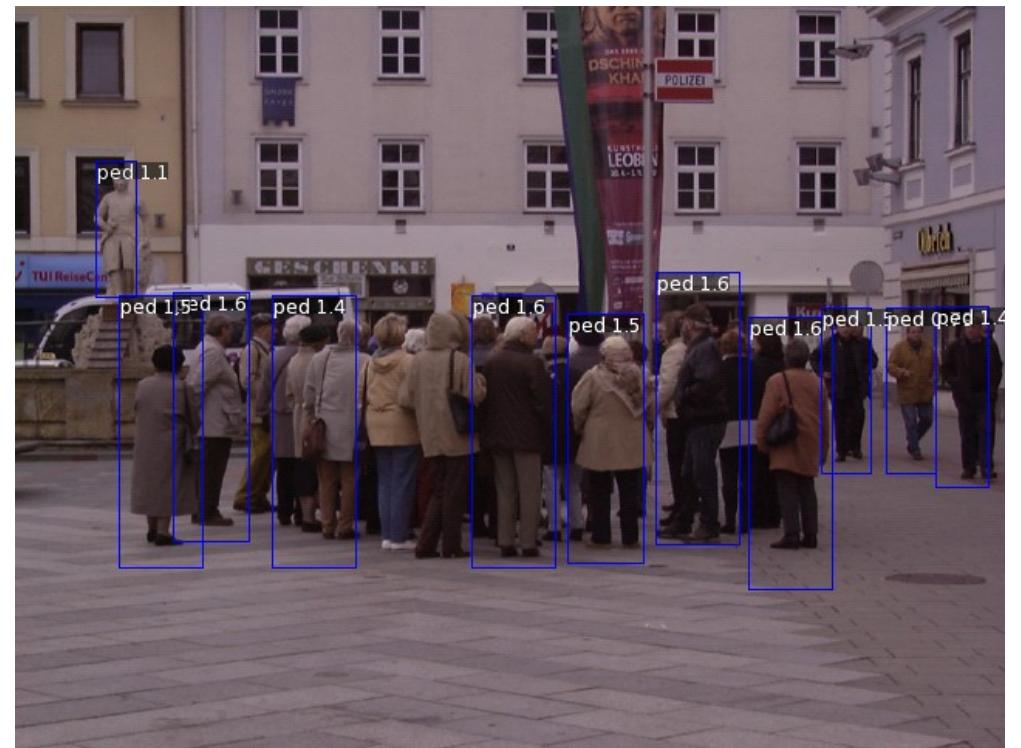
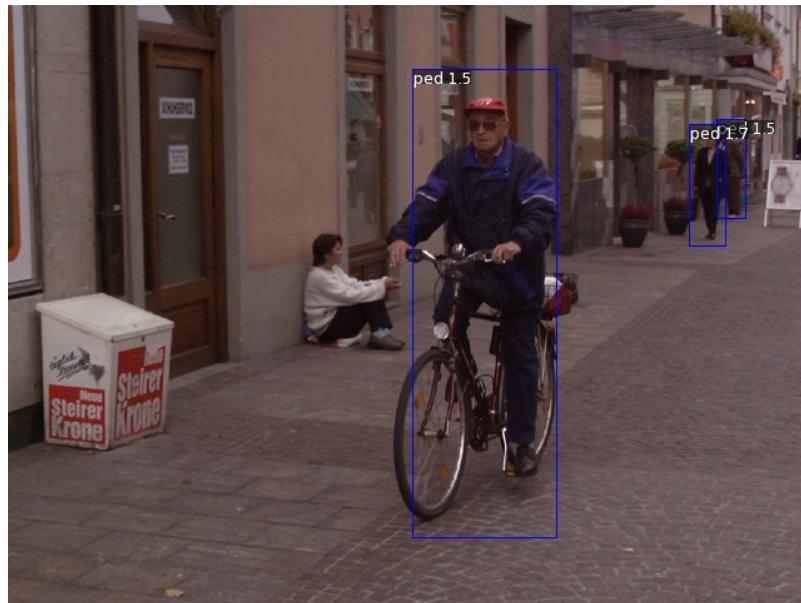
# CONV NETS: EXAMPLES

## - Texture classification



# CONV NETS: EXAMPLES

## - Pedestrian detection

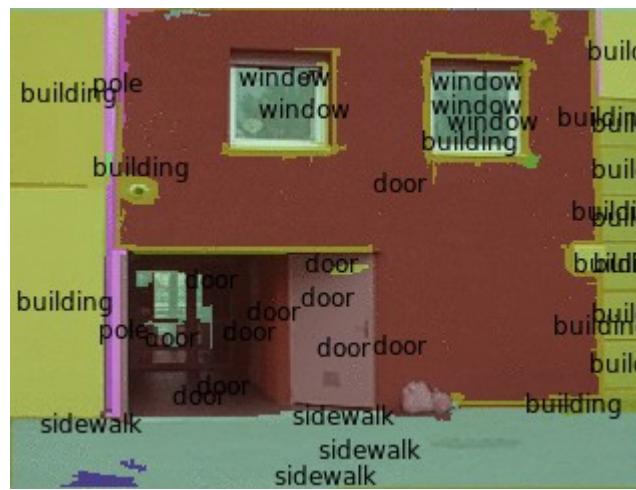


100

Sermanet et al. “Pedestrian detection with unsupervised multi-stage..” CVPR 2013

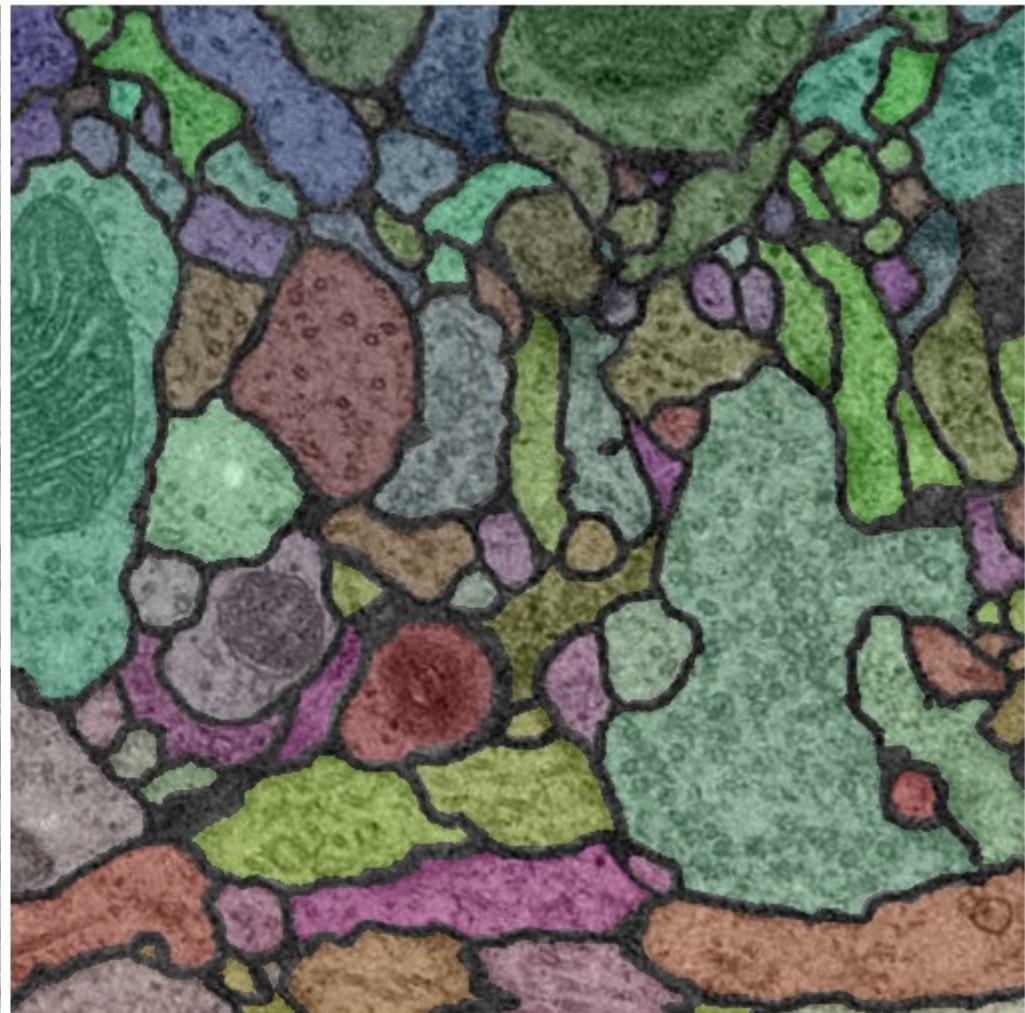
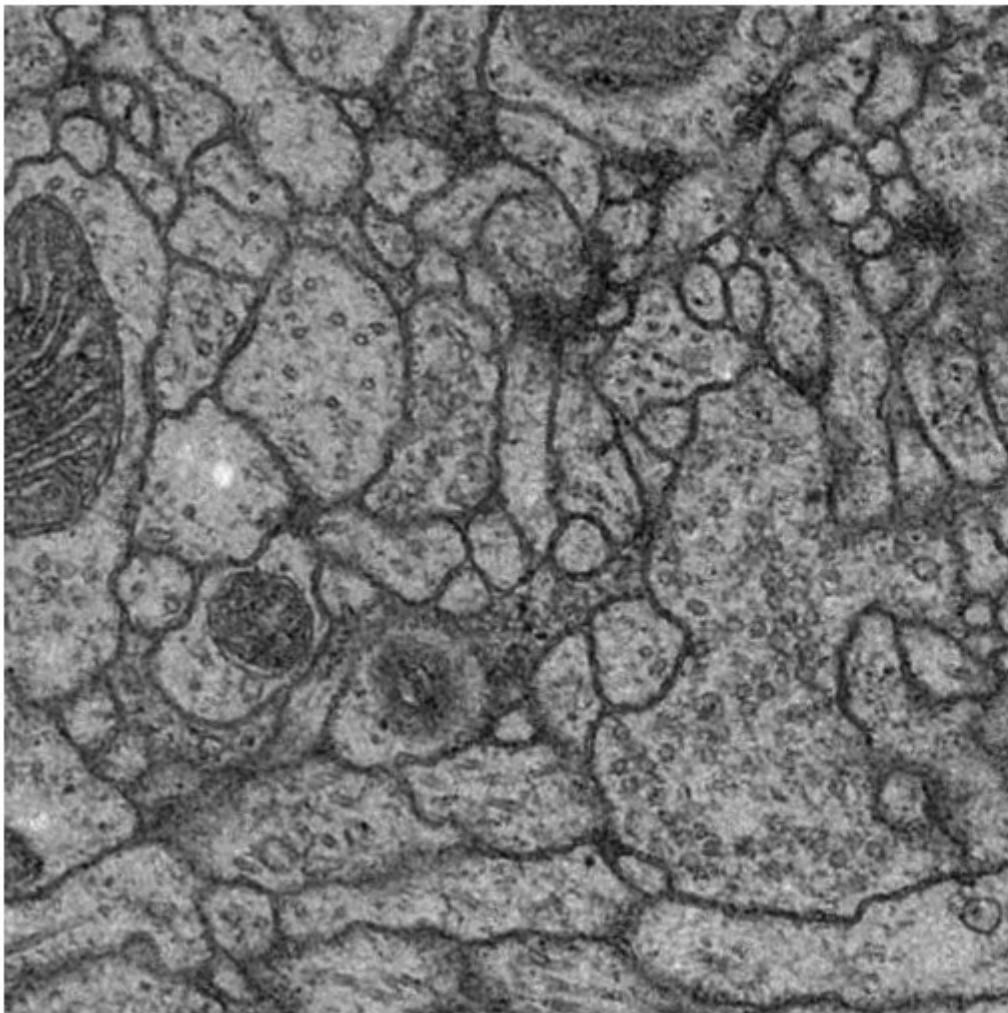
# CONV NETS: EXAMPLES

## - Scene Parsing



# CONV NETS: EXAMPLES

- Segmentation 3D volumetric images



# CONV NETS: EXAMPLES

- Action recognition from videos



103

# CONV NETS: EXAMPLES

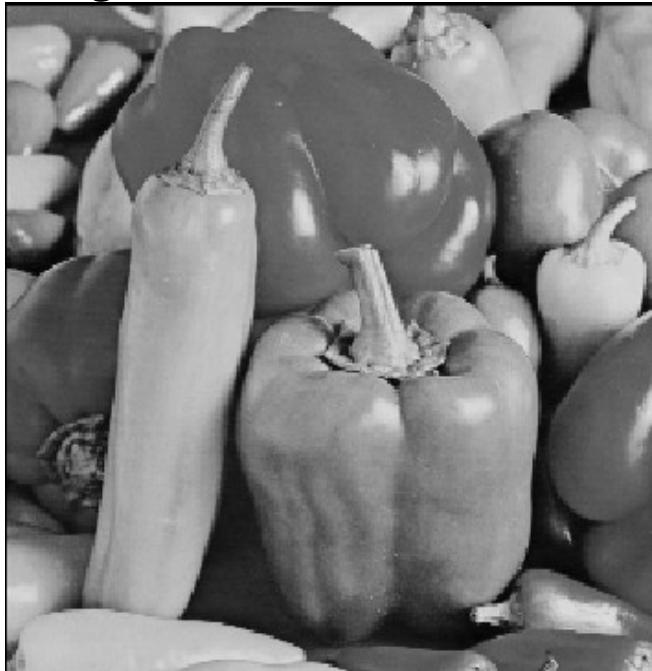
- Robotics



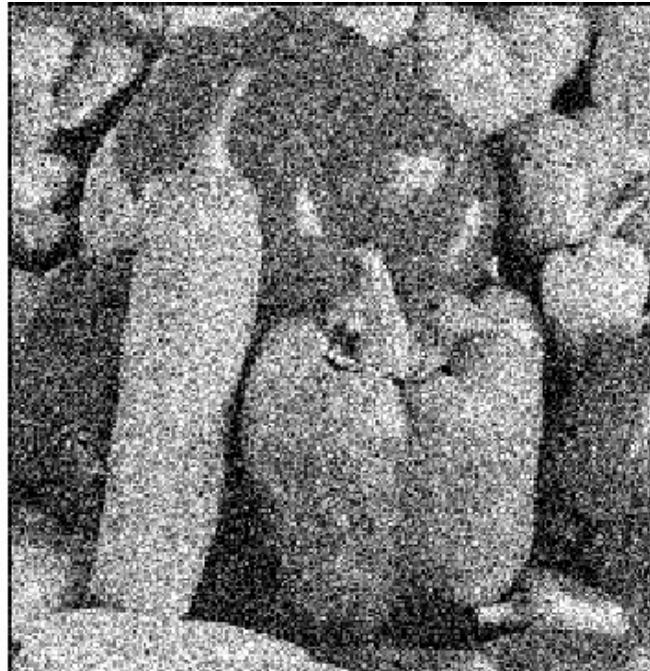
# CONV NETS: EXAMPLES

## - Denoising

original



noised

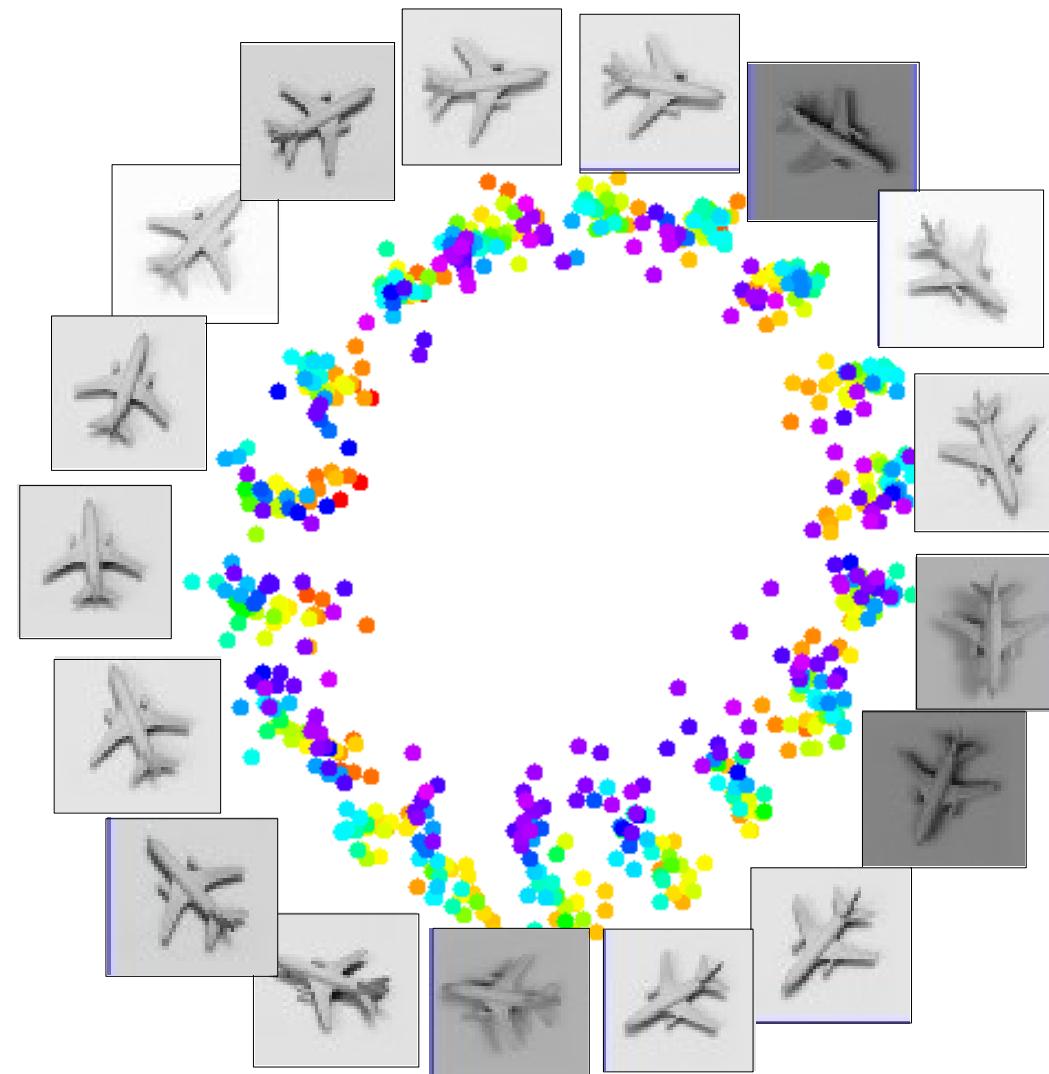


denoised



# CONV NETS: EXAMPLES

- Dimensionality reduction / learning embeddings

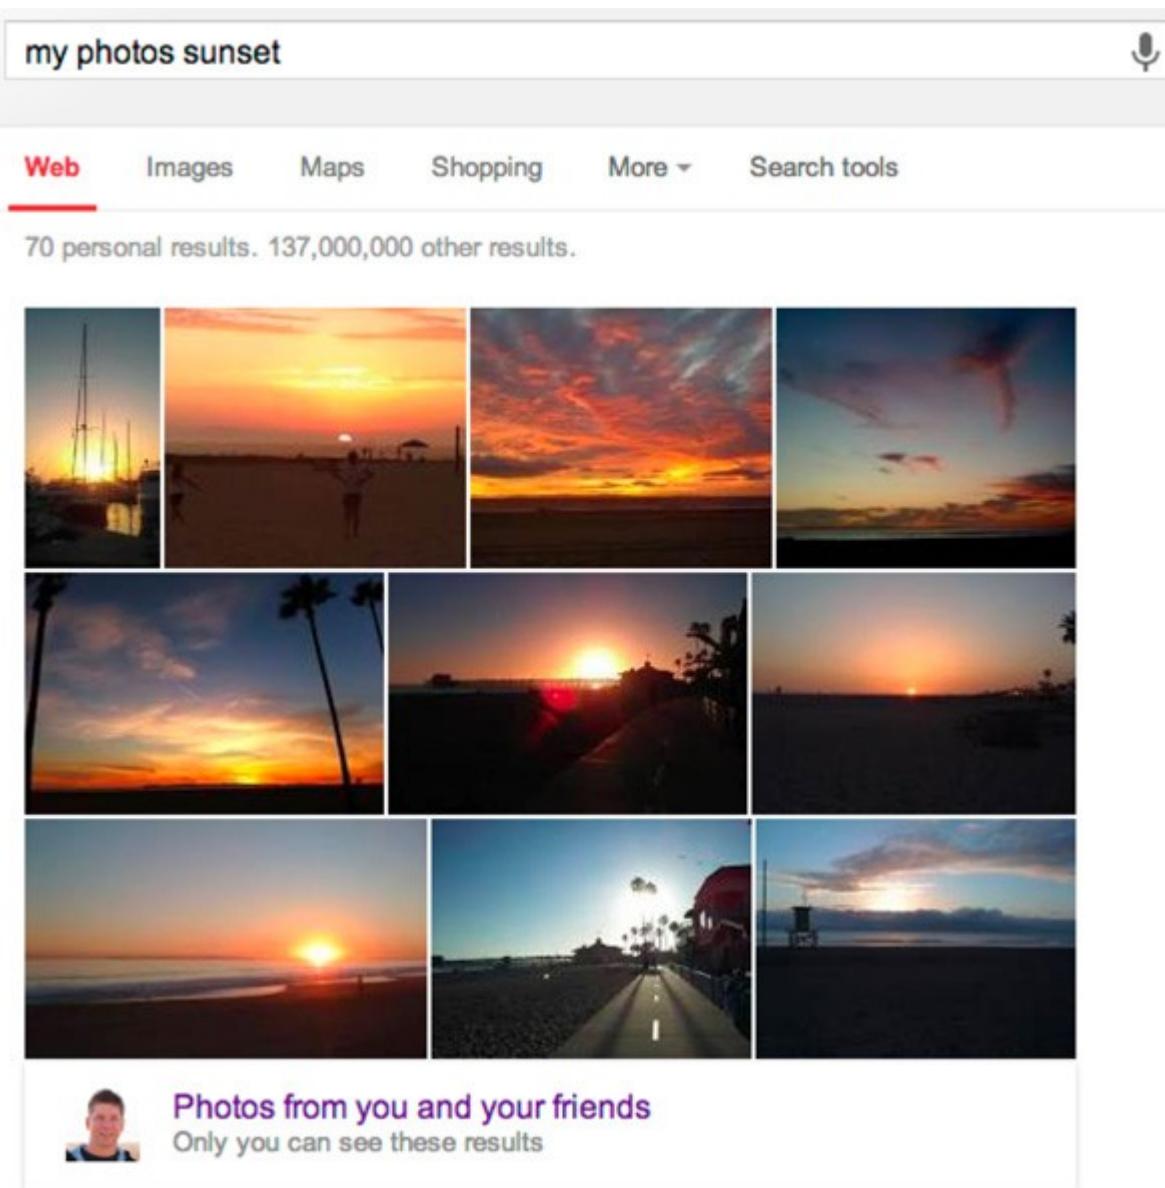


106

Hadsell et al. "Dimensionality reduction by learning an invariant mapping" CVPR 2006

# CONV NETS: EXAMPLES

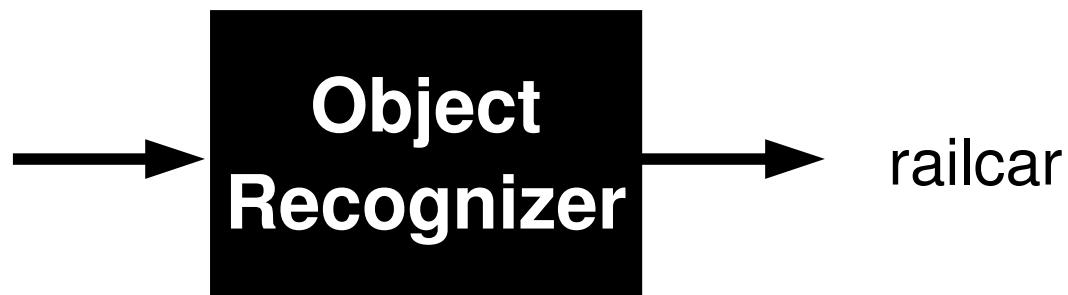
- Deployed in commercial systems (Google & Baidu, spring 2013)



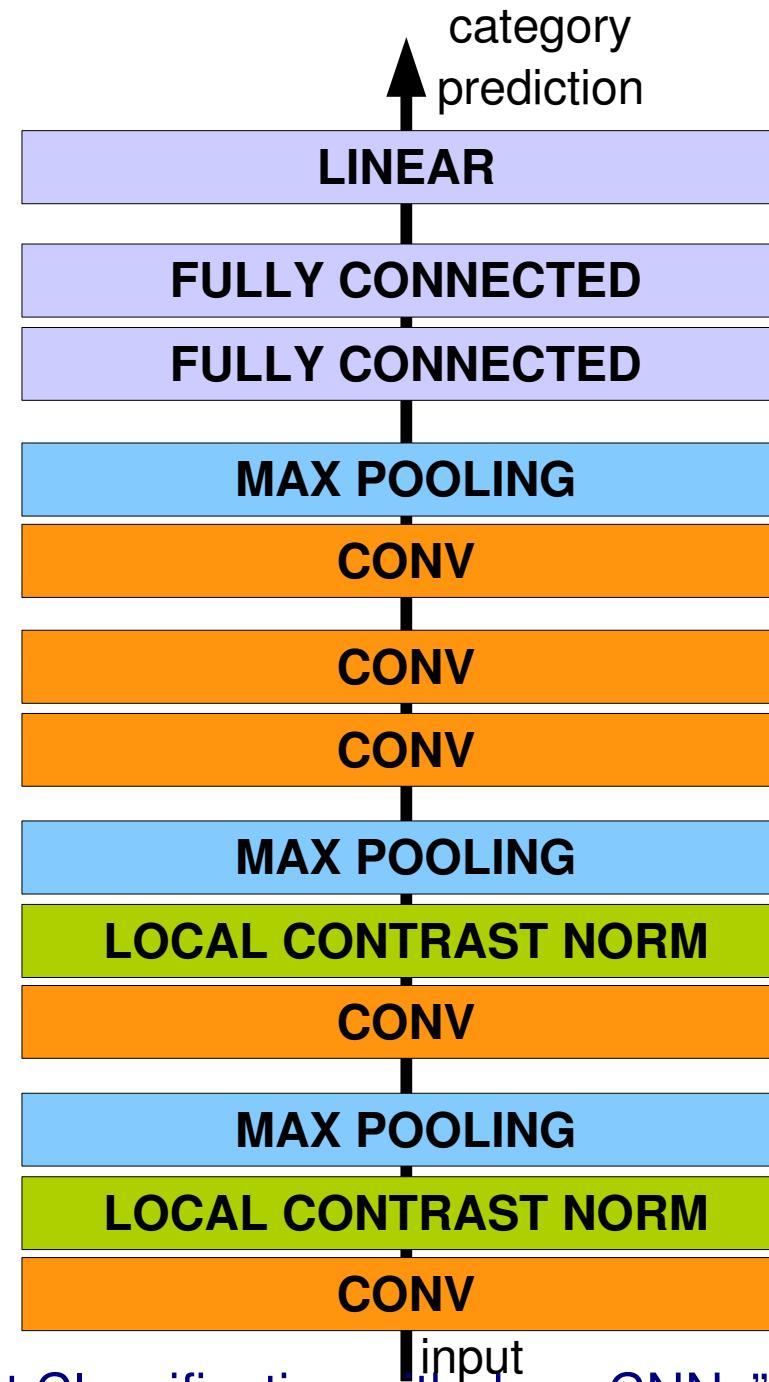
# CONV NETS: EXAMPLES

- Image classification

IMAGENET

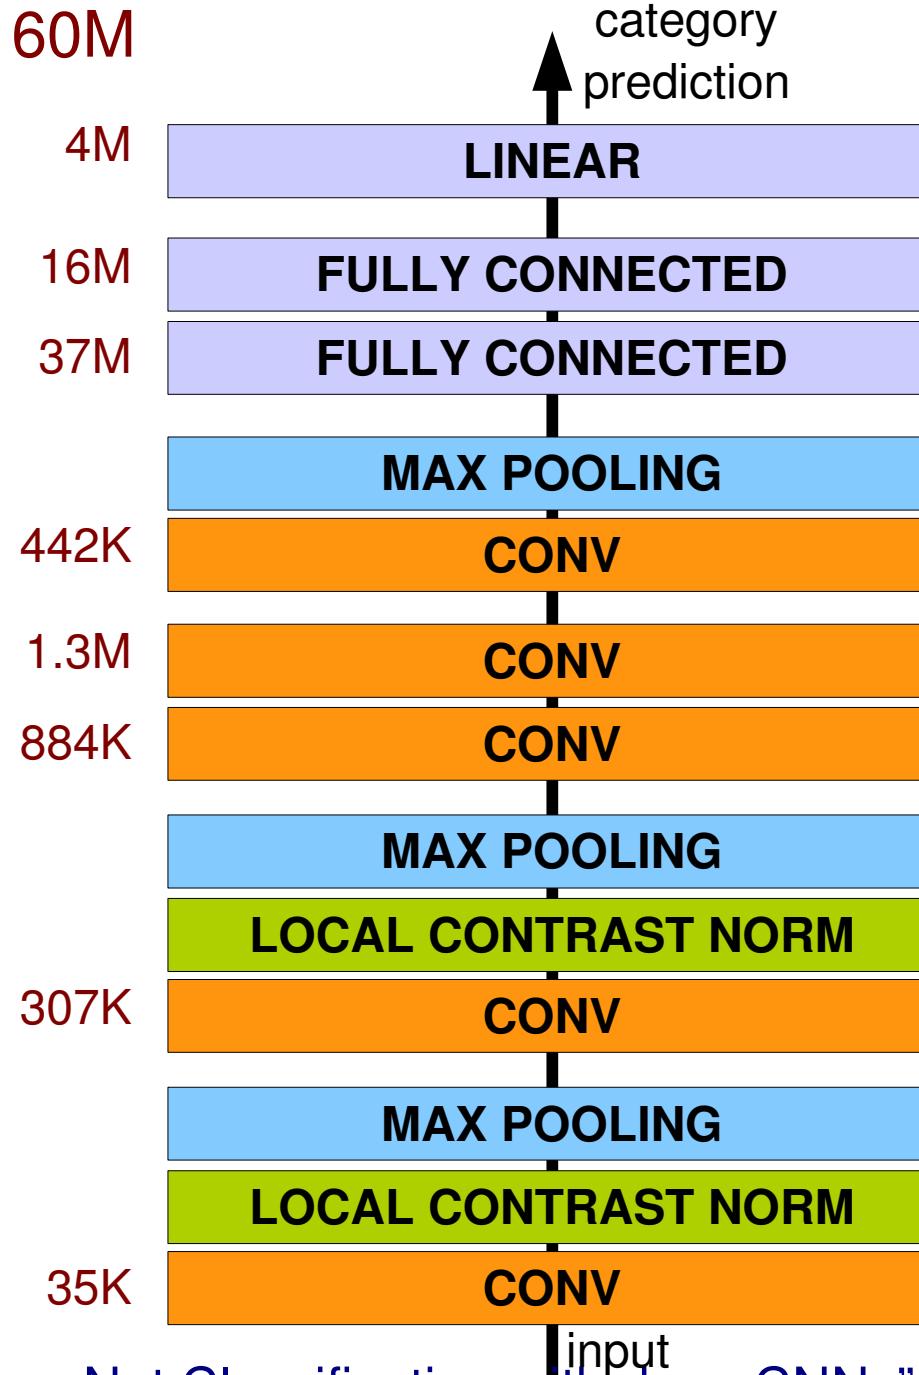


# Architecture



# Architecture

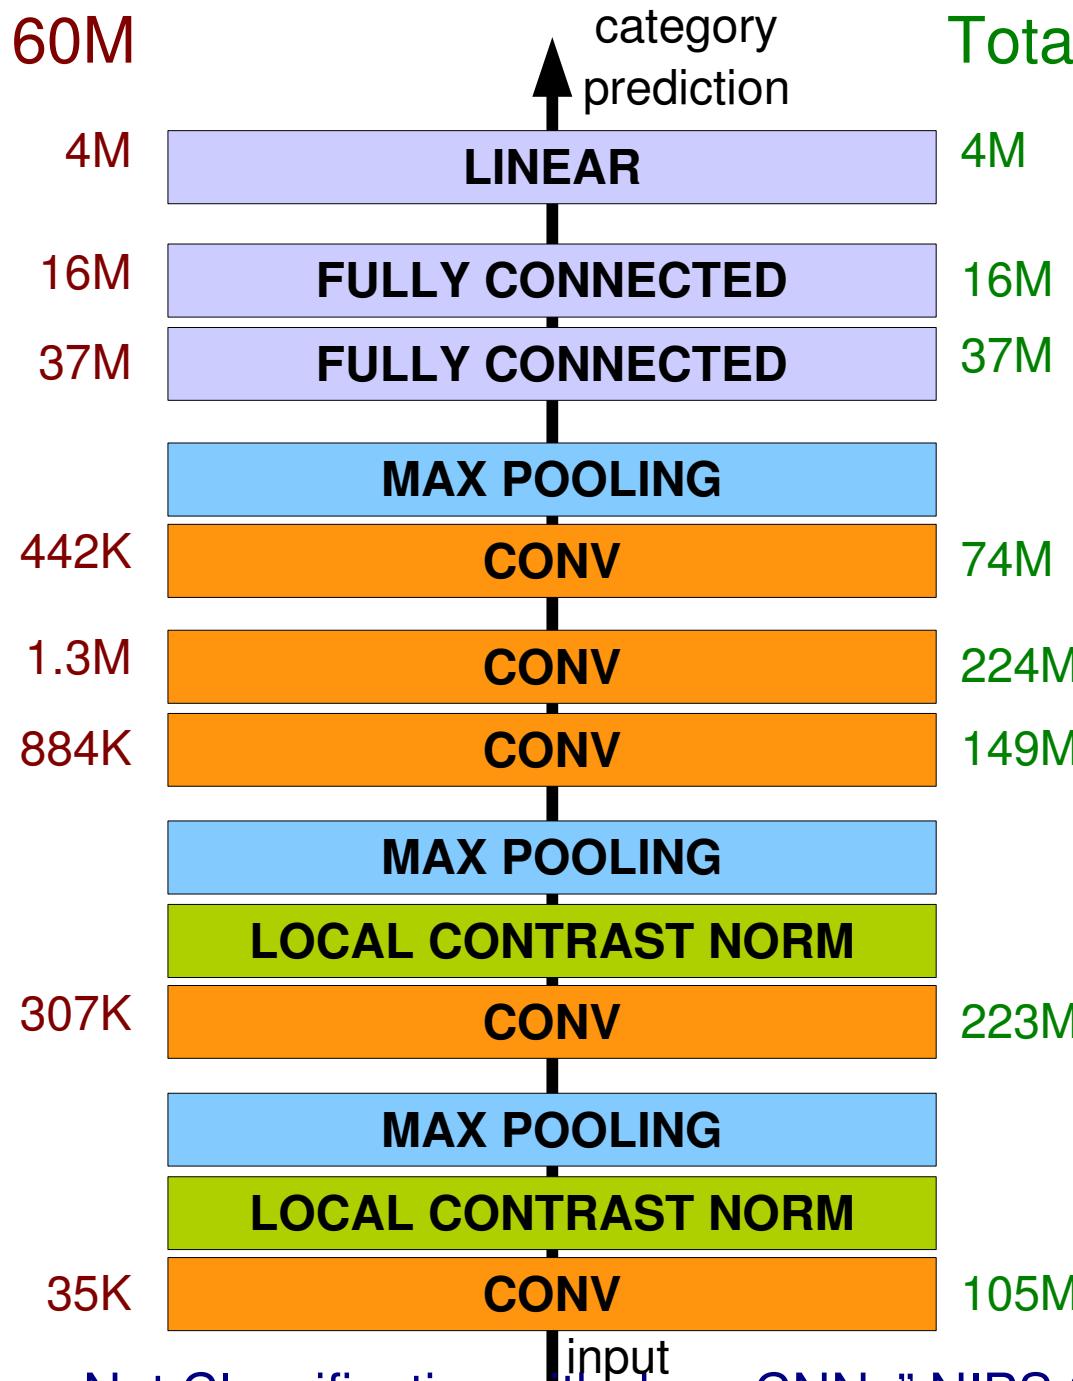
Total nr. params: 60M



# Architecture

Total nr. params: 60M

Total nr. flops: 832M



# Optimization

## SGD with momentum:

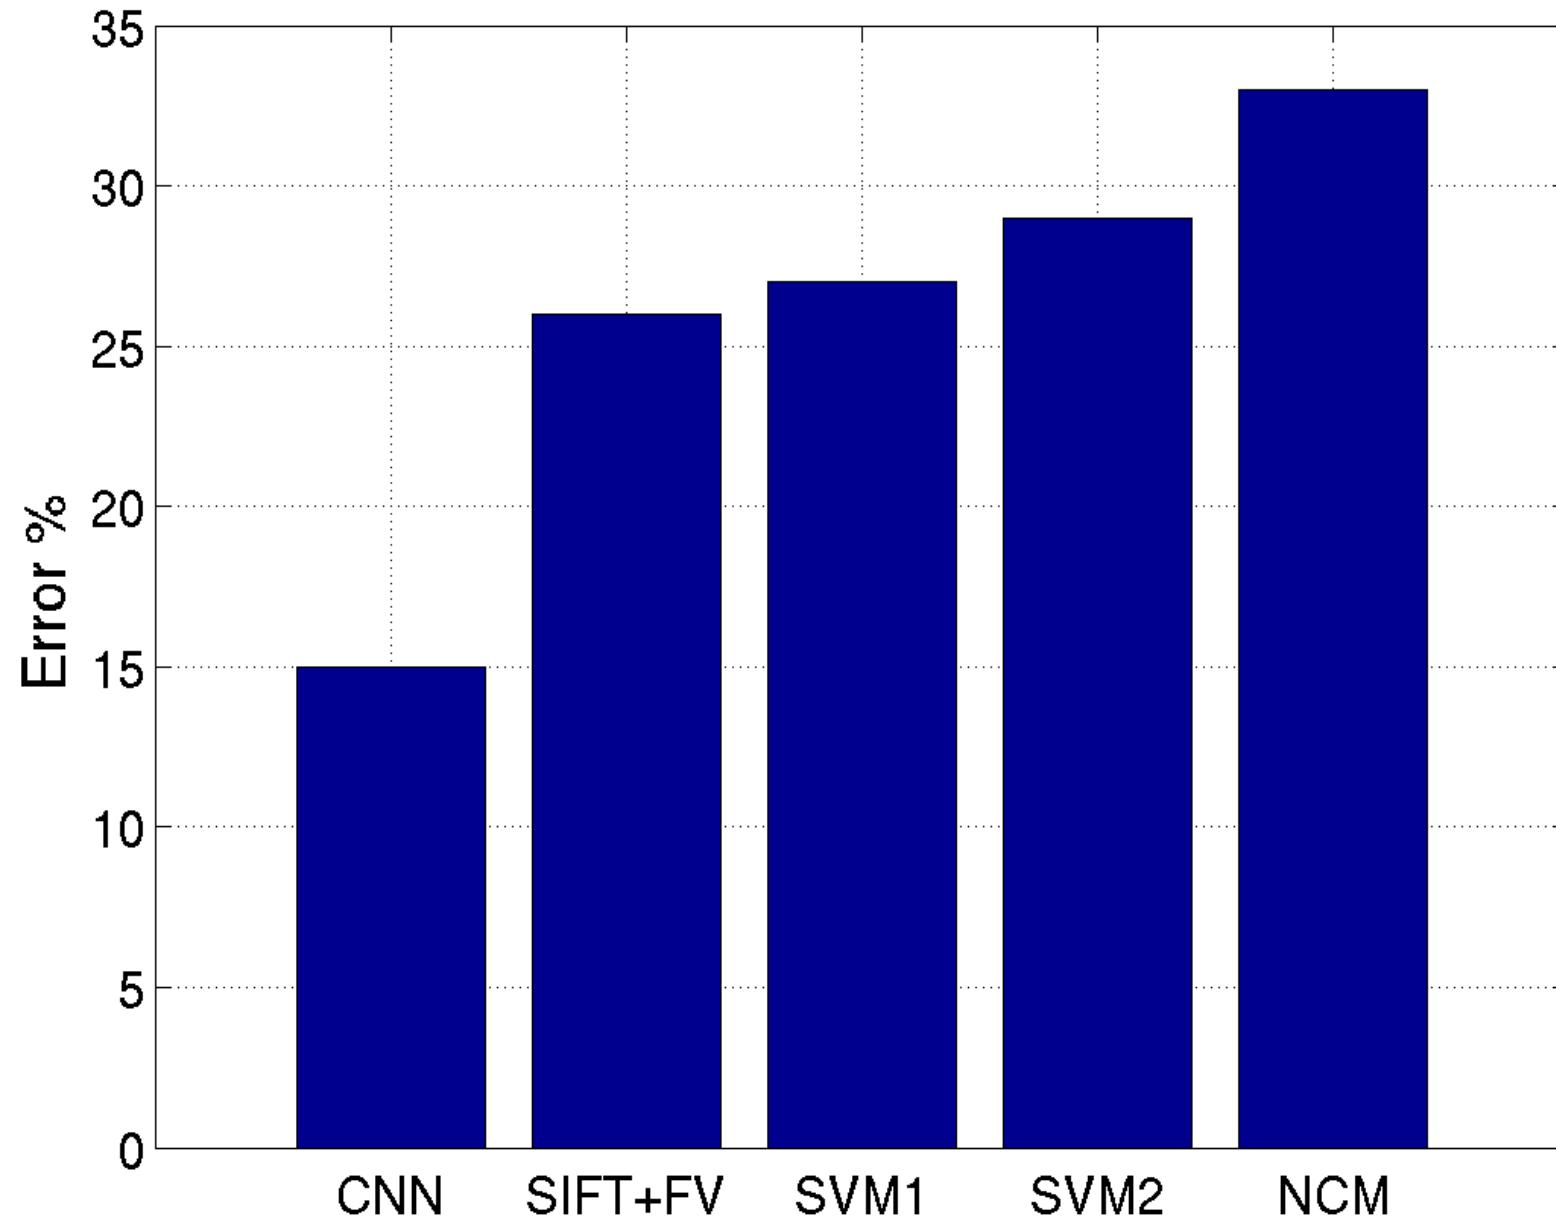
- Learning rate = 0.01
- Momentum = 0.9

## Improving generalization by:

- Weight sharing (convolution)
- Input distortions
- Dropout = 0.5
- Weight decay = 0.0005

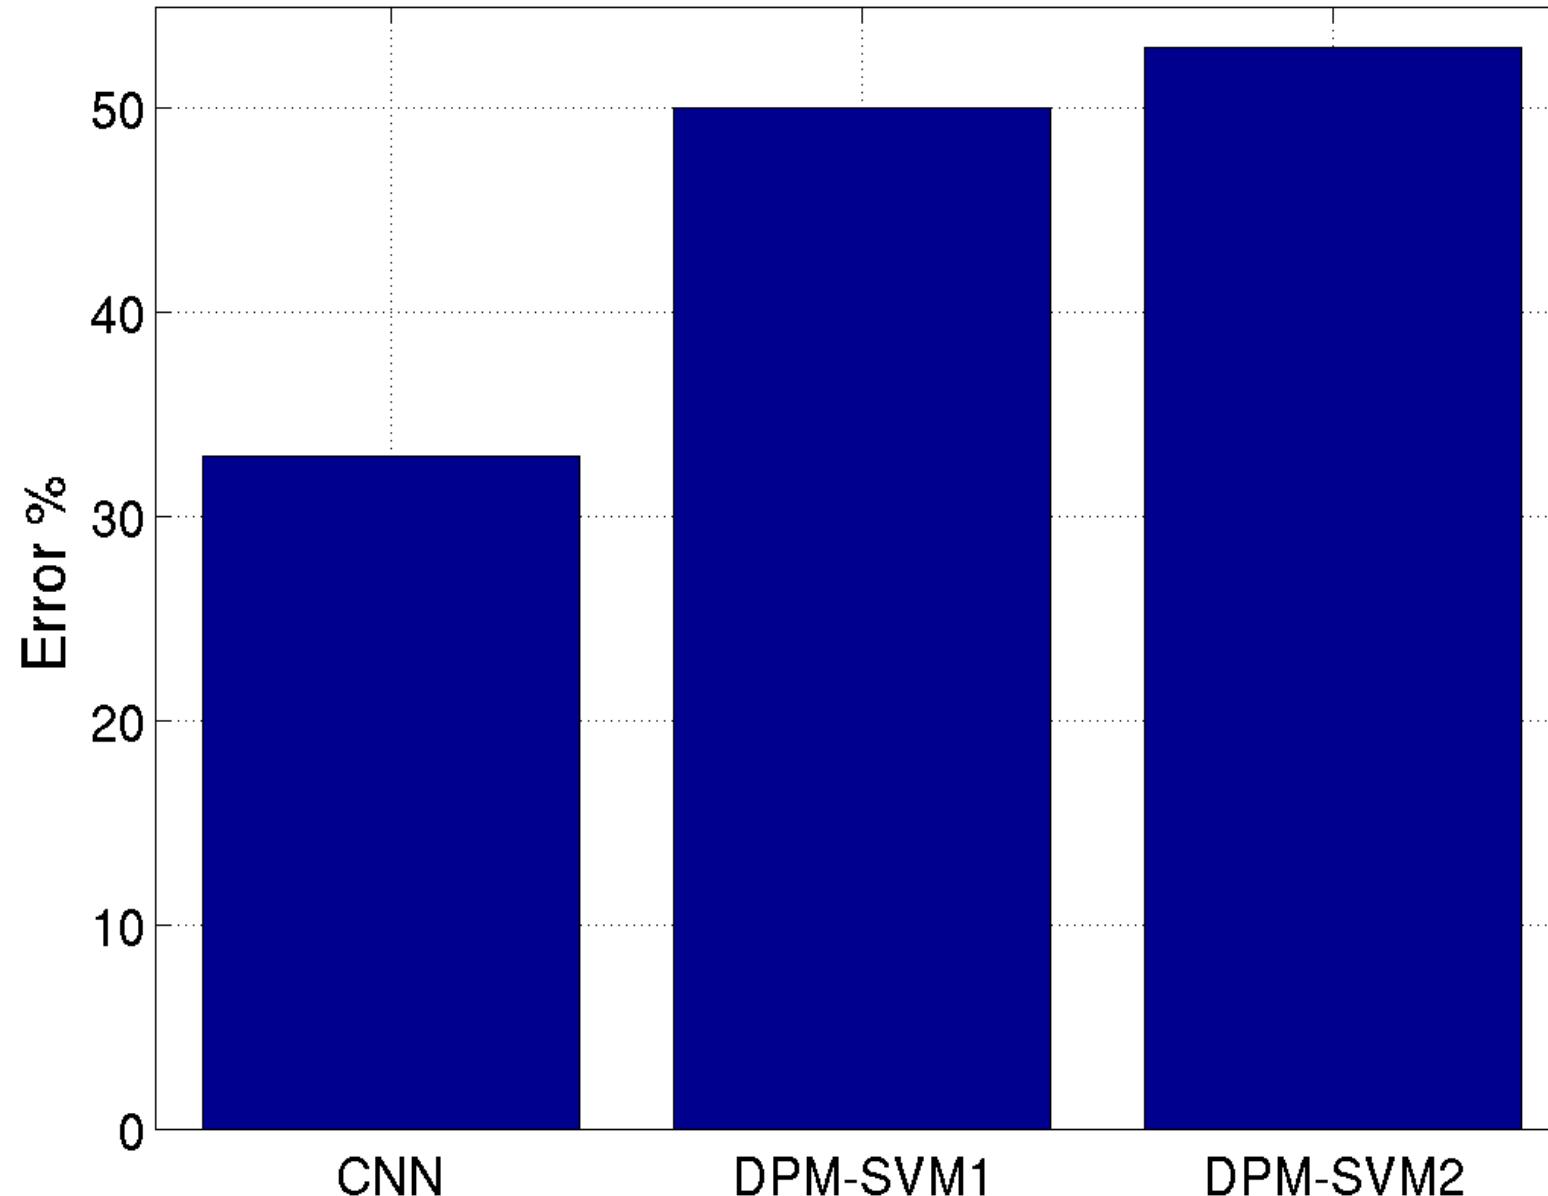
# Results: ILSVRC 2012

## TASK 1 - CLASSIFICATION



# Results: ILSVRC 2012

## TASK 2 - DETECTION



# Results



First layer learned filters (processing raw pixel values).



mite

mite  
black widow  
cockroach  
tick  
starfish

container ship

lifeboat  
amphibian  
fireboat  
drilling platform

motor scooter

go-kart  
moped  
bumper car  
golfcart

leopard

leopard  
jaguar  
cheetah  
snow leopard  
Egyptian cat



grille

convertible  
grille  
pickup  
beach wagon  
fire engine

mushroom

agaric  
mushroom  
jelly fungus  
gill fungus  
dead-man's-fingers

cherry

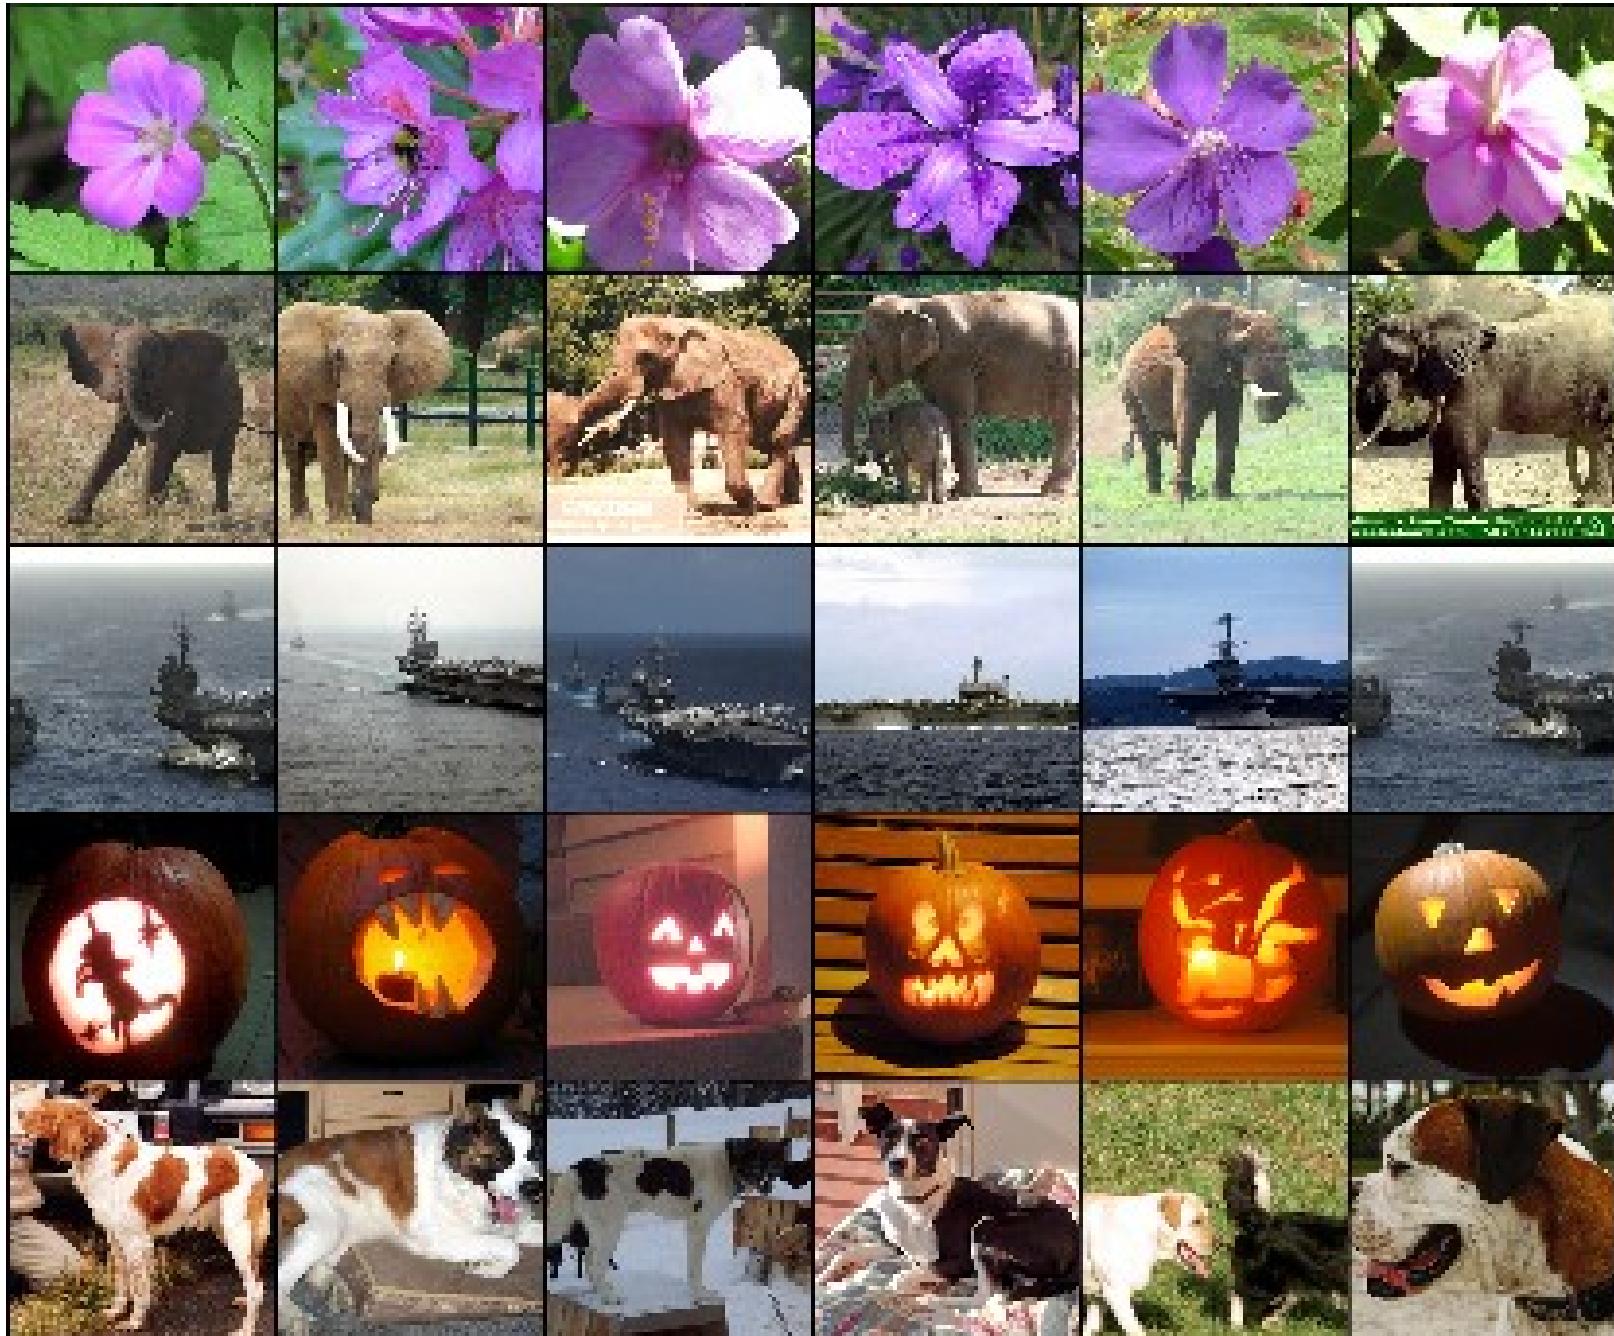
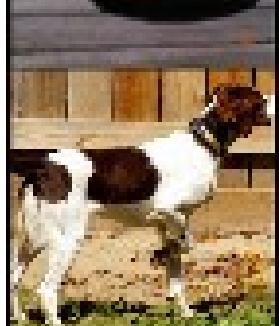
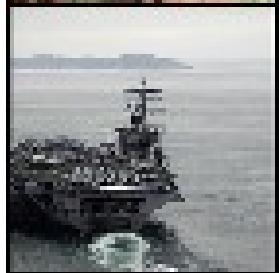
dalmatian  
grape  
elderberry  
ffordshire bullterrier  
currant

Madagascar cat

squirrel monkey  
spider monkey  
titi  
indri  
howler monkey

# TEST IMAGE

# RETRIEVED IMAGES



# Outline

- Motivation
- Deep Learning: The Big Picture
- From neural nets to convolutional nets
- Applications
- A practical guide

# CHOOSING THE ARCHITECTURE

- [Convolution → LCN → pooling]\* + fully connected layer
- Cross-validation
- Task dependent
- The more data: the more layers and the more kernels
  - Look at the number of parameters at each layer
  - Look at the number of flops at each layer
- Computational cost
- Be creative :)

# HOW TO OPTIMIZE

- SGD (with momentum) usually works very well
- Pick learning rate by running on a subset of the data  
Bottou “Stochastic Gradient Tricks” Neural Networks 2012
  - Start with large learning rate and divide by 2 until loss does not diverge
  - Decay learning rate by a factor of ~100 or more by the end of training
- Use  non-linearity
- Initialize parameters so that each feature across layers has similar variance. Avoid units in saturation.

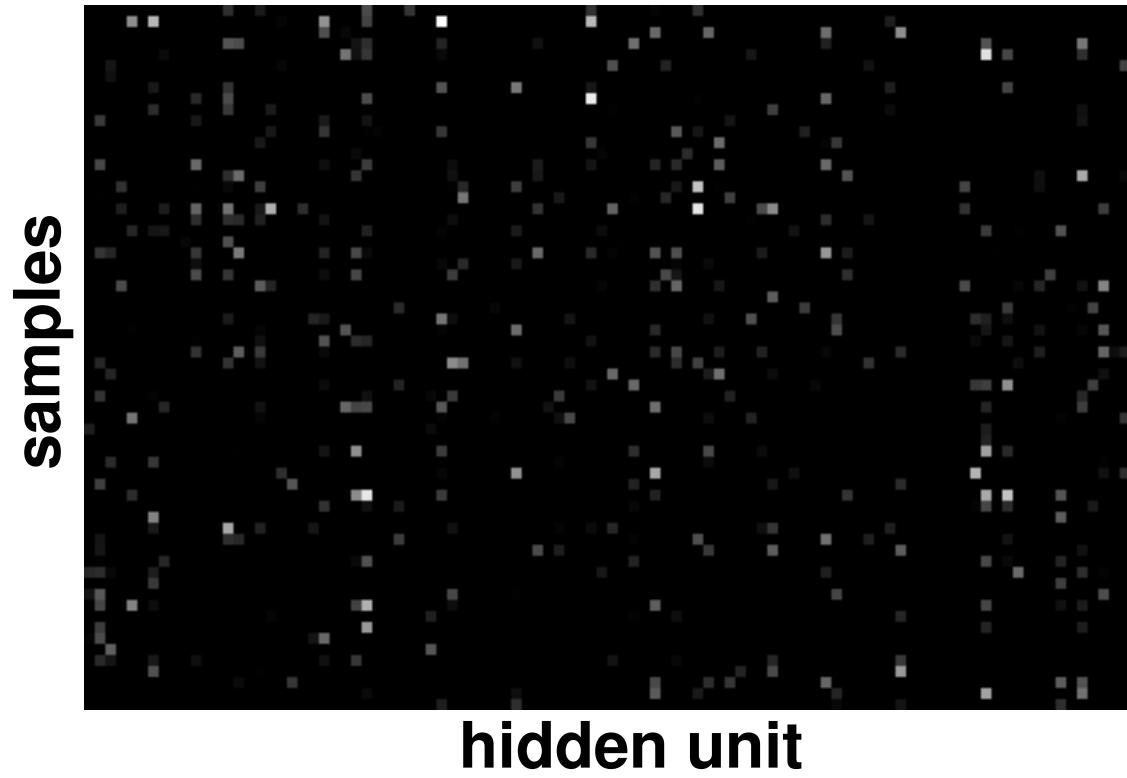
# HOW TO IMPROVE GENERALIZATION

- Weight sharing (greatly reduce the number of parameters)
- Data augmentation (e.g., jittering, noise injection, etc.)
- Dropout

Hinton et al. “Improving Nns by preventing co-adaptation of feature detectors”  
arxiv 2012
- Weight decay (L2, L1)
- Sparsity in the hidden units
- Multi-task (unsupervised learning)

# OTHER THINGS GOOD TO KNOW

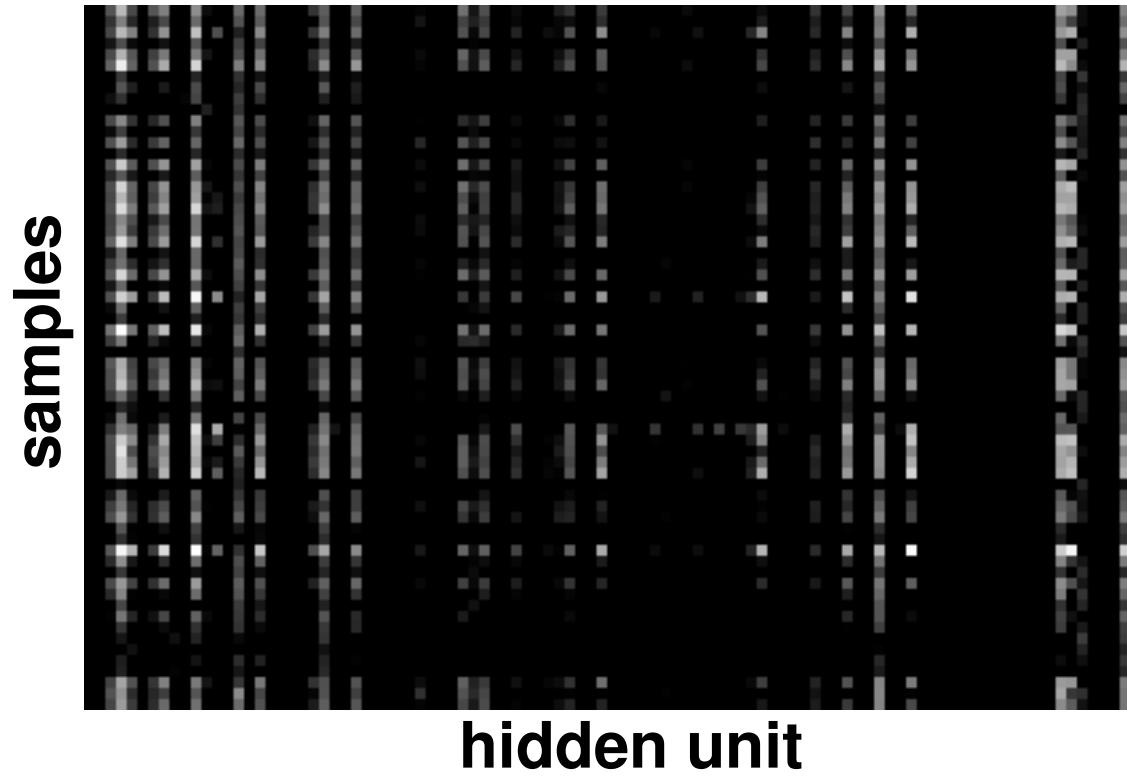
- Check gradients numerically by finite differences
- Visualize features (feature maps need to be uncorrelated) and have high variance.



**Good training:** hidden units are sparse across samples  
and across features.

# OTHER THINGS GOOD TO KNOW

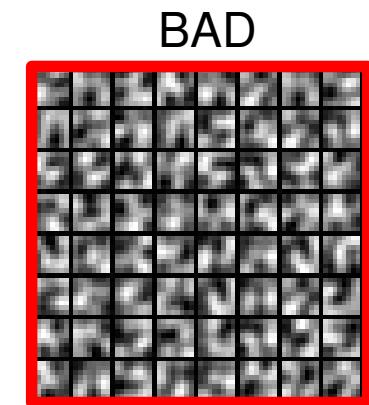
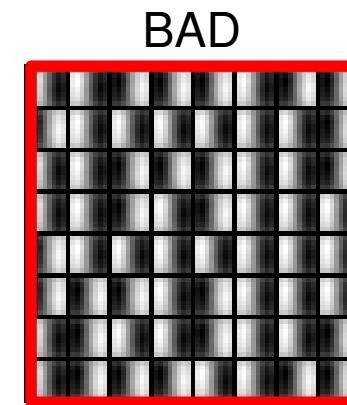
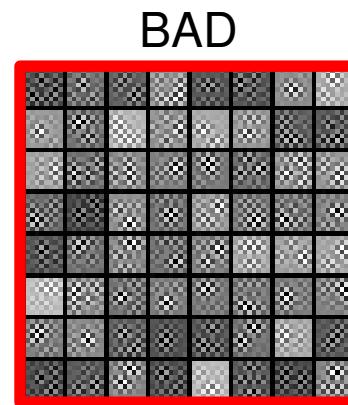
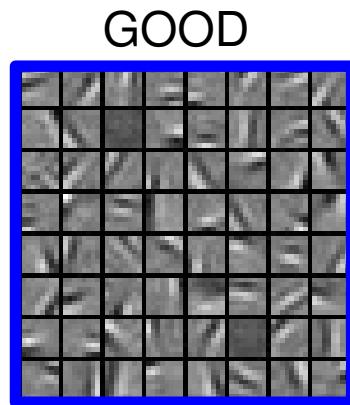
- Check gradients numerically by finite differences
- Visualize features (feature maps need to be uncorrelated) and have high variance.



**Bad training:** many hidden units ignore the input and/or exhibit strong correlations.

# OTHER THINGS GOOD TO KNOW

- Check gradients numerically by finite differences
- Visualize features (feature maps need to be uncorrelated) and have high variance.
- Visualize parameters



**Good training:** learned filters exhibit structure and are uncorrelated.

# OTHER THINGS GOOD TO KNOW

- Check gradients numerically by finite differences
- Visualize features (feature maps need to be uncorrelated) and have high variance.
- Visualize parameters
- Measure error on both training and validation set.
- Test on a small subset of the data and check the error → 0.

# WHAT IF IT DOES NOT WORK?

- Training diverges:
  - Learning rate may be too large → decrease learning rate
  - BPROP is buggy → numerical gradient checking
- Parameters collapse / loss is minimized but accuracy is low
  - Check loss function:
    - Is it appropriate for the task you want to solve?
    - Does it have degenerate solutions?
- Network is underperforming
  - Compute flops and nr. params. → if too small, make net larger
  - Visualize hidden units/params → fix optimization
- Network is too slow
  - Compute flops and nr. params. → GPU,distrib. framework, make net smaller

# FUTURE CHALLENGES

- Scalability
  - Hardware
    - GPU / distributed frameworks
  - Algorithms
    - Better losses
    - Better optimizers
- Learning better representations
  - Video
  - Unsupervised learning
  - Multi-task learning
- Feedback at training and inference time
- Structure prediction
- Black-box tool (hyper-parameters optimization)

# SUMMARY

- Want to efficiently learn non-linear adaptive hierarchical systems
- End-to-end learning
- Gradient-based learning
- Adapting neural nets to vision:
  - Weight sharing
  - Pooling and Contrast Normalization
- Improving generalization on small datasets:
  - Weight decay, dropout, sparsity, multi-task
- Training a convnet means:
  - Design architecture
  - Design loss function
  - Optimization (SGD)
- Very successful (large-scale) applications

# **SOFTWARE**

## **Torch7: learning library that supports neural net training**

---

<http://www.torch.ch>

<http://code.cogbits.com/wiki/doku.php> (tutorial with demos by C. Farabet)

## **Python-based learning library (U. Montreal)**

---

- <http://deeplearning.net/software/theano/> (does automatic differentiation)

## **C++ code for ConvNets (Sermanet)**

---

– <http://eblearn.sourceforge.net/>

## **Efficient CUDA kernels for ConvNets (Krizhevsky)**

---

– [code.google.com/p/cuda-convnet](http://code.google.com/p/cuda-convnet)

# REFERENCES

## Convolutional Nets

---

- LeCun, Bottou, Bengio and Haffner: Gradient-Based Learning Applied to Document Recognition, Proceedings of the IEEE, 86(11):2278-2324, November 1998
- Krizhevsky, Sutskever, Hinton “ImageNet Classification with deep convolutional neural networks” NIPS 2012
- Jarrett, Kavukcuoglu, Ranzato, LeCun: What is the Best Multi-Stage Architecture for Object Recognition?, Proc. International Conference on Computer Vision (ICCV'09), IEEE, 2009
- Kavukcuoglu, Sermanet, Boureau, Gregor, Mathieu, LeCun: Learning Convolutional Feature Hierachies for Visual Recognition, Advances in Neural Information Processing Systems (NIPS 2010), 23, 2010
- see [yann.lecun.com/exdb/publis](http://yann.lecun.com/exdb/publis) for references on many different kinds of convnets.
- see <http://www.cmap.polytechnique.fr/scattering/> for scattering networks (similar to convnets but with less learning and stronger mathematical foundations)

# REFERENCES

## Applications of Convolutional Nets

---

- Farabet, Couprie, Najman, LeCun. Scene Parsing with Multiscale Feature Learning, Purity Trees, and Optimal Covers”, ICML 2012
- Pierre Sermanet, Koray Kavukcuoglu, Soumith Chintala and Yann LeCun: Pedestrian Detection with Unsupervised Multi-Stage Feature Learning, CVPR 2013
- D. Ciresan, A. Giusti, L. Gambardella, J. Schmidhuber. Deep Neural Networks Segment Neuronal Membranes in Electron Microscopy Images. NIPS 2012
- Raia Hadsell, Pierre Sermanet, Marco Scoffier, Ayse Erkan, Koray Kavackuoglu, Urs Muller and Yann LeCun. Learning Long-Range Vision for Autonomous Off-Road Driving, Journal of Field Robotics, 26(2):120-144, 2009
- Burger, Schuler, Harmeling. Image Denoising: Can Plain Neural Networks Compete with BM3D?, CVPR 2012
- Hadsell, Chopra, LeCun. Dimensionality reduction by learning an invariant mapping, CVPR 2006
- Bergstra et al. Making a science of model search: hyperparameter optimization in hundred of dimensions for vision architectures, ICML 2013

# REFERENCES

## Deep Learning in general

---

- deep learning tutorial slides at ICML 2013
- Yoshua Bengio, Learning Deep Architectures for AI, Foundations and Trends in Machine Learning, 2(1), pp.1-127, 2009.
- LeCun, Chopra, Hadsell, Ranzato, Huang: A Tutorial on Energy-Based Learning, in Bakir, G. and Hofman, T. and Schölkopf, B. and Smola, A. and Taskar, B. (Eds), Predicting Structured Data, MIT Press, 2006

# ACKNOWLEDGEMENTS



**Yann LeCun - NYU**



**Alex Krizhevsky - Google**



**Jeff Dean - Google**

# **THANK YOU!**